

# 3D Vision 1

Week 7

Introduction to 3D Vision  
Model Fitting

# Announcements

- Where student lab attendance dropped below a threshold for multiple weeks, the School has taken action to merge labs

ORIGINAL				NEW				
Day	Start Time	End Time	Location	Day	Start Time	End Time	Location	
Wed	13:00	15:00	Rm N113_CSIT Bldg 108	→	Wed	13:00	15:00	Rm N109_CSIT Bldg 108
Wed	14:00	16:00	Rm N112_CSIT Bldg 108	→	Wed	14:00	16:00	Comp Lab 1.23_Hanna Neumann Bldg 145
Wed	15:00	17:00	Rm N109_CSIT Bldg 108	→	Wed	14:00	16:00	Comp Lab 1.23_Hanna Neumann Bldg 145
Fri	11:00	13:00	Rm N113_CSIT Bldg 108	→	Fri	11:00	13:00	Rm N114_CSIT Bldg 108
Fri	13:00	15:00	Comp Lab 1.23_Hanna Neumann Bldg 145	→	Fri	13:00	15:00	Rm N114_CSIT Bldg 108
Fri	13:00	15:00	Rm N113_CSIT Bldg 108	→	Fri	13:00	15:00	Rm N114_CSIT Bldg 108

# Announcements

- **Reminder:** 30-minute tutor drop-in session
  - E.g., for a lab scheduled from 2:00–4:00PM, we have:

---

Time	Activity
2:00PM–2:30PM	Tutor Drop-in Session
2:30PM–4:00PM	Lab

---

# Announcements

- Tutorial 5, on backpropagation, is available on Wattle and will be covered during the labs this week
- Assignment 2 due 11:59pm Friday 26 April (end of Week 8)
  - This includes a one week extension that has already been applied
  - **Zero marks** if either report or code submitted late (unless extension)
    - Submit early; you can always resubmit an updated version later
    - Depending on your internet connection and load on the TurnItIn servers, uploading can sometimes be slow, so please factor this into your submission schedule
  - Submit your report (PDF) and code (ZIP file) **separately under the correct tab** in the submission box
  - Follow the instructions under Submission Requirements

# Weekly Study Plan: Overview

Wk	Starting	Lecture	Lab	Assessment
1	19 Feb	Introduction	X	
2	26 Feb	Low-level Vision 1	1	
3	4 Mar	Low-level Vision 2	1	
		Mid-level Vision 1		
4	11 Mar	Mid-level Vision 2	1	CLab1 report due Friday
		High-level Vision 1		
5	18 Mar	High-level Vision 2	2	
6	25 Mar	High-level Vision 3 <sup>1</sup>	2	
	1 Apr	Teaching break	X	
	8 Apr	Teaching break	X	
7	15 Apr	3D Vision 1	2	CLab2 report due Friday
8	22 Apr	3D Vision 2	3	
9	29 Apr	3D Vision 3	3	
10	6 May	3D Vision 4	3	
		Mid-level Vision 3		
11	13 May	High-level Vision 4	X	CLab3 report due Friday
12	20 May	Course Review	X	

# Weekly Study Plan: Part B

Wk	Starting	Lecture	By
7	15 Apr	3D vision: introduction, camera model, single-view geometry	Dylan
8	22 Apr	3D vision: camera calibration, two-view geometry (homography)	Dylan
9	29 Apr	3D vision: two-view geometry (epipolar geometry, triangulation, stereo)	Dylan
10	6 May	3D vision: multiple-view geometry	Weijian
		Mid-level vision: optical flow, shape-from-X	Dylan
11	13 May	High-level vision: self-supervised learning, detection, segmentation	Dylan
12	20 May	Course review	Dylan

# Outline

1. Introduction to 3D vision
2. Model fitting (line fitting)
  1. Least squares
  2. M-estimation
  3. RANSAC
  4. Hough transform
3. Image formation (review): pinhole camera model
4. Camera projection matrix and single view geometry
5. Camera calibration
6. Resectioning and camera pose

# Introduction to 3D Vision

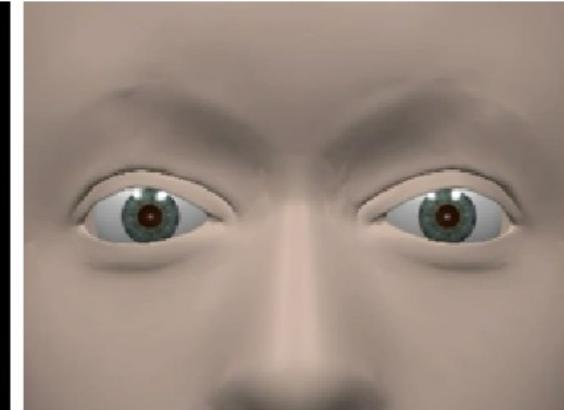
# 3D Vision: Eight Human Tricks



binocular stereo



accommodation



vergence



motion perspective



occlusion



apparent height



texture density



aerial perspective

Image Credit: Peter Corke

# 3D Vision: Eight Nine Human Tricks



Perspective Distortion

Image Credit: Pikist

# 3D Vision: Eight Human Tricks

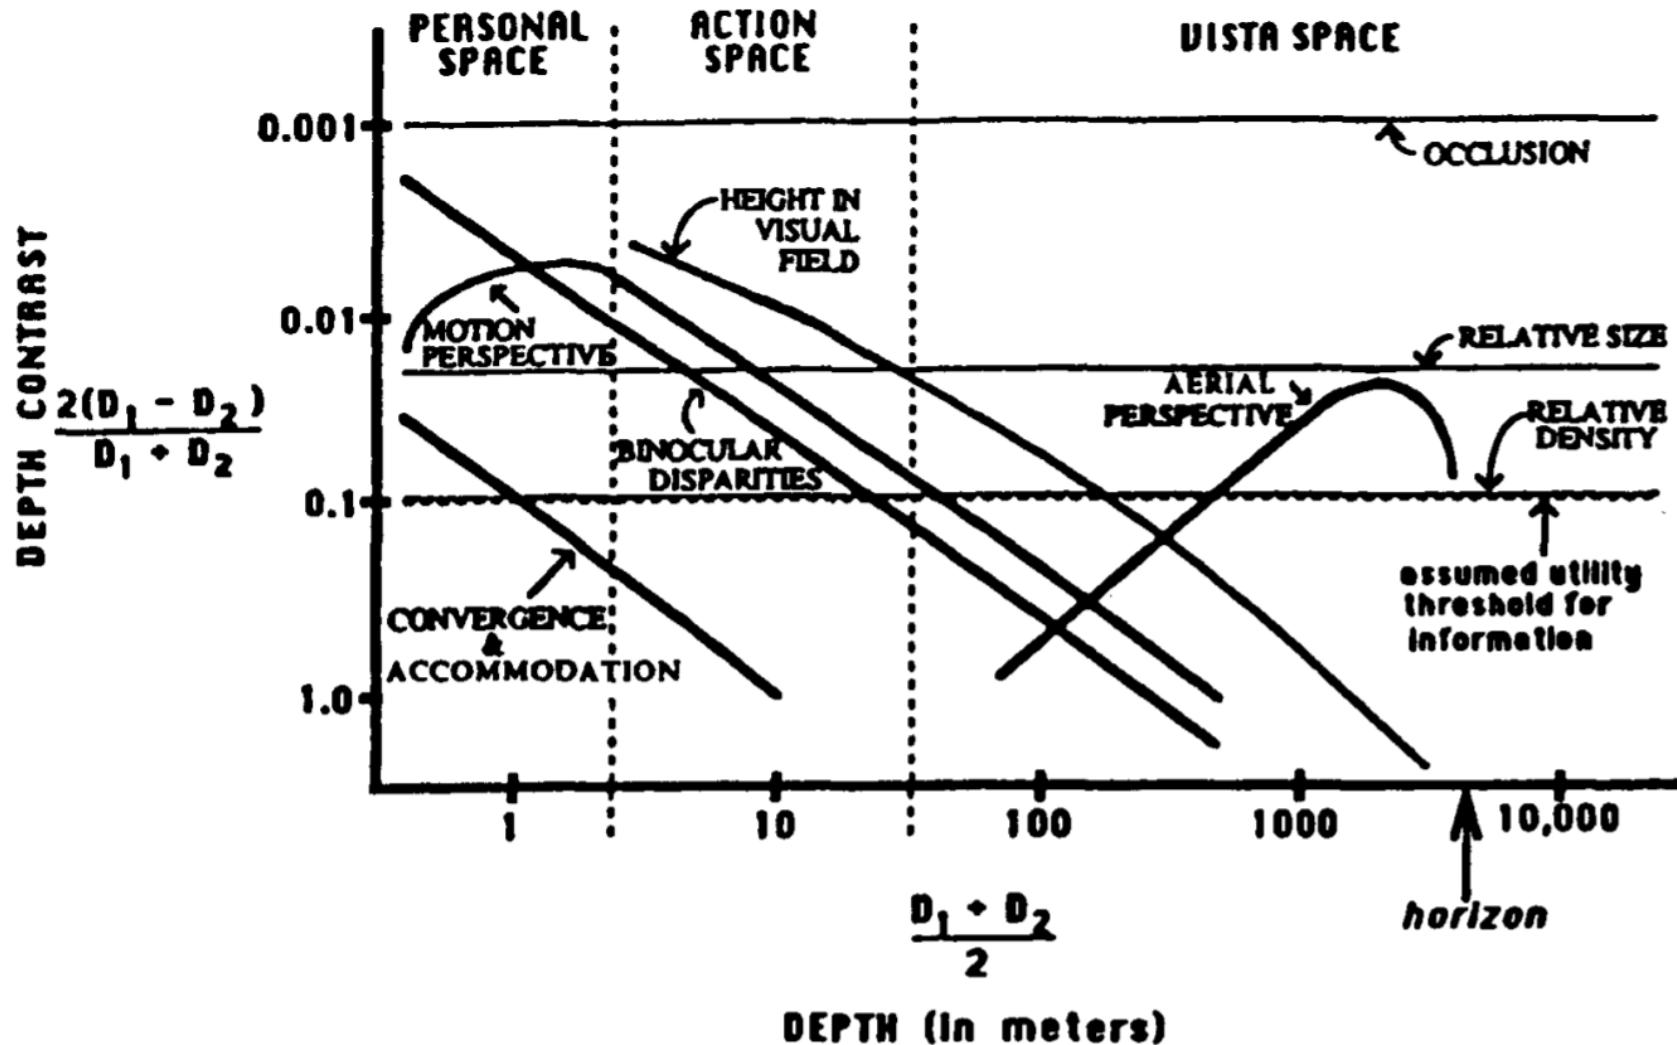
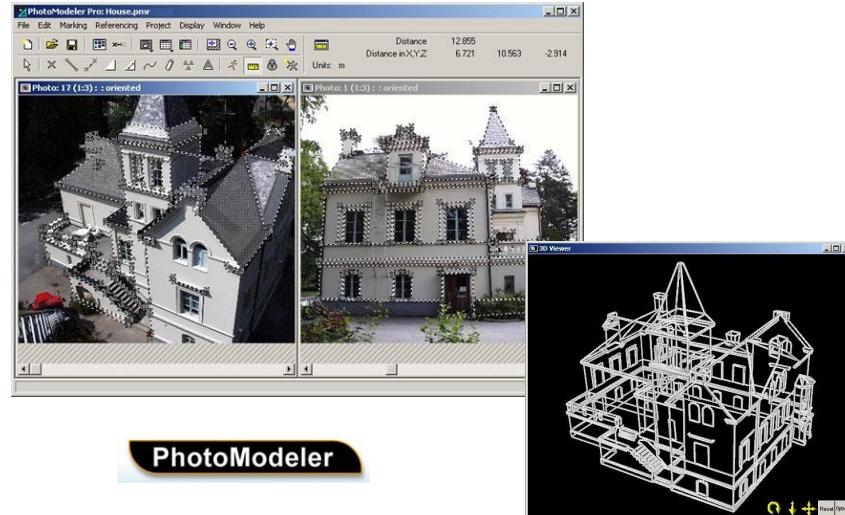


Image Credit: J. Cutting, "How the eye measures reality and virtual reality", 1997

# Why 3D Vision?

## Structure-from-Motion (SfM):

- Input: set or sequence of images
- Output: camera poses, 3D model
- Applications:
  - Architecture (survey, stability analysis, renovation planning)
  - City modelling
  - Archaeology
  - Inspection (from drones)
  - Surgery



# Building Rome in a Day



Credit: S. Agarwal et al., "Building Rome in a day", 2009

# NeRF



Credit: J. Barron et al., "Zip-NeRF", ICCV 2023

# DUST3R

- <https://dust3r.europe.naverlabs.com/>

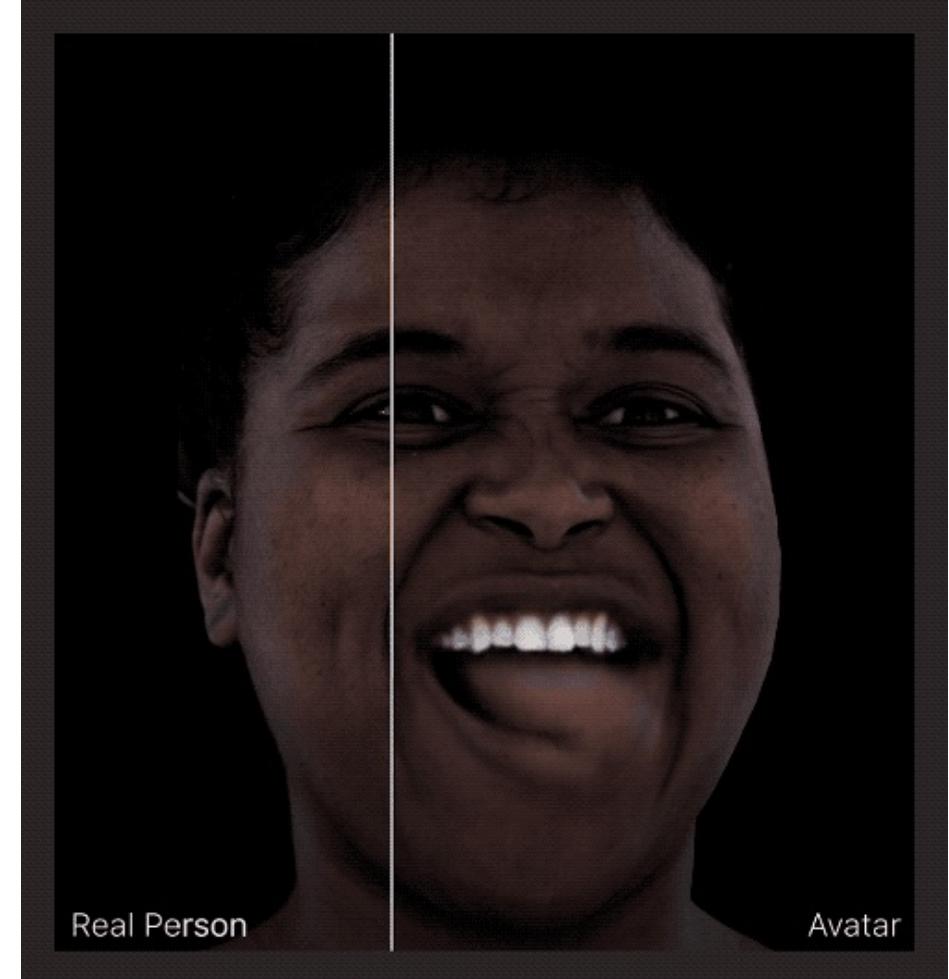
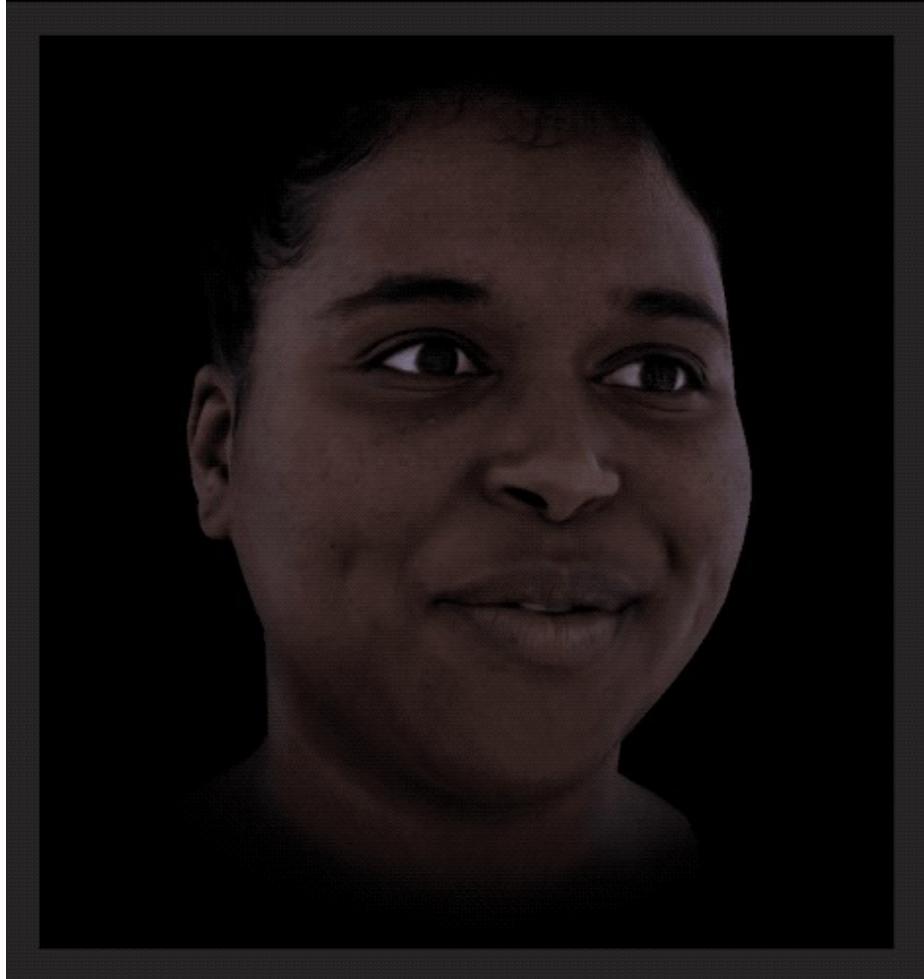
# Telepresence



Credit: Microsoft HoloLens

16

# Telepresence

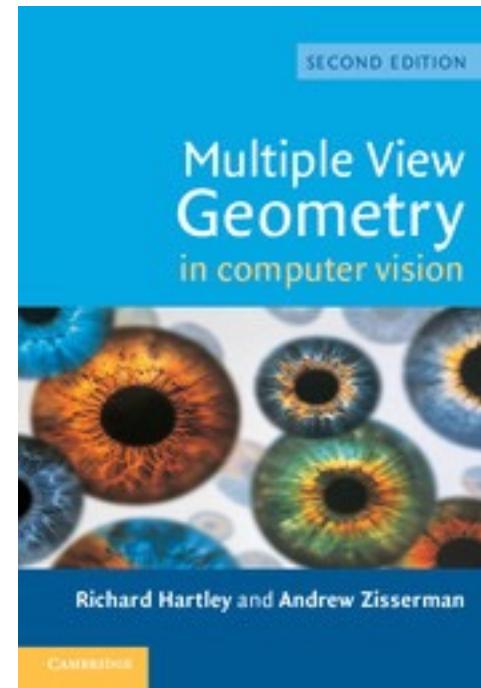
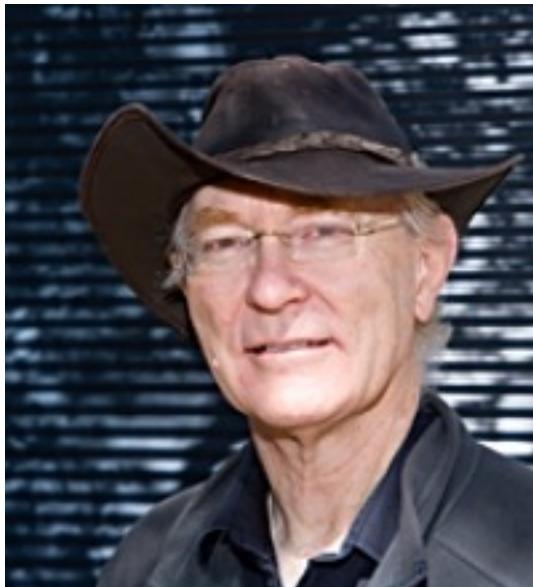


Credit: Facebook (Meta)

17

# Reference Textbook for 3D Vision

- Multiple View Geometry in Computer Vision



Richard Hartley (ANU) | Andrew Zisserman (Oxford)

# Recommended Reading

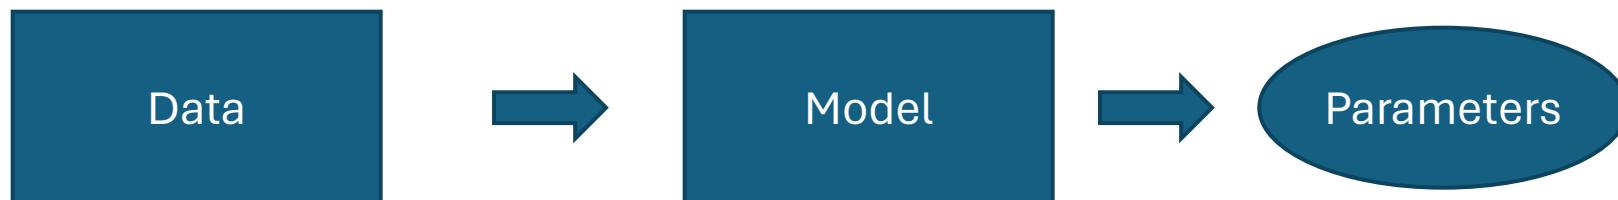
- Ch 6: Camera models
  - Ch 7: Computation of camera matrix
  - Ch 4: DLT algorithm
- 
- More in-depth than the lectures
  - Goes beyond what is required for this course

# Model Fitting

Line Fitting

# Model Fitting in Computer Vision

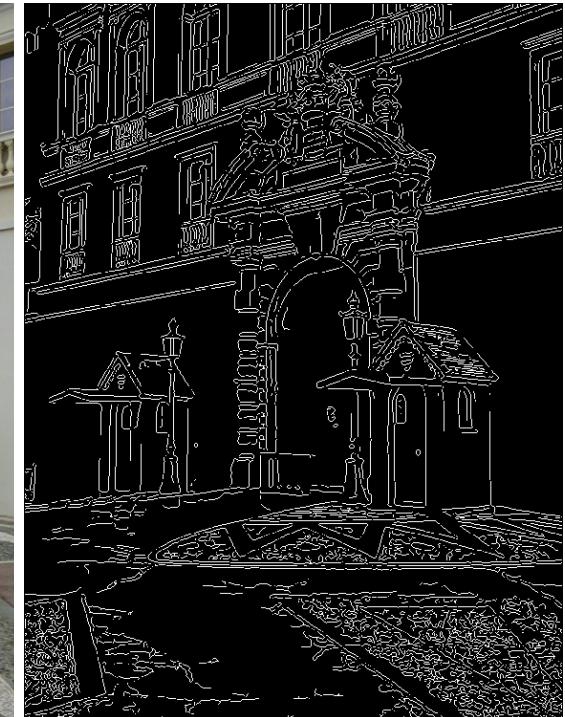
- Model fitting is a fundamental approach in computer vision



- Find model parameters that best fit the data
  - → Optimisation problem
- Model can be as simple as 2D lines, or more complex such as camera transform or 3D articulated object (e.g., skeleton model)
- We will introduce using 2D lines now, cameras later

# A Binary Edge Map is Often Insufficient

- Fragmented edges
- Many broken lines
- Noise in edge orientations
- Membership not clear: which points on which line?



# Any Feature Here?

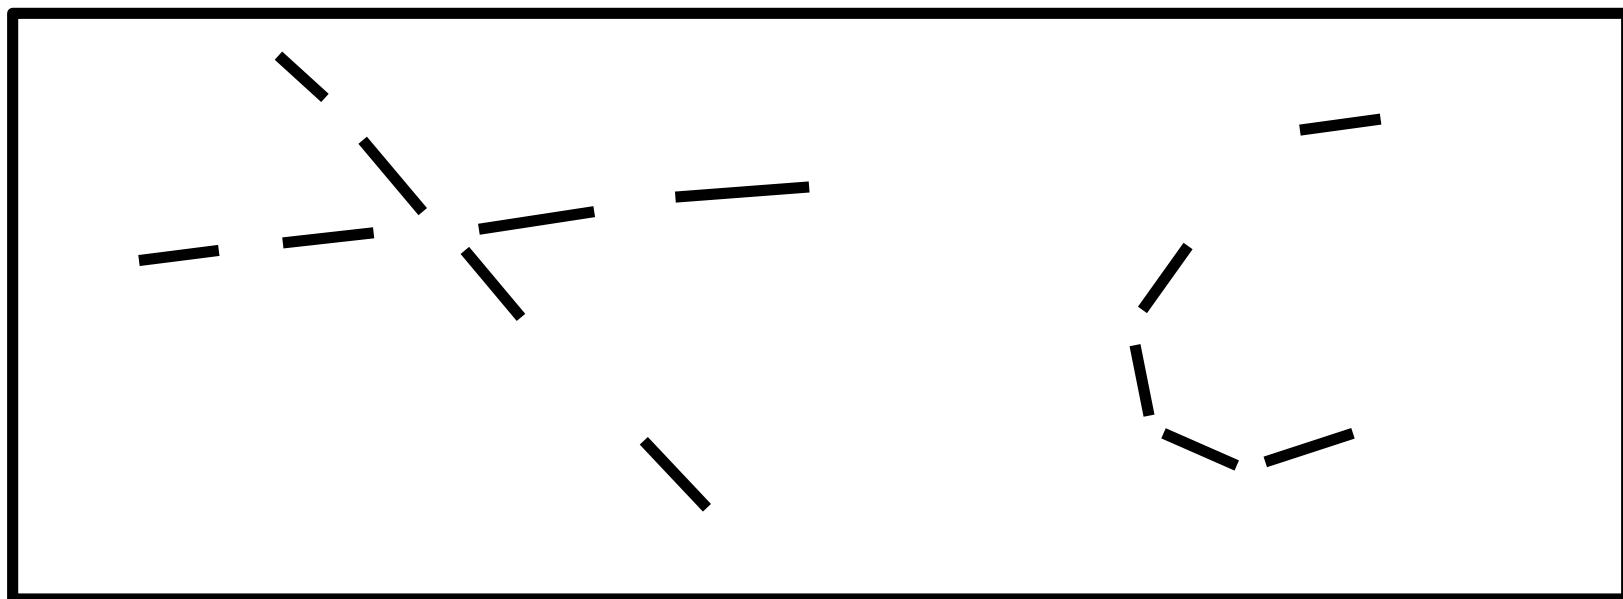


# Vanishing Points and Lines



# Goal of Line Fitting

- From fragmented edges to straight lines or curves



# Goal of Line Fitting

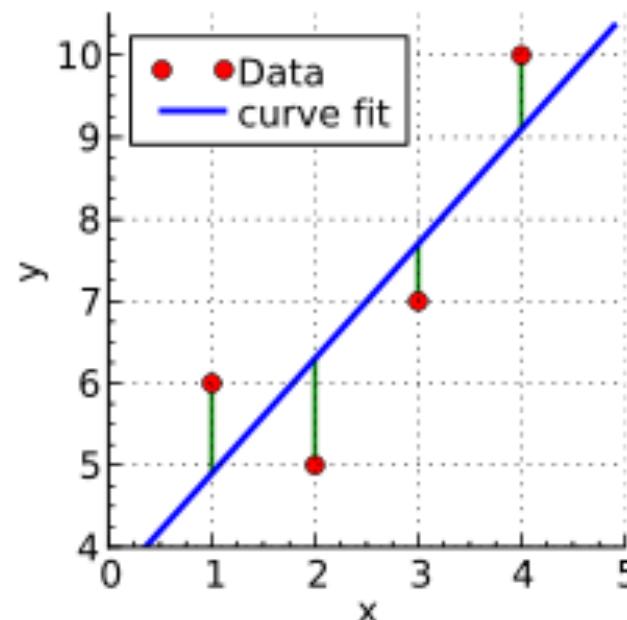


# “Fitting” in general context

- Choose a parametric model to represent a set of features
- Membership determination
  - Decide which feature belongs to which model
  - Not a local decision; must look at the full image globally
- Key questions:
  - What is the best parametric model?
  - How many models (e.g., lines) are there?
  - Which point belongs to which line?
  - Remove noise points, wrong edge points

# Basic Line Fitting: Least Squares

- Given a set of points that belong to a straight line, find the equation of the line
- Find the line parameters that yield minimal error in a least-square sense



# Basic Line Fitting: Least Squares

- Given data:  $(x_1, y_1), \dots, (x_n, y_n)$
- Line model:  $y_i = mx_i + b$
- Find  $\mathbf{p} = (m, b)$  to minimise

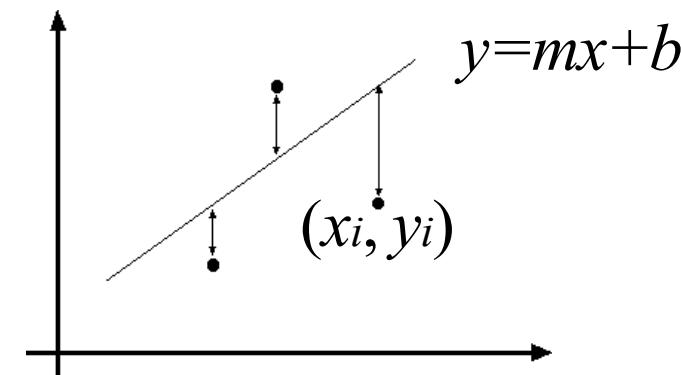
$$e = \sum_{i=1}^n (y_i - mx_i - b)^2$$

$$e = \sum_{i=1}^n \left( y_i - [x_i \ 1] \begin{bmatrix} m \\ b \end{bmatrix} \right)^2$$

$$e = \left\| \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} - \begin{bmatrix} x_1 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \right\|^2$$

$$e = \| \mathbf{y} - A\mathbf{p} \|^2$$

$$e = \mathbf{y}^T \mathbf{y} - 2(\mathbf{A}\mathbf{p})^T \mathbf{y} + (\mathbf{A}\mathbf{p})^T (\mathbf{A}\mathbf{p})$$

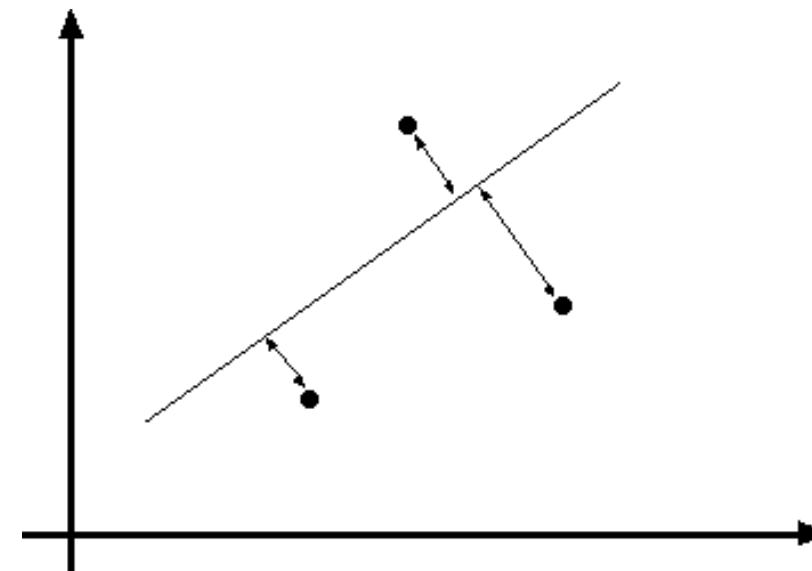
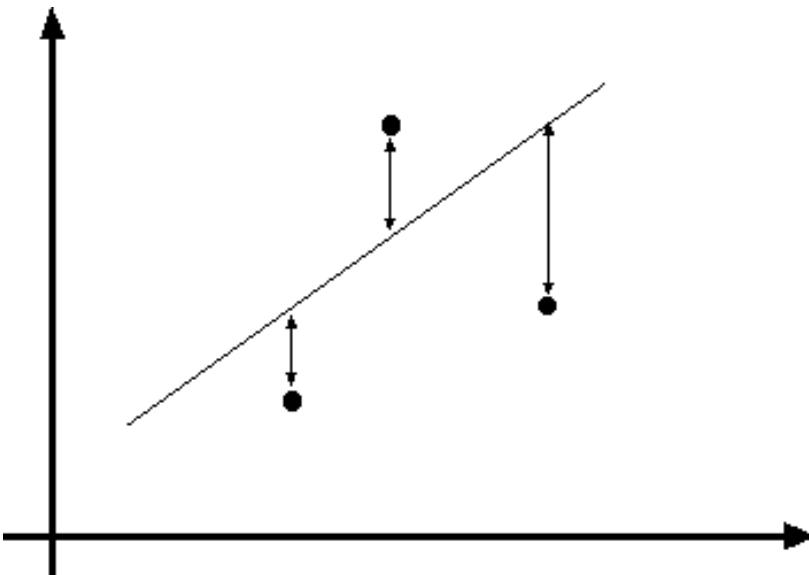


$$\frac{de}{d\mathbf{p}} = 2\mathbf{A}^T \mathbf{A}\mathbf{p} - 2\mathbf{A}^T \mathbf{y} = 0$$

$$\therefore \mathbf{p} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}$$

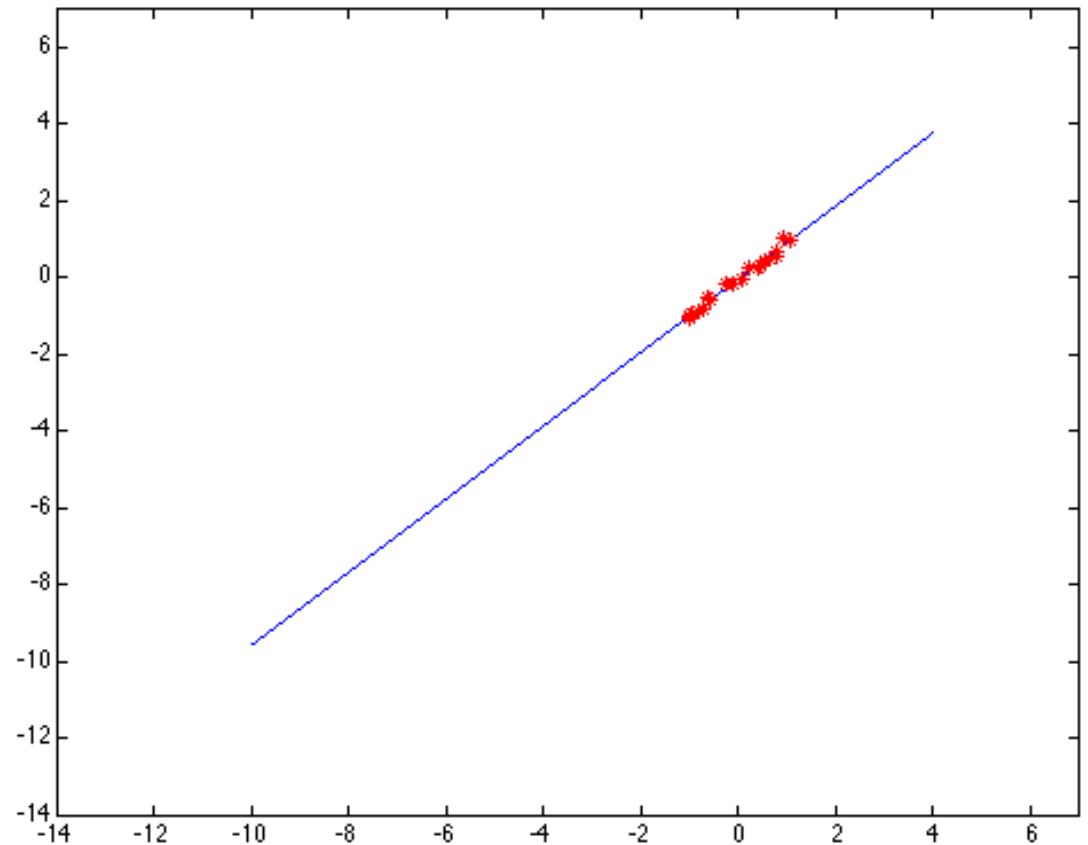
# Basic Line Fitting: Issues

- Fitting performance is affected by the slope
- Fails completely for vertical lines
- Modified least squares using a perpendicular distance



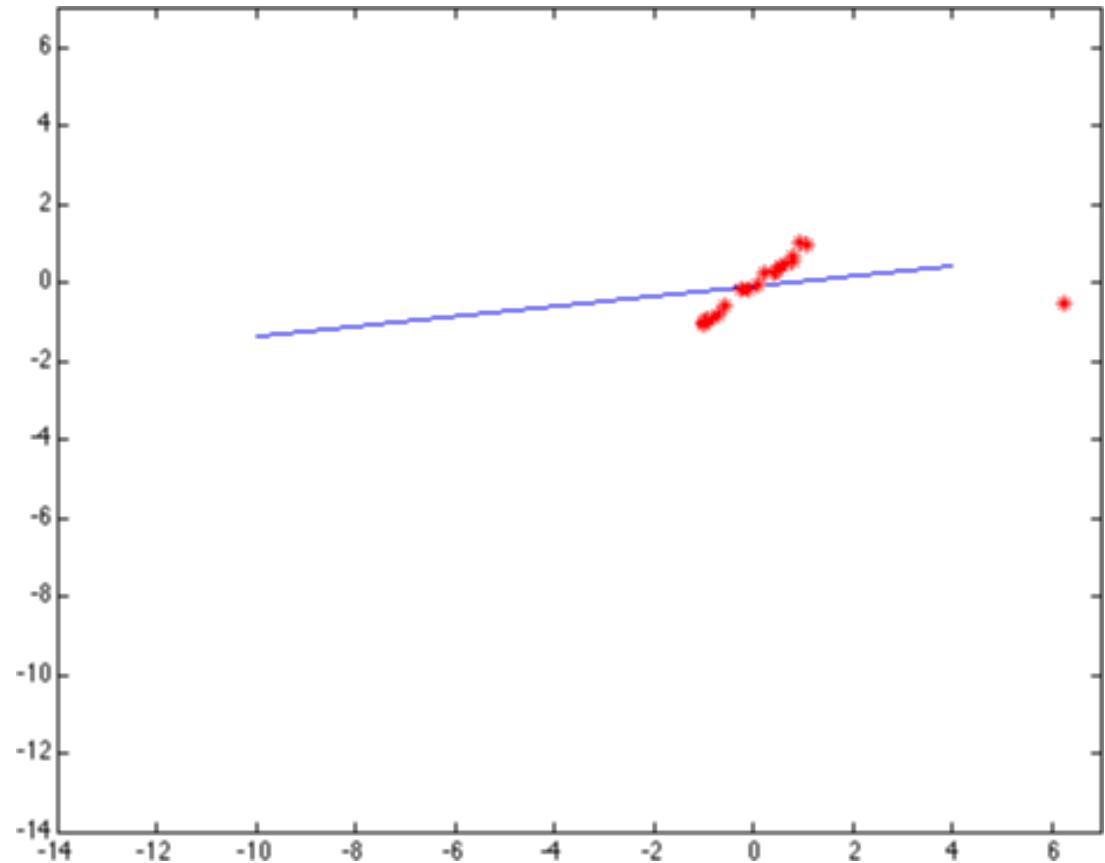
# Least Squares: Robustness to Noise

- Least squares fit to the red points



# Least Squares: Robustness to Noise

- Least squares fit with an outlier
- Problem: squared error heavily penalises outliers



# Robust M-estimators

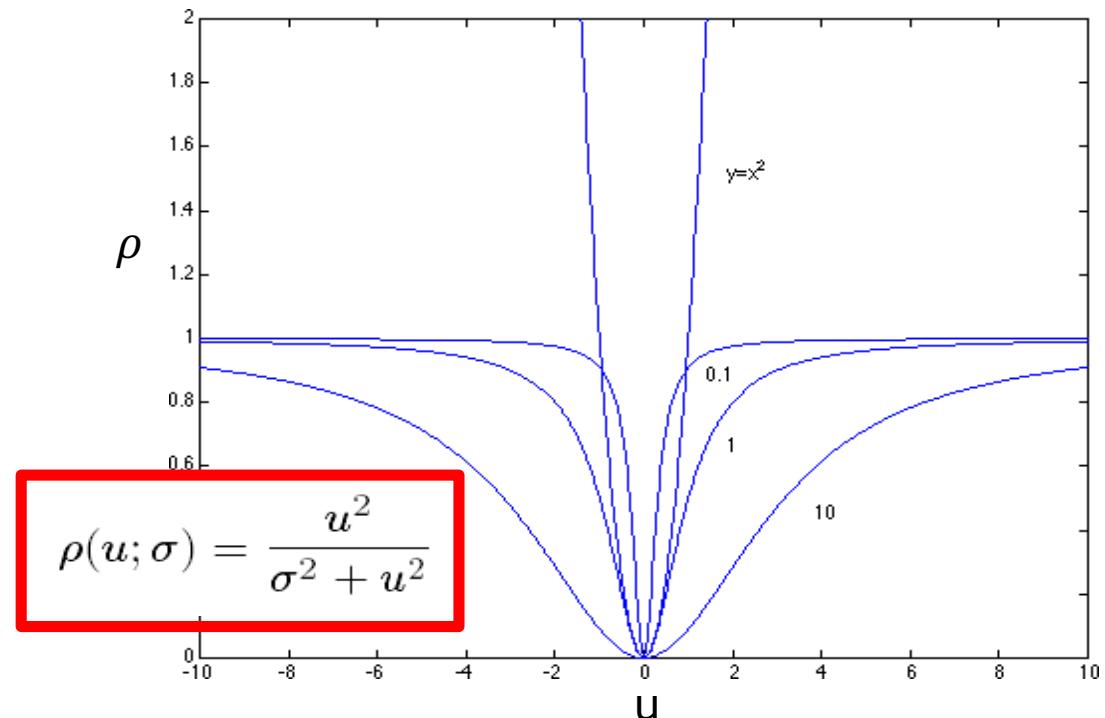
- General approach: find model parameters  $\theta$  that minimise

$$\sum_i \rho(r_i(x_i, \theta); \sigma)$$

where:

- $r_i$  is the residual/loss of the  $i$ th point w.r.t. model parameters  $\theta$
- $\rho$  is a robust function with scale parameter  $\sigma$

- Robust function behaves like squared distance for small values of the residual but saturates for larger values



# Robust M-estimators

- General approach: find model parameters  $\theta$  that minimise

$$\sum_i \rho(r_i(x_i, \theta); \sigma)$$

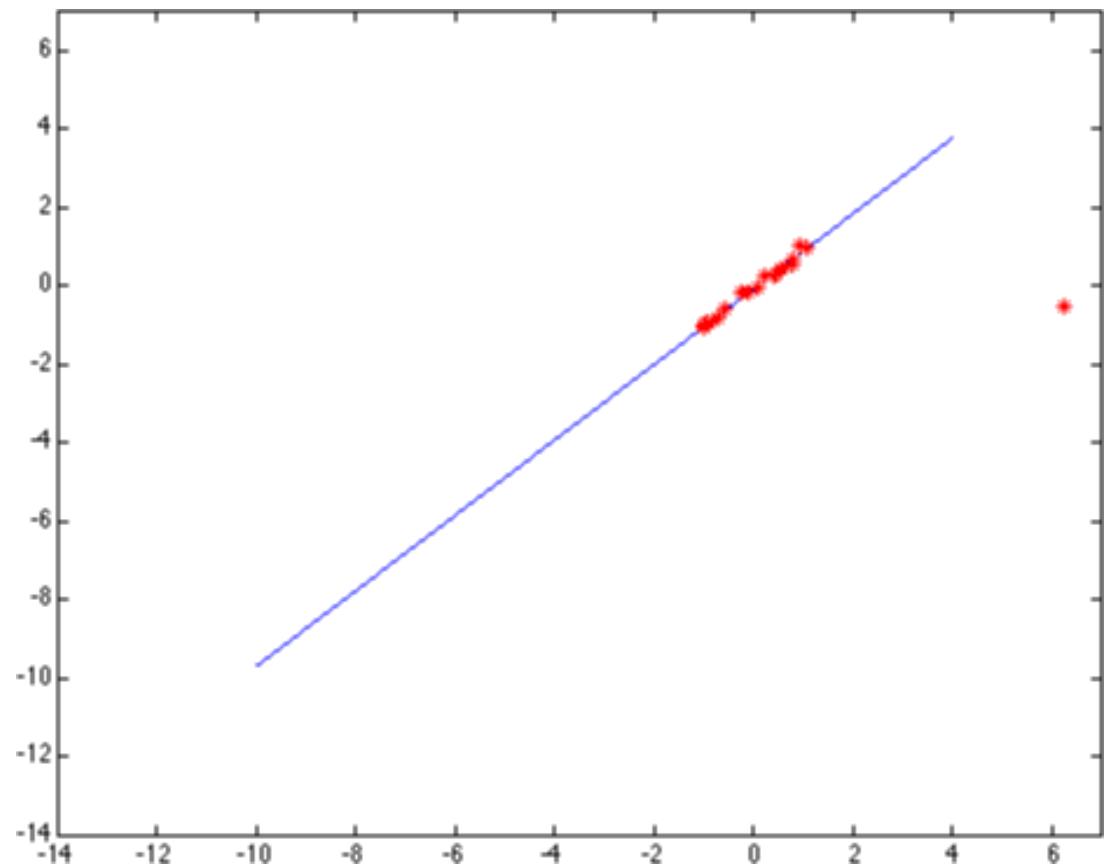
where:

- $r_i$  is the residual/loss of the  $i$ th point w.r.t. model parameters  $\theta$
- $\rho$  is a robust function with scale parameter  $\sigma$

- Robust fitting is a nonlinear optimisation problem that must be solved iteratively
- Least squares solution can be used for initialisation
- Scale of the robust function should be chosen carefully
  - You need to know something about the noise!

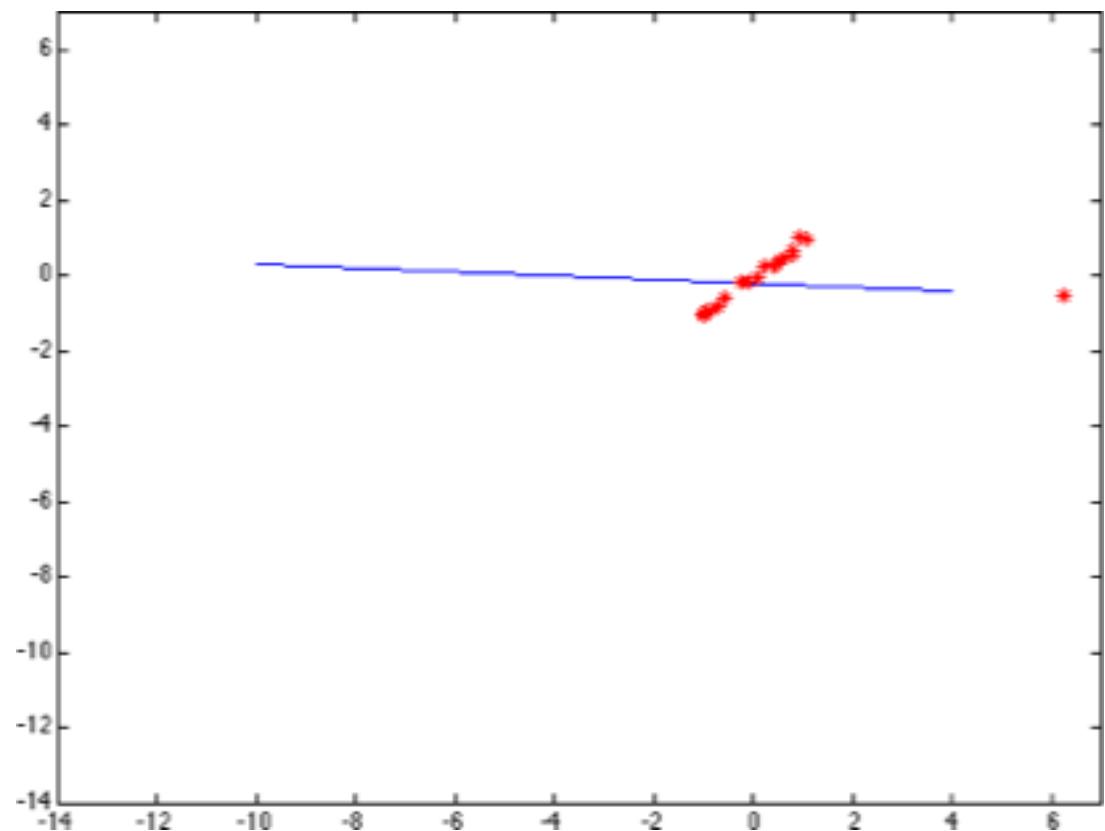
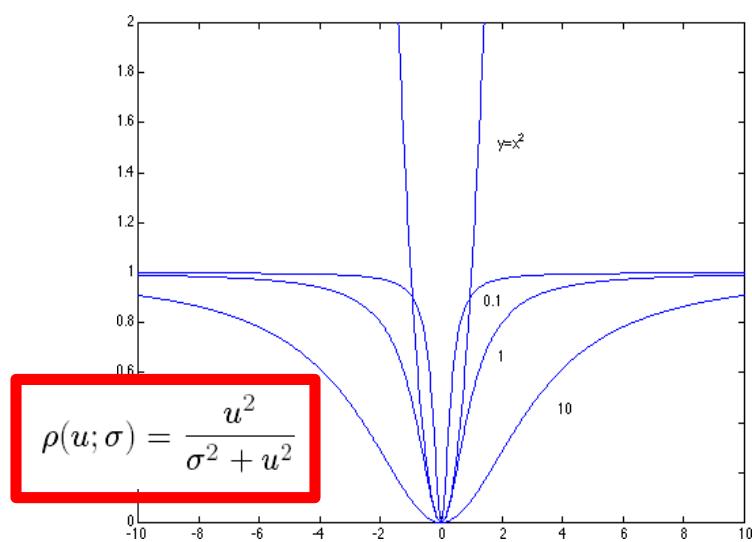
# Choosing the Scale: Just Right

- The effect of the outlier is minimised



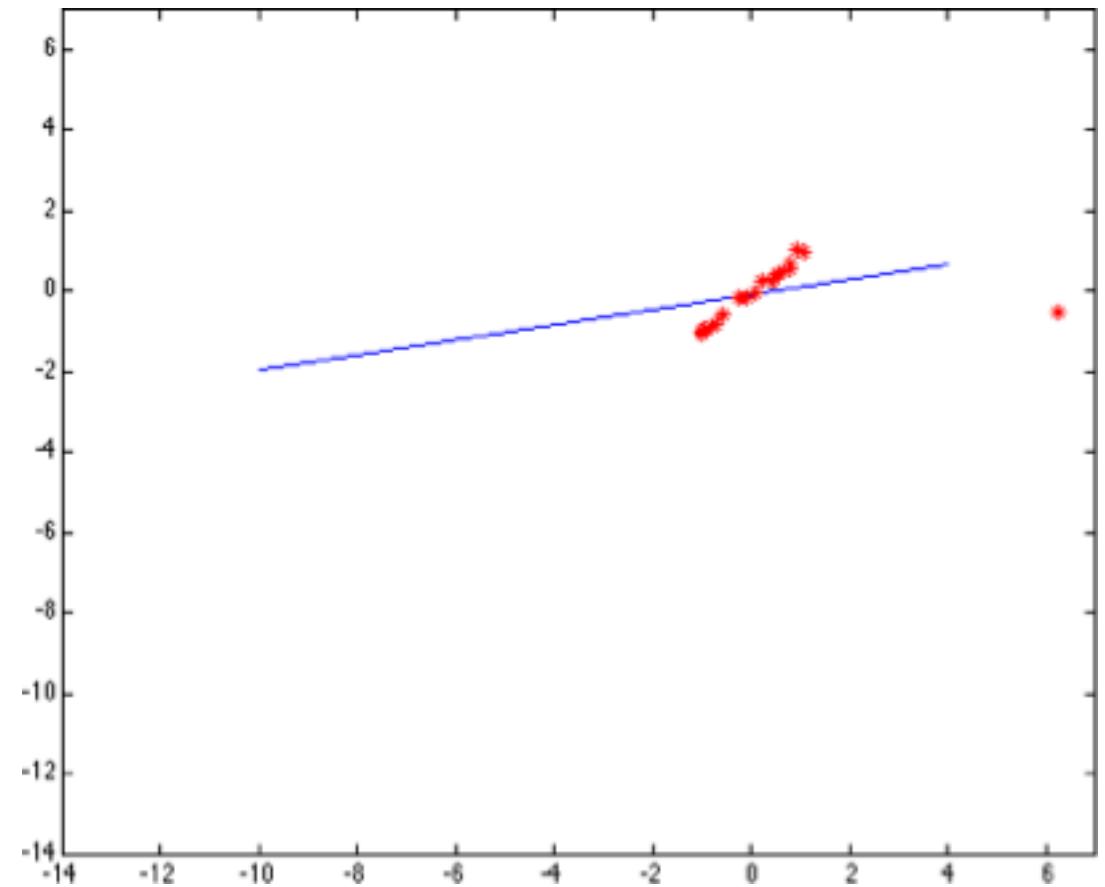
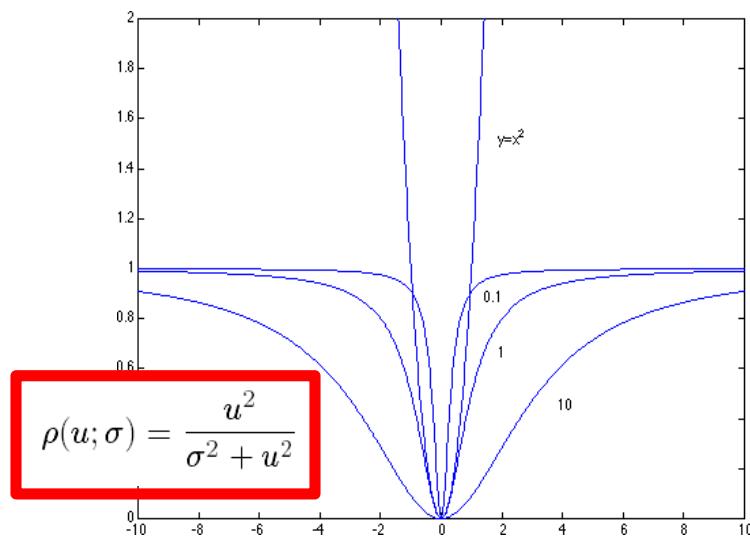
# Choosing the Scale: Too Small

- The error value is almost the same for every point and the fit is very poor



# Choosing the Scale: Too Large

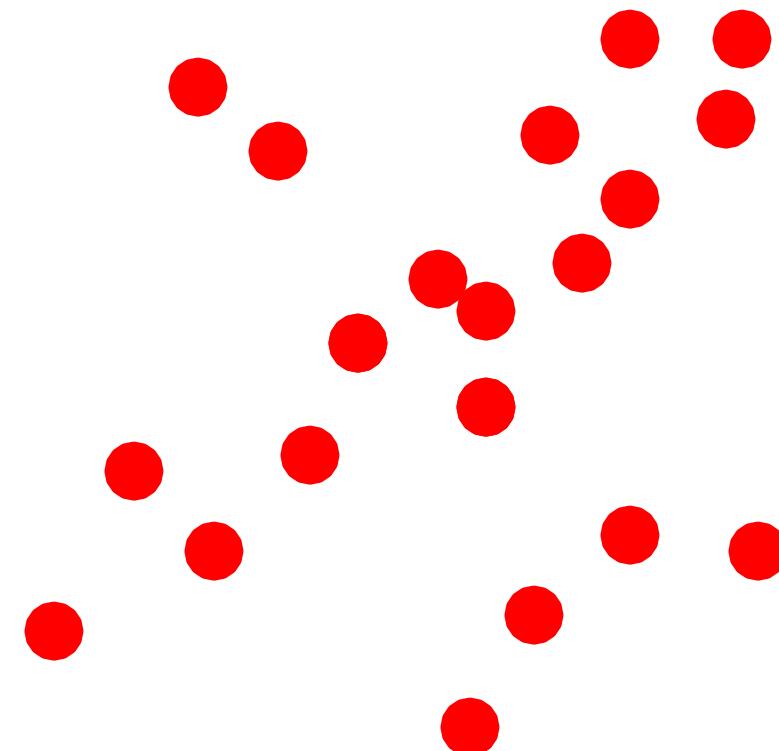
- Behaves much the same as least squares



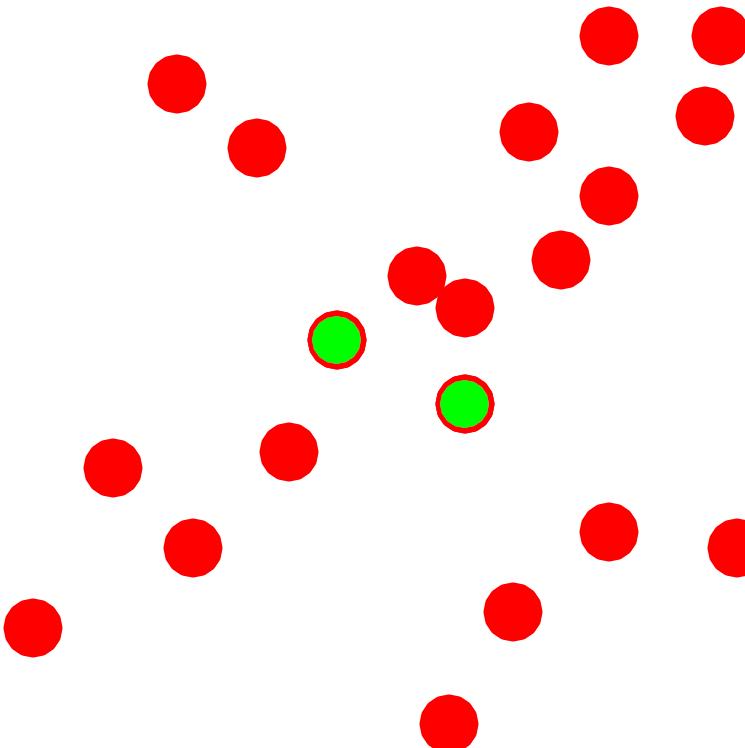
# RANSAC

# RANSAC: RANdom SAmple Consensus

- Fischler & Bolles 1981
- Basic idea (for line fitting):
  1. Randomly select two samples (minimal number for a line)
  2. Check how many other samples fall close to the line (within a distance threshold)
  3. Choose the line which has the largest number (“consensus”)



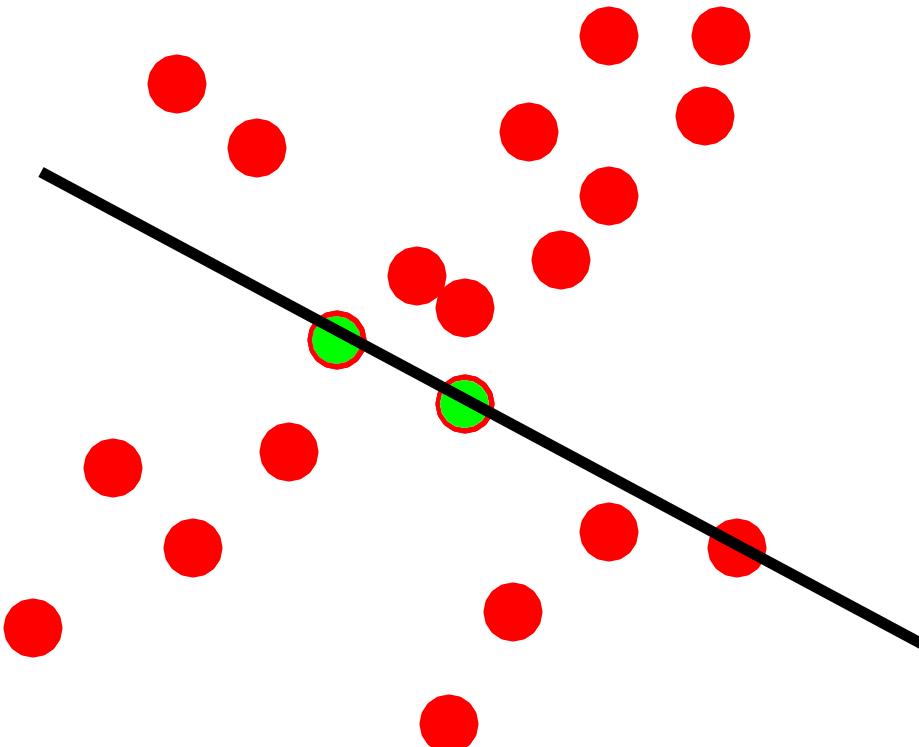
# RANSAC



Algorithm:

- 1. Sample** (randomly) the minimum number of points required to fit the model ( $n=2$ )
  - 2. Solve** for model parameters
  - 3. Score** by the fraction of inliers within a preset threshold
- Repeat** 1–3 until the best model is found

# RANSAC



Algorithm:

1. **Sample** (randomly) the minimum number of points required to fit the model ( $n=2$ )
  2. **Solve** for model parameters
  3. **Score** by the fraction of inliers within a preset threshold
- Repeat** 1–3 until the best model is found

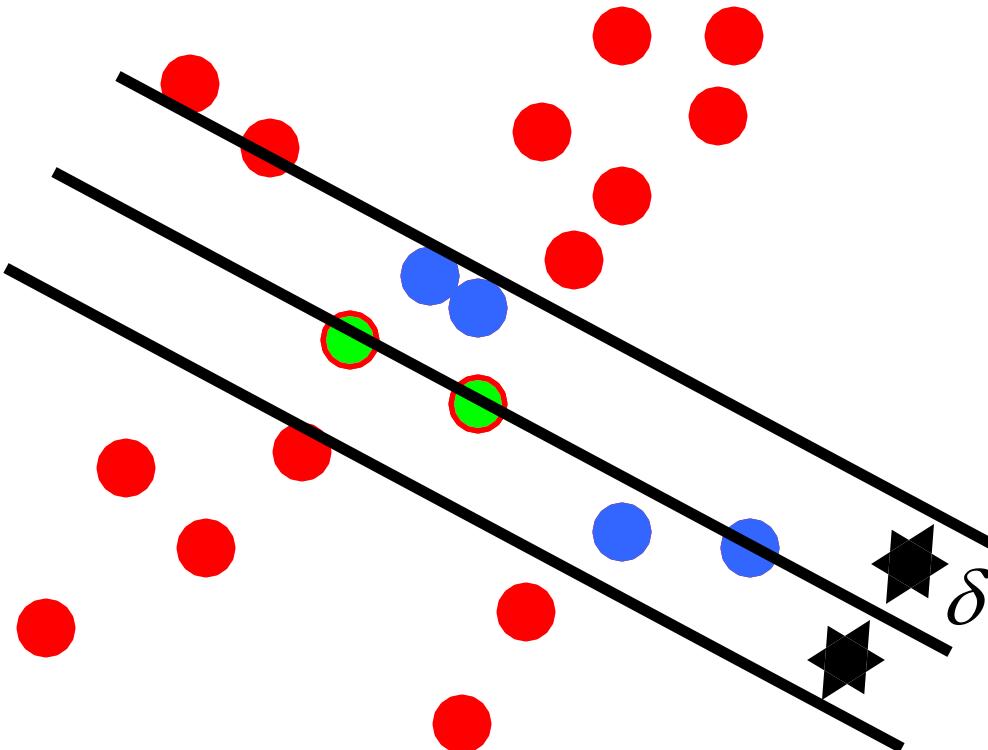
# RANSAC

$$N_I = 6$$

Algorithm:

1. **Sample** (randomly) the minimum number of points required to fit the model ( $n=2$ )
2. **Solve** for model parameters
3. **Score** by the fraction of inliers within a preset threshold

**Repeat** 1–3 until the best model is found



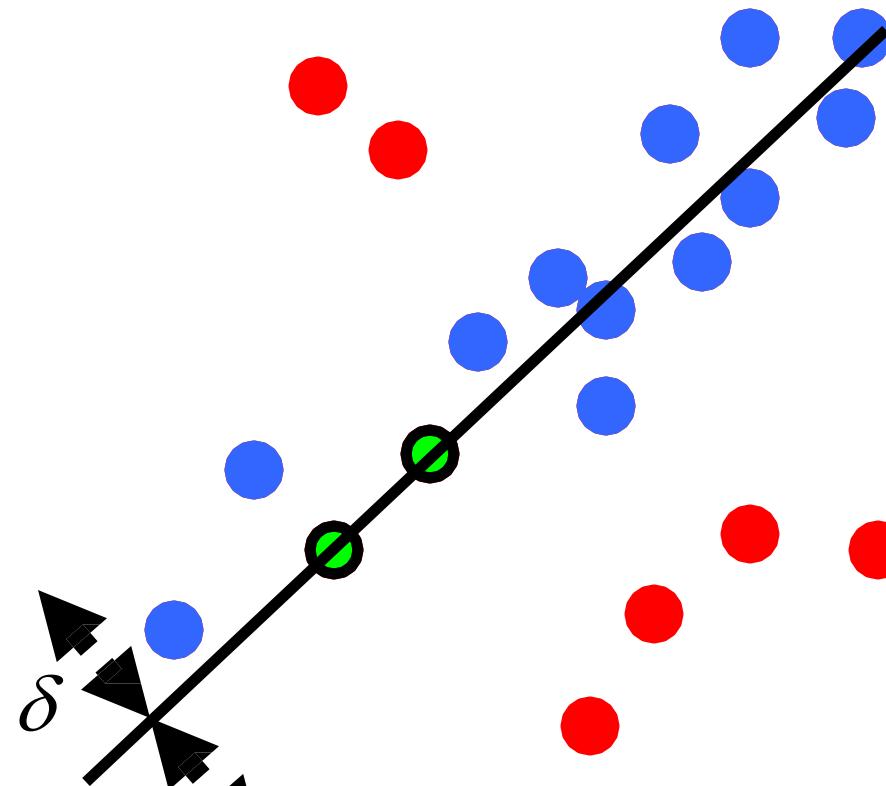
# RANSAC

$$N_I = 14$$

Algorithm:

1. **Sample** (randomly) the minimum number of points required to fit the model ( $n=2$ )
2. **Solve** for model parameters
3. **Score** by the fraction of inliers within a preset threshold

**Repeat** 1–3 until the best model is found



# RANSAC: How Many Samples to Choose?

- $e$ : probability that a point is an outlier  
 $s$ : number of points in a sample  
 $N$ : number of samples (we want to compute this)  
 $p$ : desired probability that we get a good sample

$$1 - \left(1 - (1 - e)^s\right)^N = p$$

- Probability of choosing  $s$  inliers (sample only contains inliers)

# RANSAC: How Many Samples to Choose?

- $e$ : probability that a point is an outlier  
 $s$ : number of points in a sample  
 $N$ : number of samples (we want to compute this)  
 $p$ : desired probability that we get a good sample

$$1 - (1 - (1 - e)^s)^N = p$$

- Probability that one or more points in the sample were outliers (sample is contaminated)

# RANSAC: How Many Samples to Choose?

- $e$ : probability that a point is an outlier  
 $s$ : number of points in a sample  
 $N$ : number of samples (we want to compute this)  
 $p$ : desired probability that we get a good sample

$$1 - (1 - (1 - e)^s)^N = p$$

- Probability that all the  $N$  samples were contaminated

# RANSAC: How Many Samples to Choose?

- $e$ : probability that a point is an outlier  
 $s$ : number of points in a sample  
 $N$ : number of samples (we want to compute this)  
 $p$ : desired probability that we get a good sample

$$1 - (1 - (1 - e)^s)^N = p$$

- Probability that at least one sample is good

# RANSAC: How Many Samples to Choose?

- $e$ : probability that a point is an outlier  
 $s$ : number of points in a sample  
 $N$ : number of samples (we want to compute this)  
 $p$ : desired probability that we get a good sample

$$\begin{aligned}1 - (1 - (1 - e)^s)^N &= p \\ \therefore (1 - (1 - e)^s)^N &= 1 - p \\ \therefore N \log(1 - (1 - e)^s) &= \log(1 - p) \\ \therefore N = \frac{\log(1 - p)}{\log(1 - (1 - e)^s)} &= 16.008 \rightarrow 17\end{aligned}$$

- For a 99% probability of obtaining a good model, with a 50% outlier probability and  $s=2$  samples per model

Sample size	Proportion of outliers						
	<b>N</b>	<b>5%</b>	<b>10%</b>	<b>20%</b>	<b>25%</b>	<b>30%</b>	<b>40%</b>
<b>2</b>	2	3	5	6	7	11	17
<b>3</b>	3	4	7	9	11	19	35
<b>4</b>	3	5	9	13	17	34	72
<b>5</b>	4	6	12	17	26	57	146
<b>6</b>	4	7	16	24	37	97	293
<b>7</b>	4	8	20	33	54	163	588
<b>8</b>	5	9	26	44	78	272	1177

**Figure Credit:** Hartley & Zisserman

# RANSAC: Pros and Cons

- Pros:
  - Robust to outliers
  - Applicable for large number of parameters
- Cons:
  - Computational time grows quickly with the fraction of outliers and the number of parameters
  - Not good for getting multiple line fits
  - Can rapidly find a good model, not necessarily a great model
- Common applications:
  - Computing a rotation and translation between two images (will be used later for fundamental matrix in 2- view geometry)

# Hough Transform

# Fitting Multiple Lines Using a Hough Transform

- Given a binary edge image, find the lines (or curves) that explain the data points best in the parameter space
- This parameter space is called a Hough space

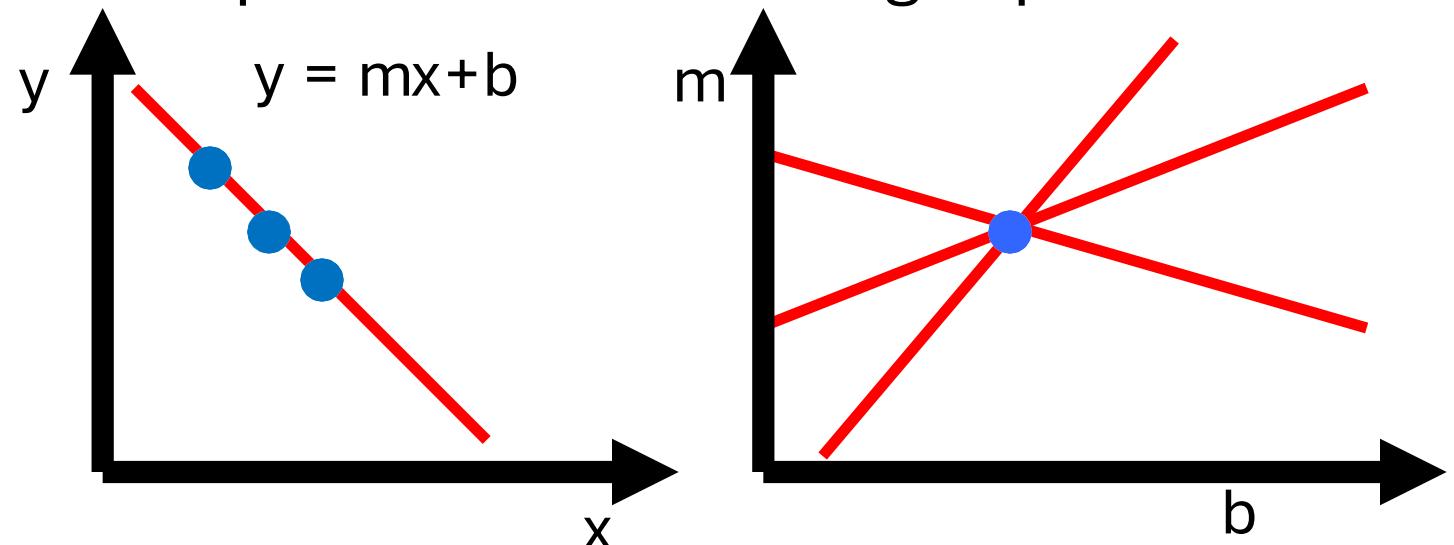


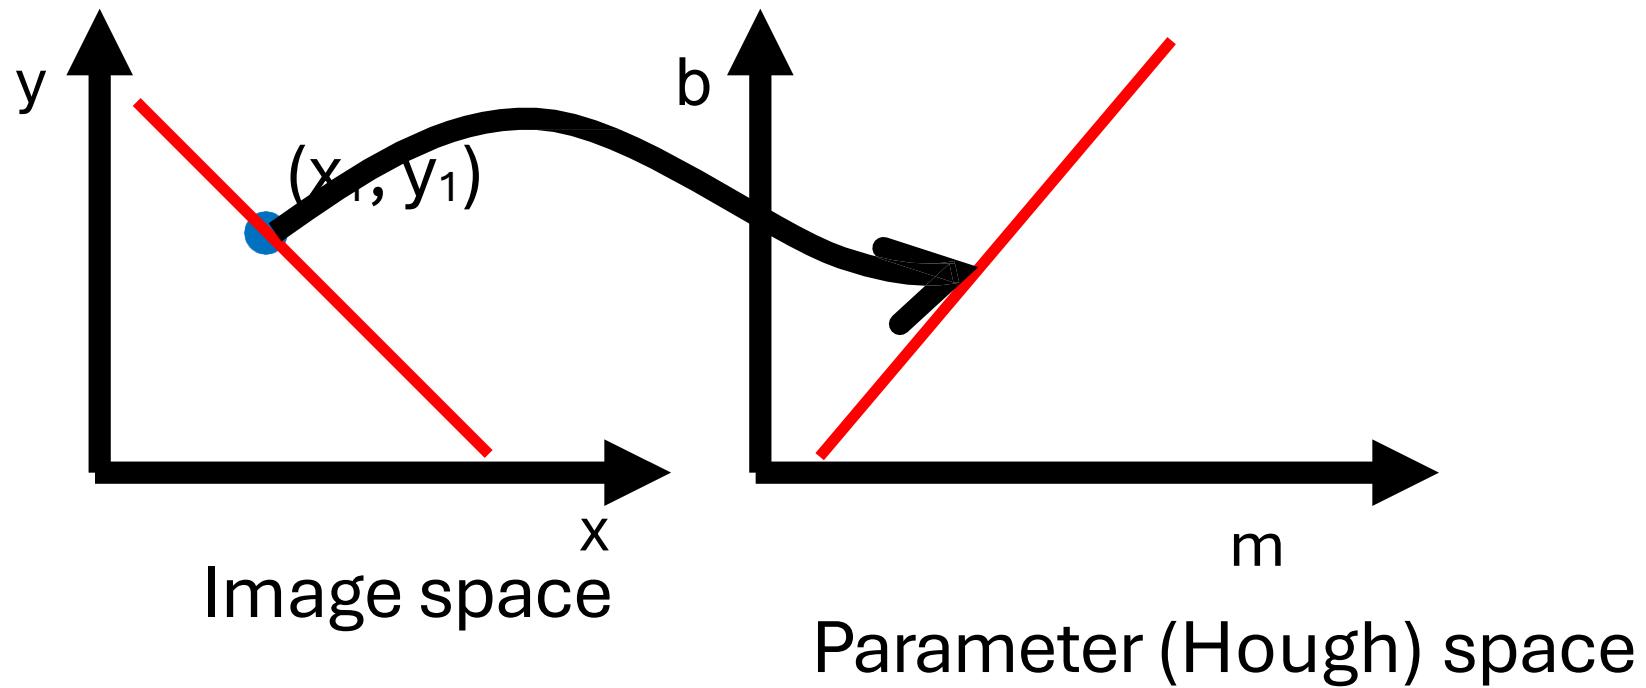
Image space

Parameter (Hough) space

A point  $(x_1, y_1)$  is mapped to a line in Hough space

$$y_1 = mx_1 + b$$

$$b = -x_1 m + y_1 \quad (\text{rewrite the equation in terms of } m \text{ and } b)$$

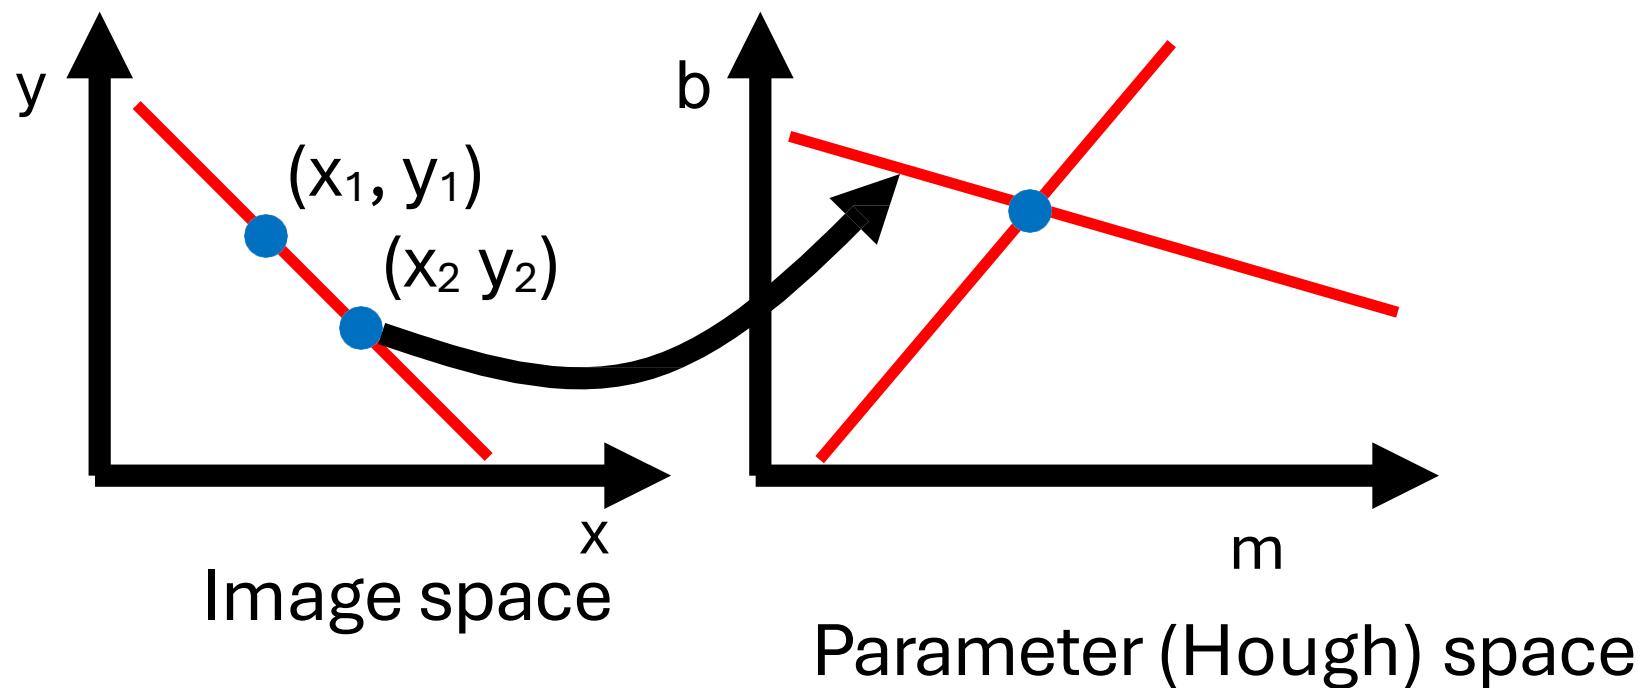


# A point $(x_2, y_2)$ is mapped to a line in Hough space

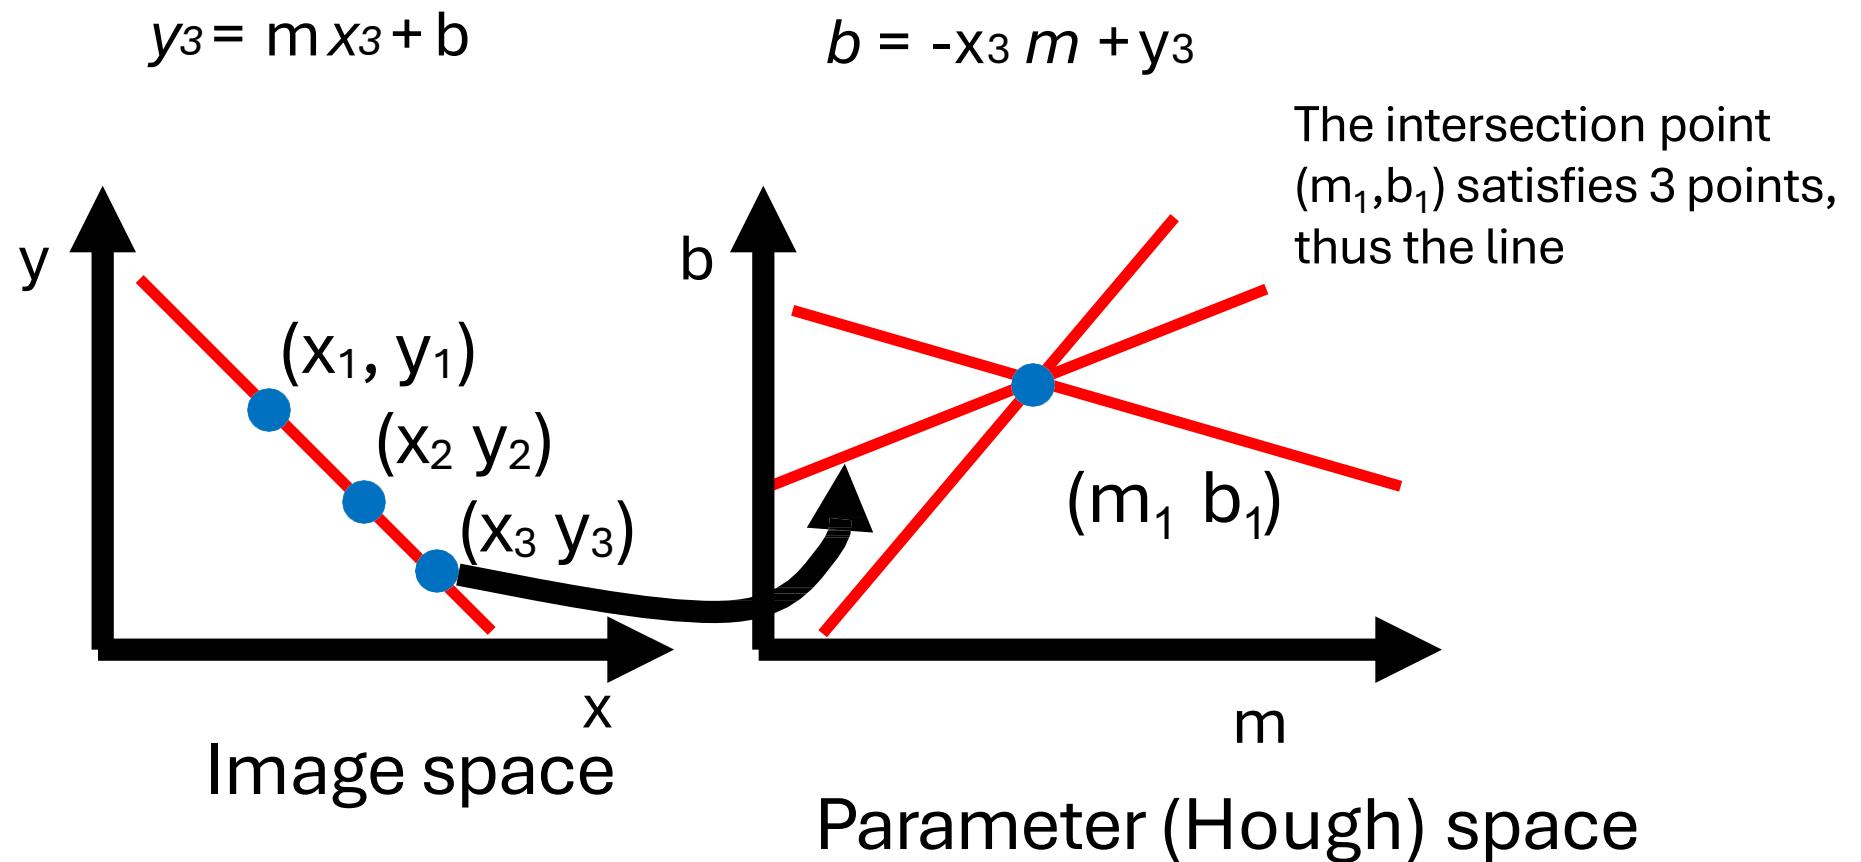
$$y_2 = mx_2 + b$$

$$b = -x_2 m + y_2$$

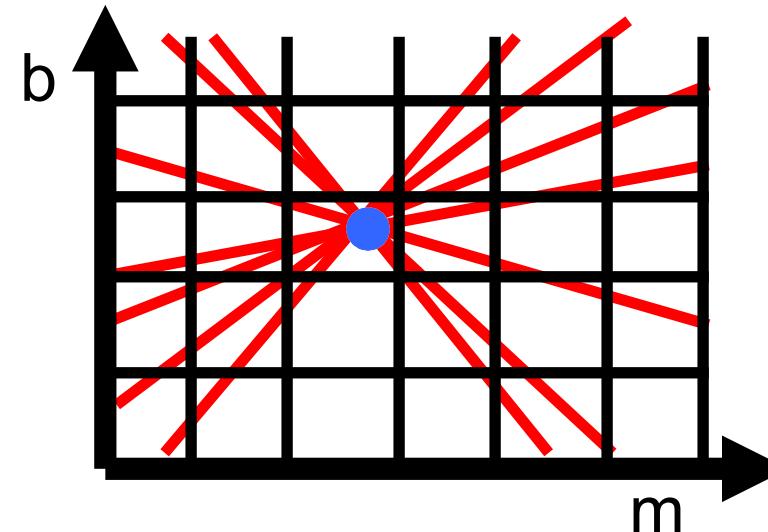
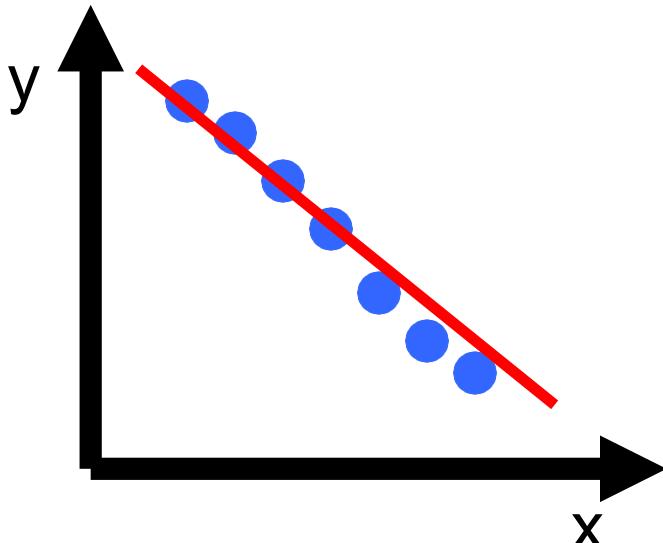
The intersection point satisfies both  $(x_1, y_1)$  and  $(x_2, y_2)$ , thus the line



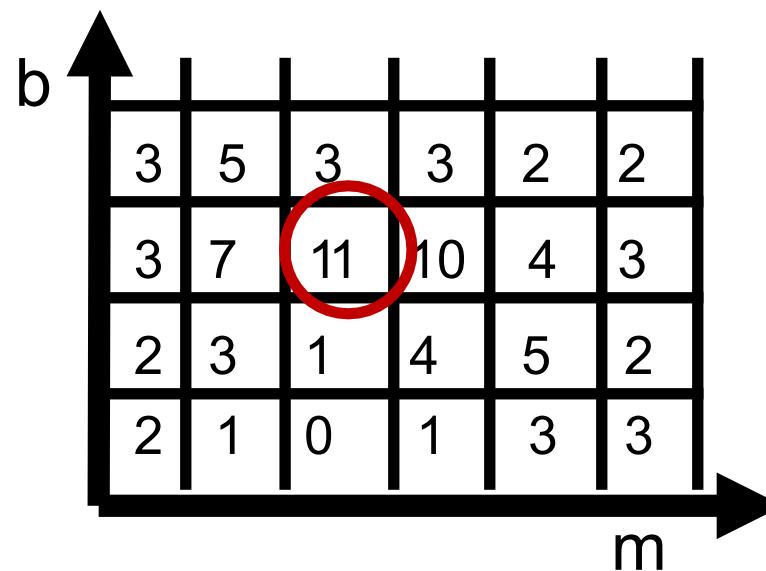
# A point $(x_3, y_3)$ is mapped to a line in Hough space



# Hough Transform: Accumulator



- For each pixel, draw a line in the discretised Hough space, assigning a value of one to all the discretised positions it passes through
- Process all image pixels

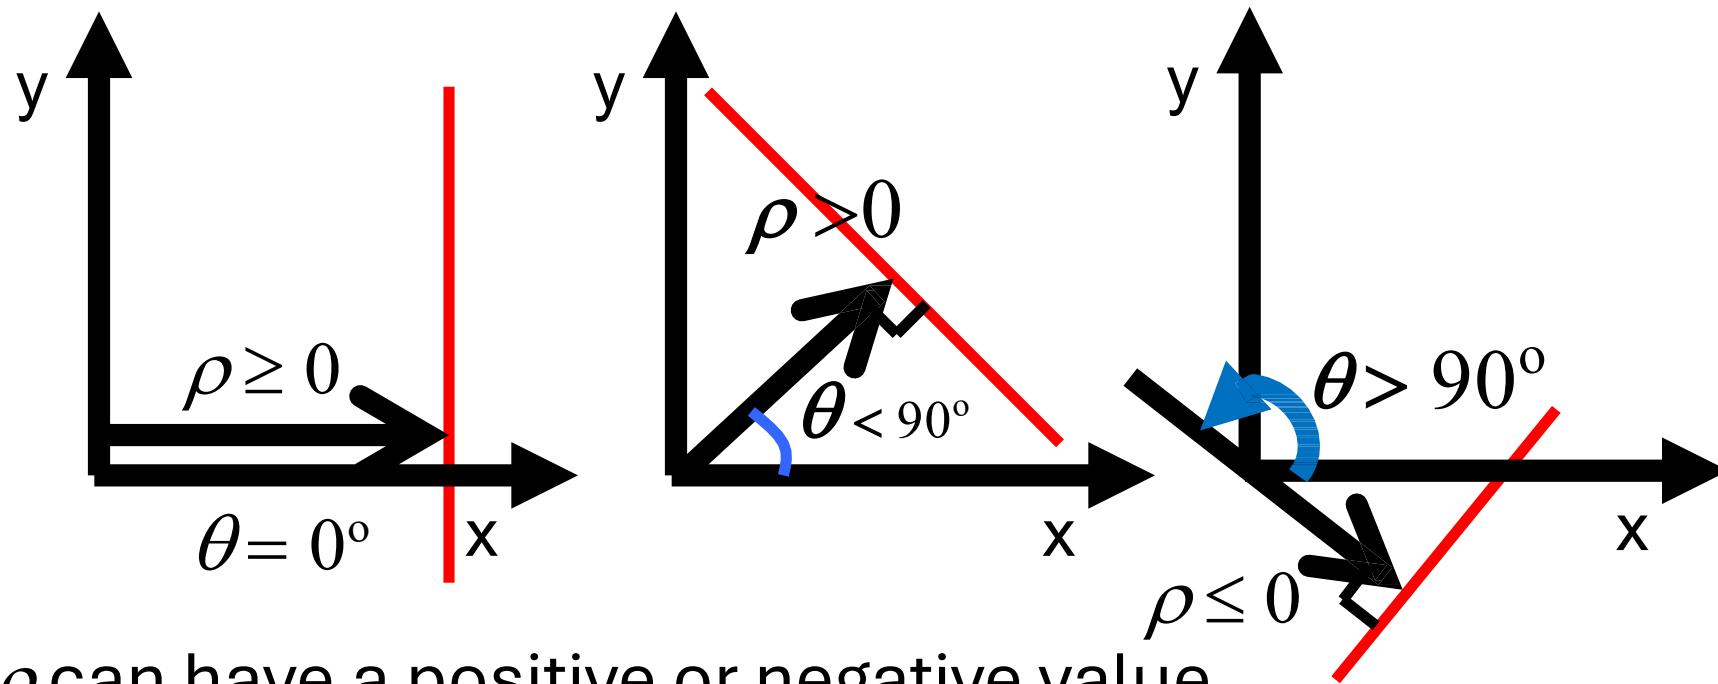


# Hough Transform

- Can detect multiple lines in an image
  - from multiple local maxima
- Can easily be extended for circles and ellipses
- Computationally efficient
- Problem:  $(m, b)$  are unbounded
  - E.g., the slope parameter  $m$  can have an infinite value

# Hough Transform: Polar Form

- Use a polar representation for the parameter space
- $x \cos \theta + y \sin \theta = \rho \rightarrow y = -\frac{\cos \theta}{\sin \theta} x + \frac{\rho}{\sin \theta}$   $(0 \leq \theta \leq 180^\circ)$

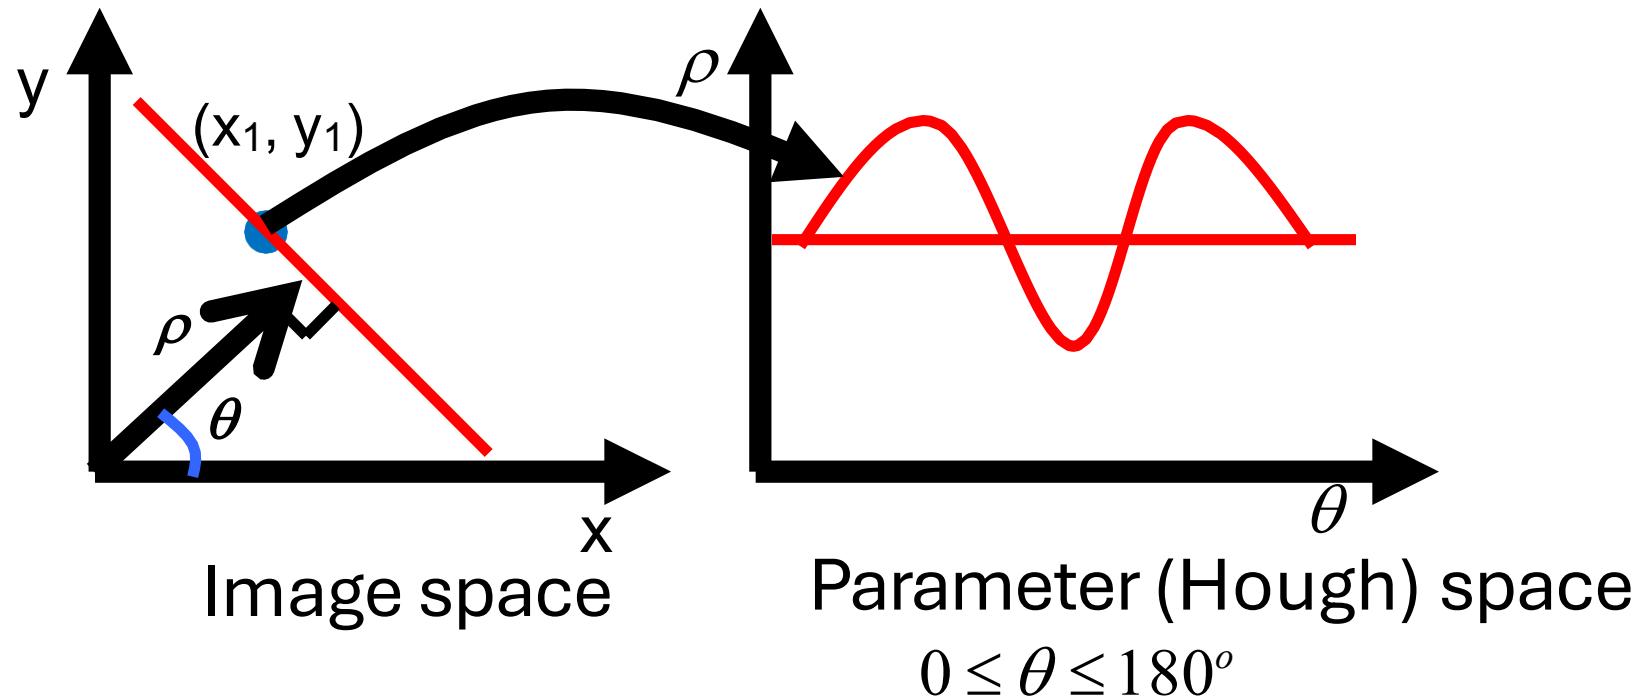


- Note:  $\rho$  can have a positive or negative value

A point  $(x_1, y_1)$  is mapped to a sinusoid in the polar parameter space

$$x_1 \cos \theta + y_1 \sin \theta = \rho$$

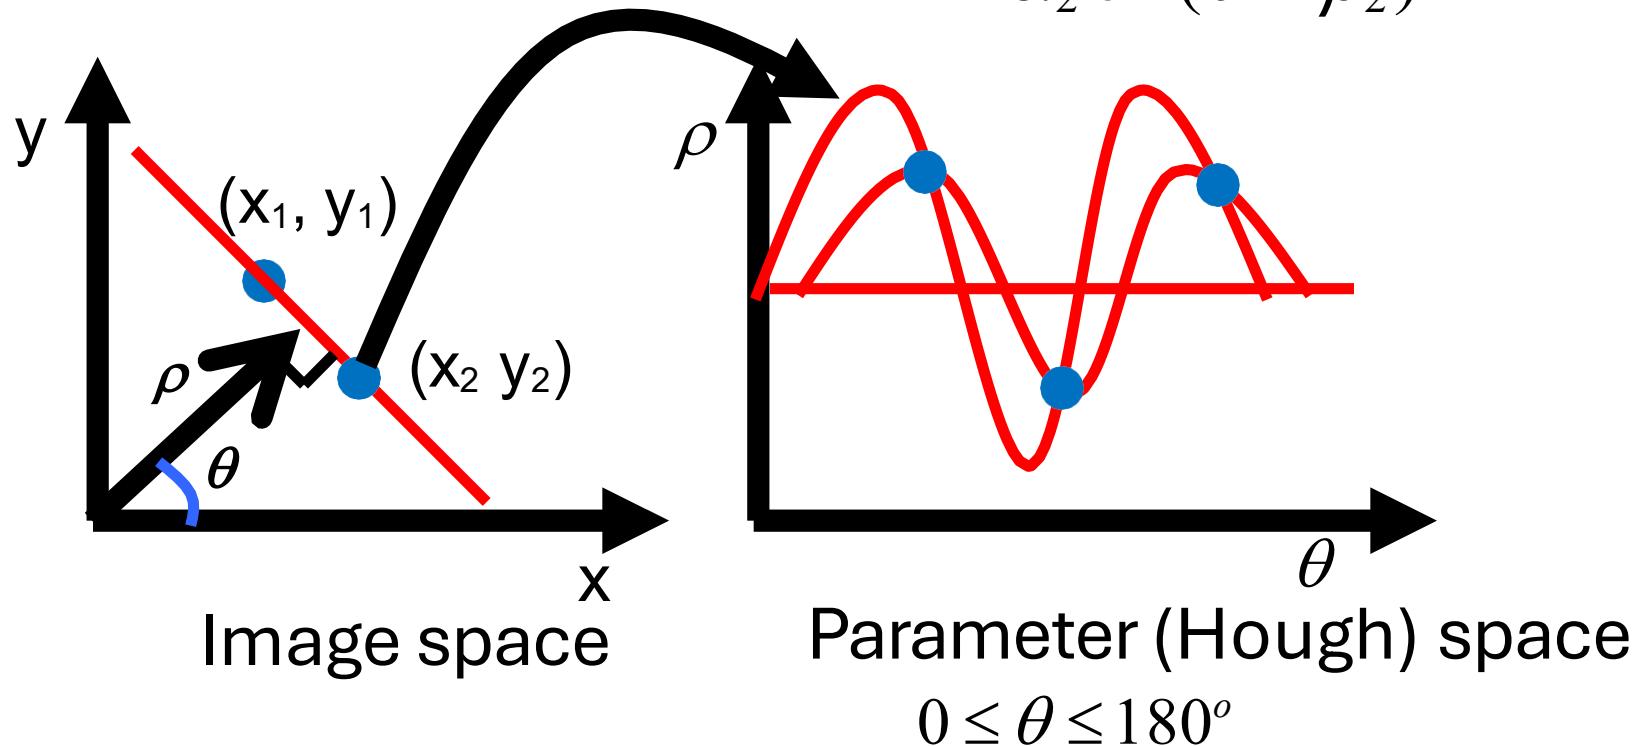
$$\begin{aligned}\rho &= x_1 \cos \theta + y_1 \sin \theta \\ &= \alpha_1 \sin(\theta + \beta_1)\end{aligned}$$



How to map: just divide the right-hand side by  $\sqrt{x_1^2 + y_1^2}$  and multiply by the same

A point  $(x_2, y_2)$  is mapped to a sinusoid in the polar parameter space

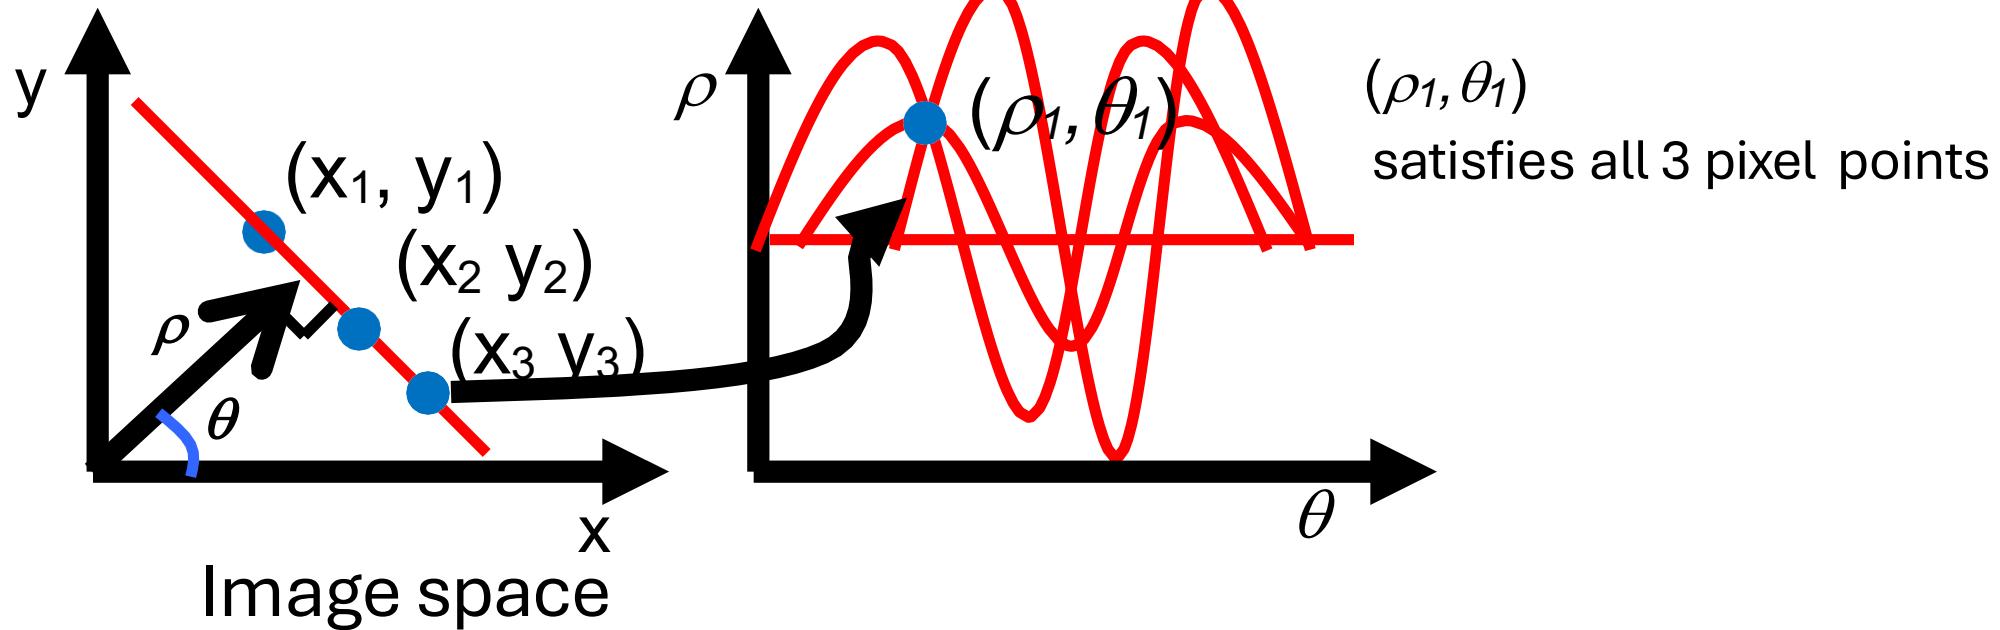
$$x_2 \cos \theta + y_2 \sin \theta = \rho$$
$$\rho = x_2 \cos \theta + y_2 \sin \theta$$
$$= \alpha_2 \sin(\theta + \beta_2)$$



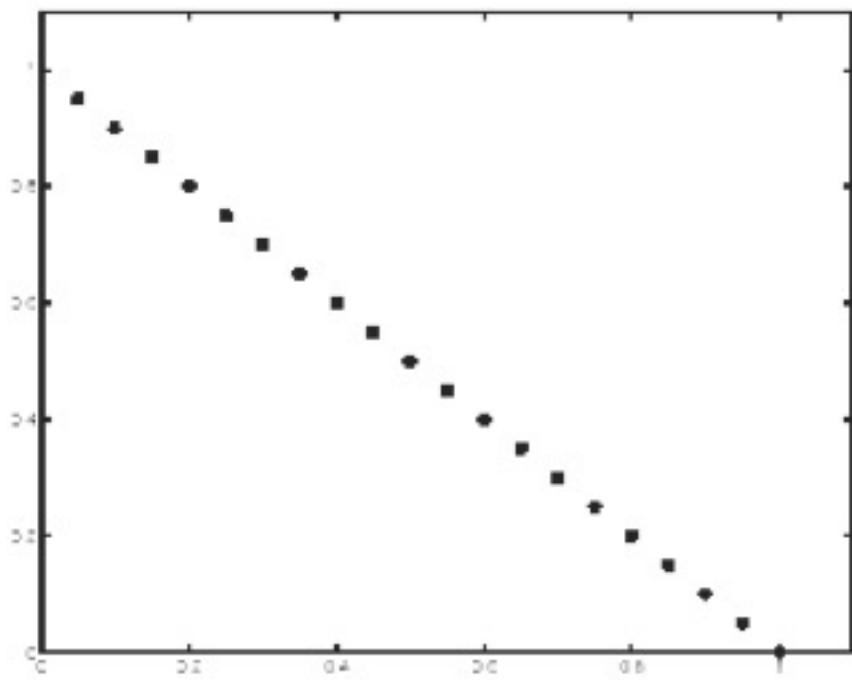
A point  $(x_3, y_3)$  is mapped to a sinusoid in the polar parameter space

$$x_3 \cos \theta + y_3 \sin \theta = \rho$$

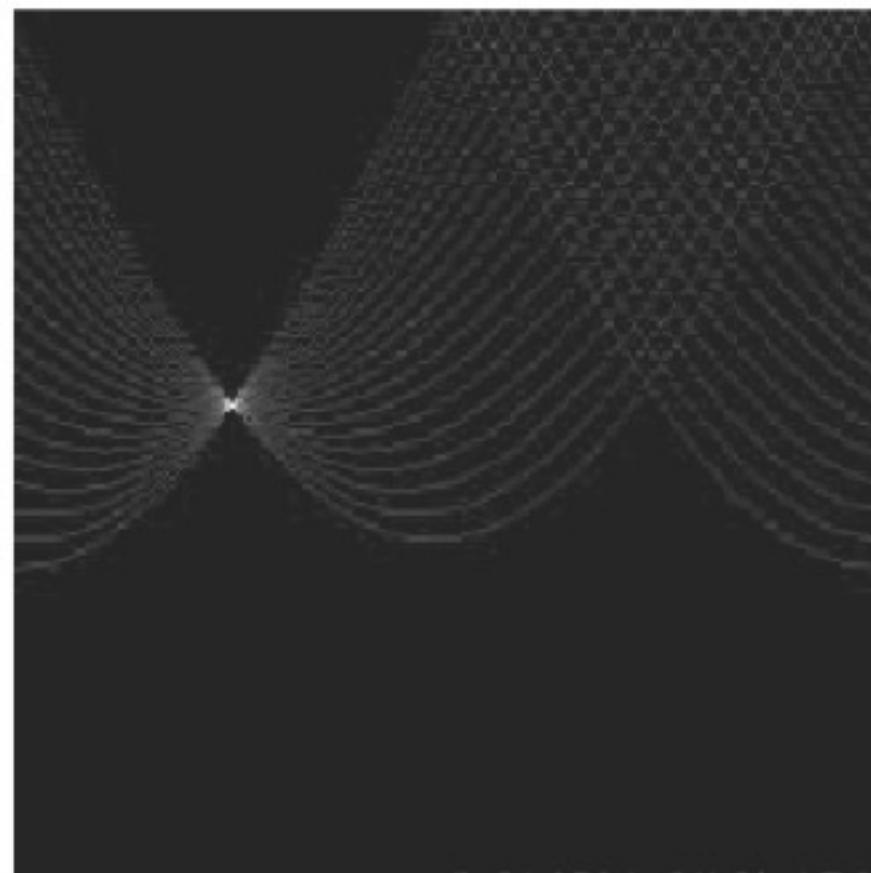
$$\begin{aligned}\rho &= x_3 \cos \theta + y_3 \sin \theta \\ &= \alpha_3 \sin(\theta + \beta_3)\end{aligned}$$



# Hough Transform: Example

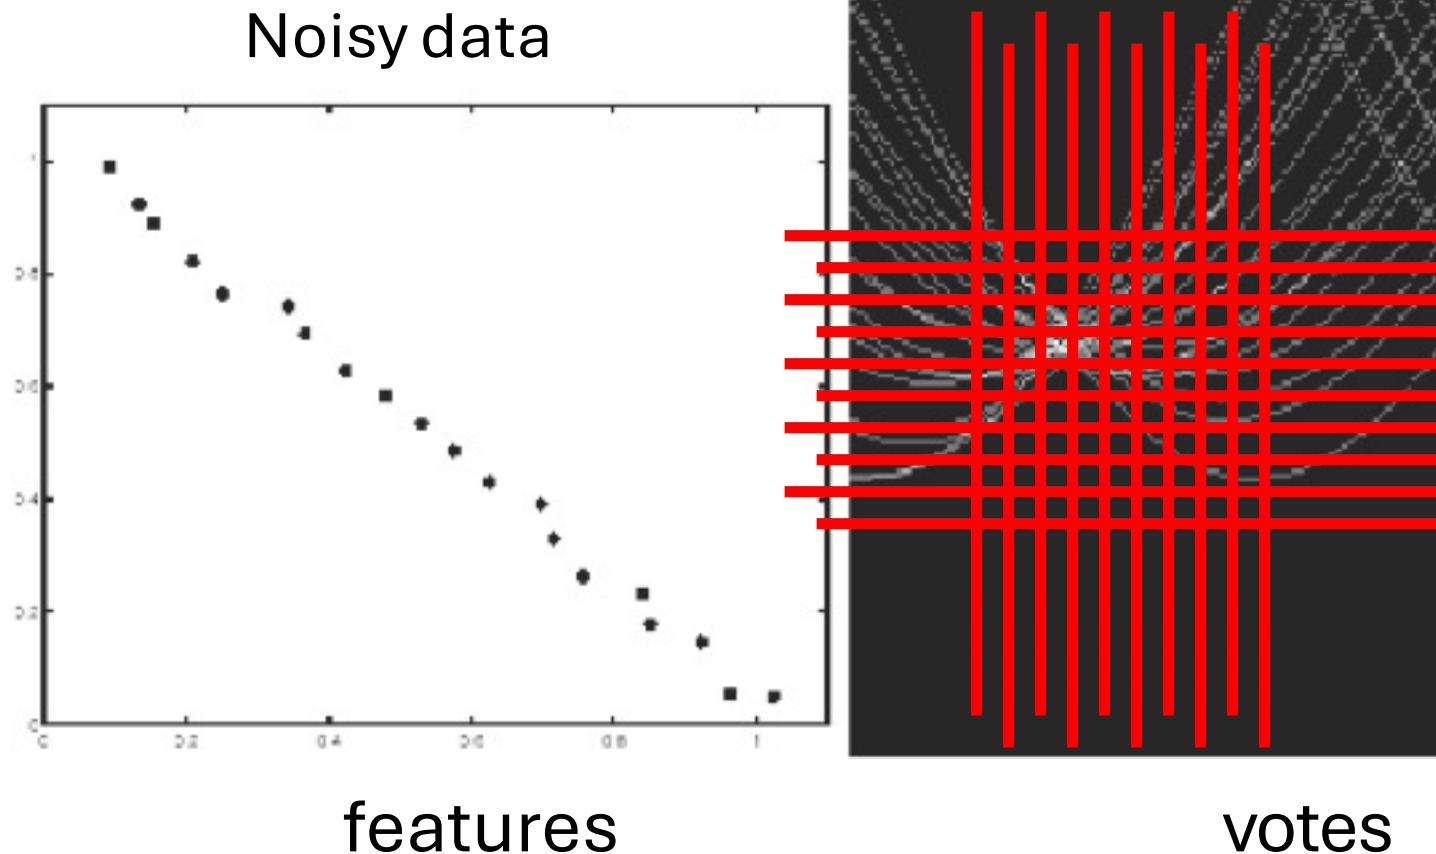


features



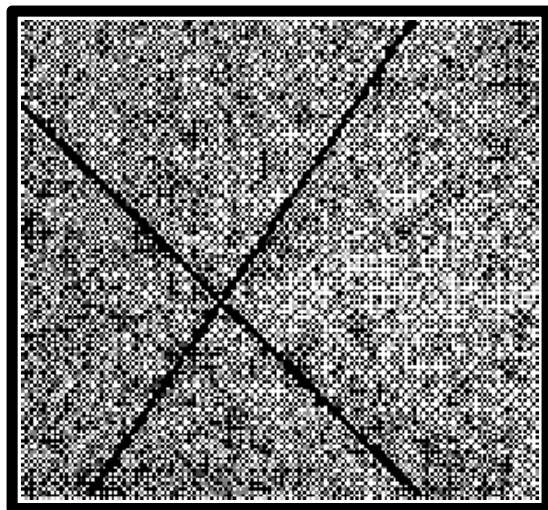
votes

# Hough Transform: Example



Issue: the grid size needs to be adjusted...

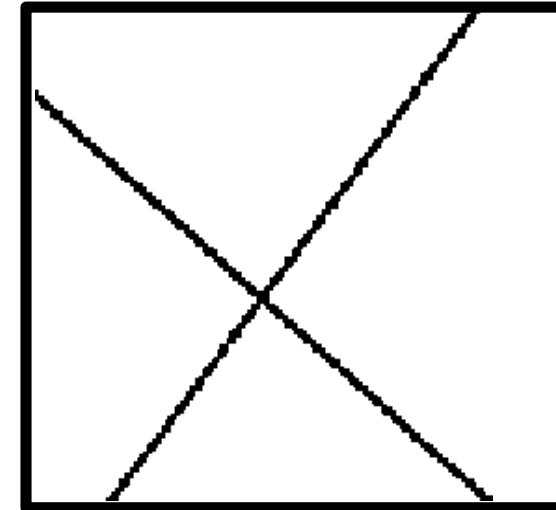
# Hough Transform: Example



**Image**



**Edge Detection**

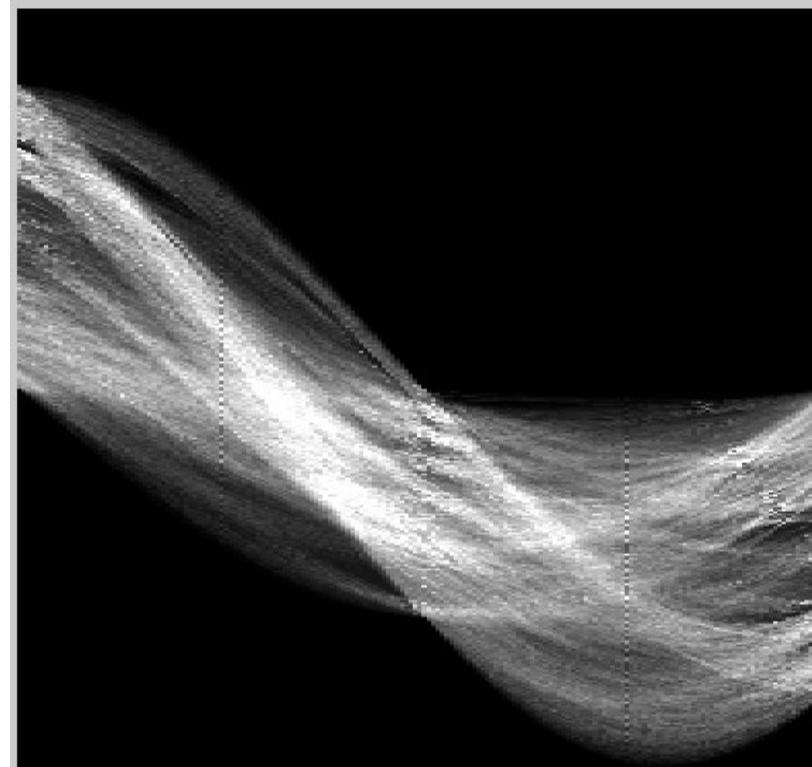


**Hough Transform**

# Hough Transform: 1) Image → Canny Edges

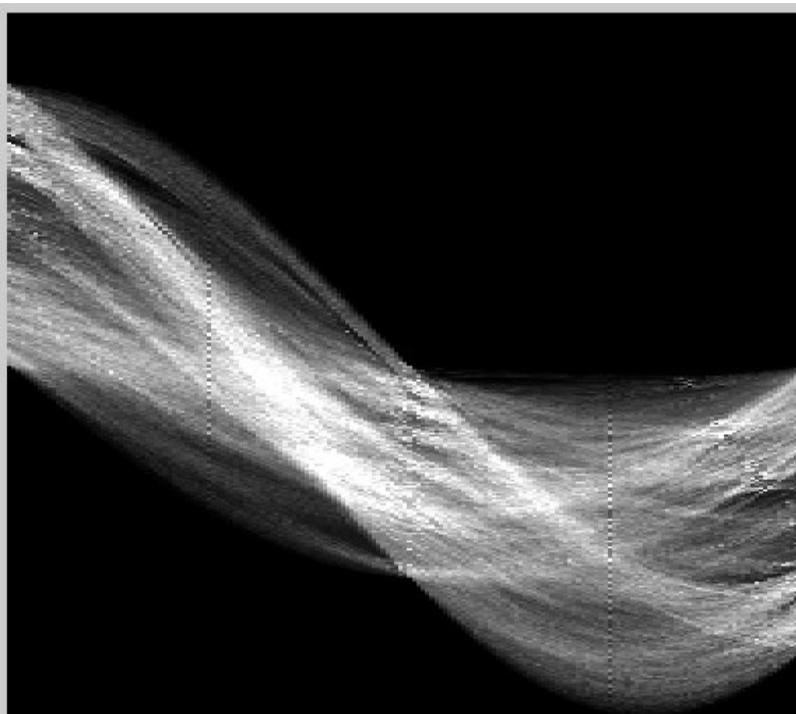


# Hough Transform: 2) Canny Edges → Votes

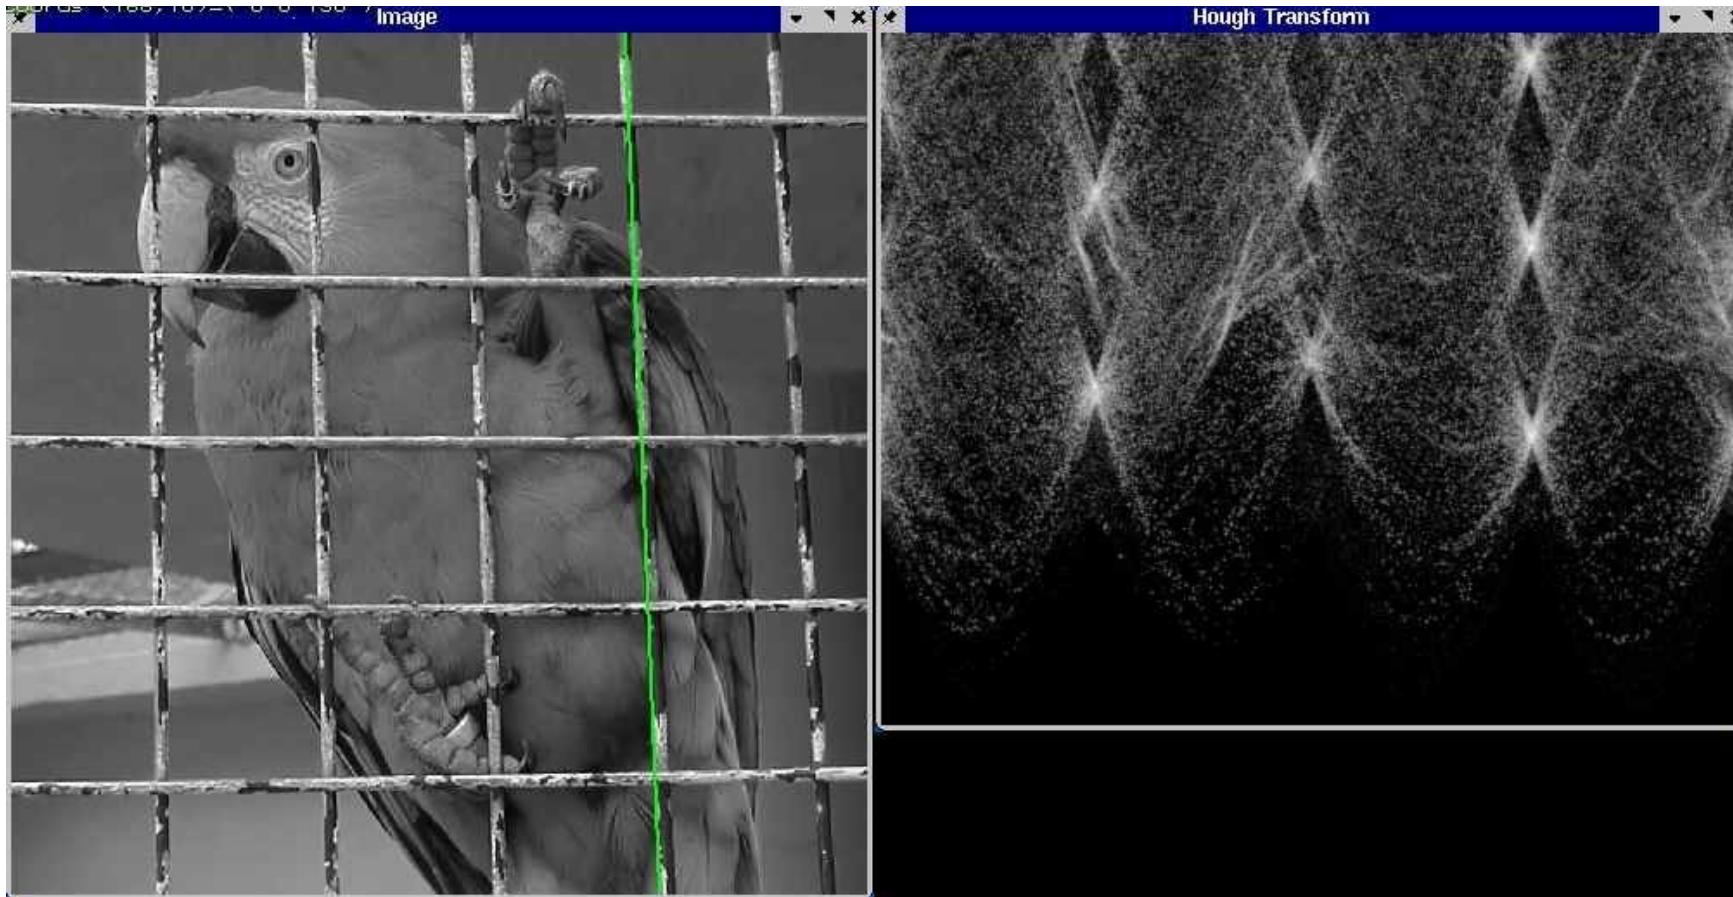


# Hough Transform: 3) Votes → Lines

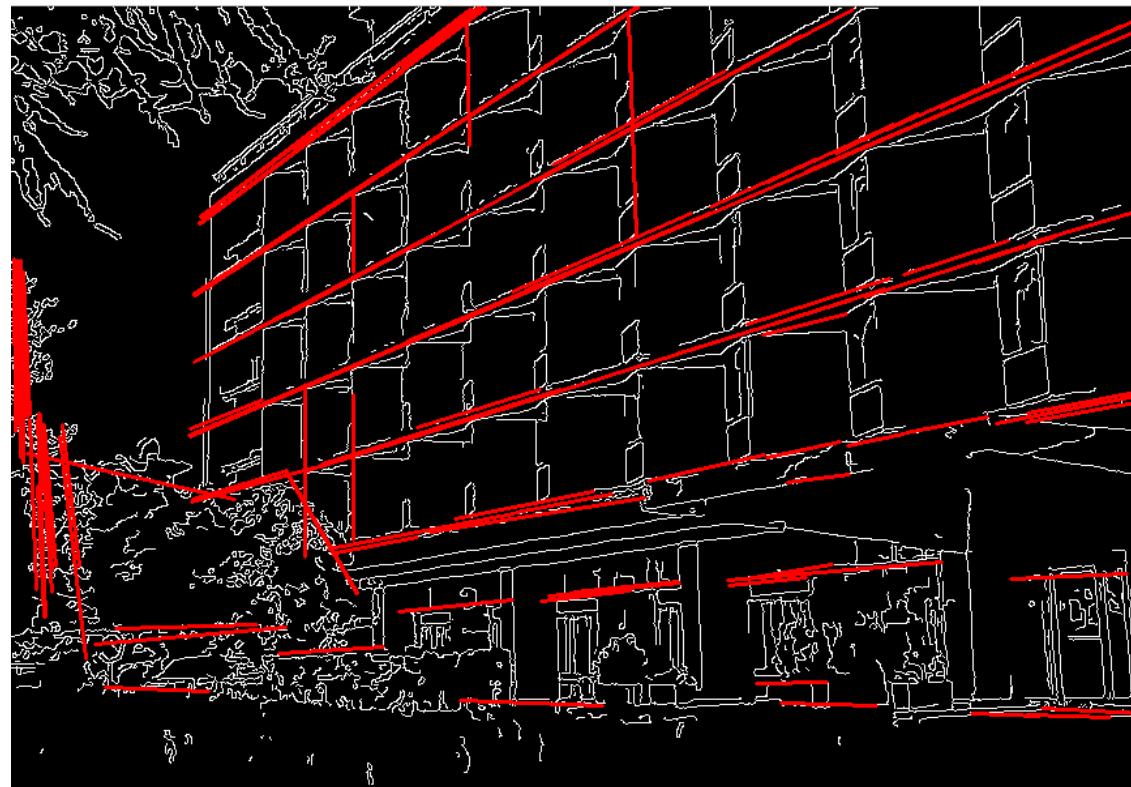
- Find peaks and post-process



# Hough Transform: Example



# Hough Transform: Example



# Hough Transform: Pros and Cons

## Pros:

- All points processed independently, so can cope with occlusion
- Some robustness to noise: noisy points are unlikely to contribute consistently to any single bin
- Can detect multiple instances of line/circle/ellipse in a single pass

## Cons:

- Complexity of search time increases exponentially with the number of model parameters
- Non-target shapes can produce spurious peaks in parameter space
- Quantization: hard to pick a grid size

# Summary

- Line features can be extracted using least squares or RANSAC fitting methods
- Both have difficulties in fitting multiple lines
- The Hough transform can handle multiple lines and the polar form is efficient due to the bounded search region
- Further reading: Chapter 4.3 in the book
- We will apply RANSAC to 3D Vision in future lectures