

3D Vision 2

Week 8

Single-view Geometry: Camera Calibration

Single-view Geometry: Resectioning and Camera Pose

Announcements

- Assignment 2 due Friday (11:59pm Friday 26 April)
 - This includes a one week extension that has already been applied
 - **Zero** marks if either report or code submitted late (unless extension)
 - Submit early; you can always resubmit an updated version later
 - Depending on your internet connection and load on the TurnItIn servers, uploading can sometimes be slow, so please factor this into your submission schedule
 - Submit your report (PDF) and code (ZIP file) **separately under the correct tab** in the submission box
 - Follow the instructions under Submission Requirements

Announcements

- Public Holiday on Thursday 25 April:
 - Thursday lab rescheduled to **13:00-15:00 Tuesday Rm 109 CSIT Building**

Weekly Study Plan: Overview

Wk	Starting	Lecture	Lab	Assessment
1	19 Feb	Introduction	X	
2	26 Feb	Low-level Vision 1	1	
3	4 Mar	Low-level Vision 2	1	
		Mid-level Vision 1		
4	11 Mar	Mid-level Vision 2	1	CLab1 report due Friday
		High-level Vision 1		
5	18 Mar	High-level Vision 2	2	
6	25 Mar	High-level Vision 3 ¹	2	
	1 Apr	Teaching break	X	
	8 Apr	Teaching break	X	
7	15 Apr	3D Vision 1	2	CLab2 report due Friday
8	22 Apr	3D Vision 2	3	
9	29 Apr	3D Vision 3	3	
10	6 May	3D Vision 4	3	
		Mid-level Vision 3		
11	13 May	High-level Vision 4	X	CLab3 report due Friday
12	20 May	Course Review	X	



Weekly Study Plan: Part B

Wk	Starting	Lecture	By
7	15 Apr	3D vision: introduction, camera model, single-view geometry	Dylan
8	22 Apr	3D vision: camera calibration, two-view geometry (homography)	Dylan
9	29 Apr	3D vision: two-view geometry (epipolar geometry, triangulation, stereo)	Dylan
10	6 May	3D vision: multiple-view geometry	Weijian
		Mid-level vision: optical flow, shape-from-X	Dylan
11	13 May	High-level vision: self-supervised learning, detection, segmentation	Dylan
12	20 May	Course review	Dylan


Outline

1. Single-view Geometry: Camera Calibration
2. Single-view Geometry: Resectioning and Absolute Camera Pose
3. Two-view Geometry: Homography Estimation

Vector Operations (Review)


Vectors

- $\mathbf{x} = (x_1, x_2, x_3, \dots, x_n)$
- Magnitude: $\|\mathbf{x}\| = (x_1^2 + x_2^2 + x_3^2 + \dots + x_n^2)^{\frac{1}{2}}$
- Unit vector (magnitude is one): $\hat{\mathbf{x}} = \frac{\mathbf{x}}{\|\mathbf{x}\|}$
- Orientation (for a 2D vector): $\theta = \tan^{-1} \frac{x_2}{x_1}$
- Homogeneous vectors (in P^n): $\tilde{\mathbf{x}}$
 - One fewer degrees-of-freedom than the number of dimensions
 - Known up to scale: only the ratios between coordinates are significant $x_1 : x_2 : x_3 : \dots$
 - $\tilde{\mathbf{x}} = k\tilde{\mathbf{x}}$




Inner (Dot) Product

- $\mathbf{x} = (x_1, x_2, x_3, \dots, x_n)$
- $\mathbf{y} = (y_1, y_2, y_3, \dots, y_n)$
- $\mathbf{x} \cdot \mathbf{y} = x_1y_1 + x_2y_2 + x_3y_3 + \dots + x_ny_n \rightarrow \text{scalar}$
- $\mathbf{x} \cdot \mathbf{y} = \|\mathbf{x}\| \|\mathbf{y}\| \cos \theta$
- If $\mathbf{x} \perp \mathbf{y}$, $\mathbf{x} \cdot \mathbf{y} = 0$




Vector (Cross) Product

- $\mathbf{x} = (x_1, x_2, x_3, \dots, x_n)$
- $\mathbf{y} = (y_1, y_2, y_3, \dots, y_n)$
- $\mathbf{z} = \mathbf{x} \times \mathbf{y} \rightarrow$ vector
- Magnitude: $\|\mathbf{z}\| = \|\mathbf{x}\| \|\mathbf{y}\| \sin \theta$
- Orientation:
 - $\mathbf{z} \perp \mathbf{x} \Rightarrow \mathbf{x} \cdot \mathbf{z} = \mathbf{x} \cdot (\mathbf{x} \times \mathbf{y}) = 0$
 - $\mathbf{z} \perp \mathbf{y} \Rightarrow \mathbf{y} \cdot \mathbf{z} = \mathbf{y} \cdot (\mathbf{x} \times \mathbf{y}) = 0$
- If $\mathbf{x} \parallel \mathbf{y}$, $\mathbf{z} = \mathbf{0}$




Vector (Cross) Product: Computation (3D)

- $\mathbf{x} = (x_1, x_2, x_3)$
- $\mathbf{y} = (y_1, y_2, y_3)$
- Unit basis vectors: $\hat{\mathbf{i}} = (1, 0, 0), \hat{\mathbf{j}} = (0, 1, 0), \hat{\mathbf{k}} = (0, 0, 1)$
 - $\hat{\mathbf{i}} = \hat{\mathbf{j}} \times \hat{\mathbf{k}}, \hat{\mathbf{j}} = \hat{\mathbf{k}} \times \hat{\mathbf{i}}, \hat{\mathbf{k}} = \hat{\mathbf{i}} \times \hat{\mathbf{j}}$
- $\mathbf{z} = (x_2 y_3 - x_3 y_2) \hat{\mathbf{i}} + (x_3 y_1 - x_1 y_3) \hat{\mathbf{j}} + (x_1 y_2 - x_2 y_1) \hat{\mathbf{k}}$



Vector (Cross) Product: Alternative Notation

- $\mathbf{x} = (x_1, x_2, x_3)$
- $\mathbf{y} = (y_1, y_2, y_3)$
- $\mathbf{x} \times \mathbf{y} = [\mathbf{x}]_{\times} \mathbf{y} = -[\mathbf{y}]_{\times} \mathbf{x}$
- Cross product matrix:
 - $[\mathbf{x}]_{\times} = \begin{pmatrix} 0 & -x_3 & x_2 \\ x_3 & 0 & -x_1 \\ -x_2 & x_1 & 0 \end{pmatrix}$
 - Anti-symmetric: $A = -A^T$




Camera Calibration

Recovering the Projection Matrix

Objectives

- To **calibrate** a perspective camera:
 - To estimate the camera matrix $P = K[R|t]$
 - To estimate the camera calibration (intrinsics) matrix K
 - To estimate the camera extrinsic parameters $R, t/C$
- To understand the **Direct Linear Transformation** (DLT) algorithm



$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \propto \underbrace{\begin{bmatrix} p_{11} & p_{12} & p_{13} & p_{14} \\ p_{21} & p_{22} & p_{23} & p_{24} \\ p_{31} & p_{32} & p_{33} & p_{34} \end{bmatrix}}_{P \{3 \times 4\}} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Image Projection

- Project world point $\mathbf{X} = (x, y, z)$ to image point $\mathbf{x} = (u, v)$
- Assuming a pinhole camera model


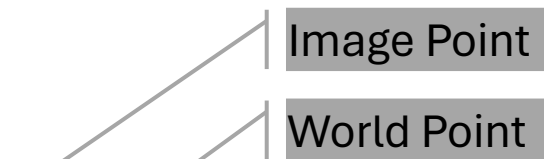



Image Projection


$$\begin{aligned}x &= PX \\&= KR[I \mid -C]X = K[R \mid t]X \\&= \begin{bmatrix} \alpha & \gamma & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{bmatrix} X\end{aligned}$$



- How many parameters?
 - 11 total:
 - 5 intrinsic ($\alpha, \beta, \gamma, u_0, v_0$)
 - 6 extrinsic (R, t)
- How to compute?
 - Camera calibration!
- But what about scale?
 - Inherent ambiguity between the focal length and the magnitude of the translation vector

Two Approaches to Calibration

1. Calibrate to a meaningful fixed world coordinate system
 - Solve for P directly
 - Good for a fixed camera, specific applications
 - E.g., tracking vehicles on a highway, a mobile robot on the ground plane
 - But gives no insight into internal calibration parameters
 - If camera–world relationship changes, calibration must start from scratch
2. Compute internal and external parameters separately:
 - $P = K[R|t]$
 - Internal parameters turn camera into a metric device
 - Can now be used for computing 3D rays in Euclidean space
 - Necessary for SFM

Camera Calibration

- Determine the camera parameters from known 3D points or a calibration object(s)
 1. Internal or intrinsic parameters (i.e., focal length, principal point, aspect ratio)
 2. External or extrinsic (pose) parameters (i.e., position and orientation of the camera)
- *3D points cannot all be on the same plane*



Basic Procedure

1. Prepare a calibration target/object
 - E.g., 2 orthogonal planes with a checkerboard pattern
 - Important: **3D point coordinates are known**
2. Position camera in front of target
 - Capture image of the calibration target
3. Find corners of the target in the image
 - Obtain **2D–3D correspondences**
4. Derive constraints on camera matrix
 - Estimate the intrinsic and extrinsic parameters


Basic Procedure: (3) Find Corners in Image

1. Option 1:

1. Detect edges with Canny detector
2. Fit straight lines to detected linked edges (Hough)
3. Intersect lines to obtain 2D image coordinates
4. Match image corners & 3D target checkerboard corners
 - E.g., count corners along the line and find corresponding edge in target
5. Result: 2D–3D correspondences

2. Option 2:

- Apply a corner detector directly (Harris, Susan, FAST)



Camera Projection Matrix

- Fold intrinsic calibration matrix **K** and extrinsic pose parameters (**R**,**t**) together into a camera matrix **P**
- **P = K [R | t]**

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \sim \begin{bmatrix} m_{00} & m_{01} & m_{02} & m_{03} \\ m_{10} & m_{11} & m_{12} & m_{13} \\ m_{20} & m_{21} & m_{22} & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

- Put 1 in the lower right corner to remove a degree of freedom (DoF)
 - 11 DoFs

Inhomogeneous Projection Equation

- Directly estimate 11 unknowns in the M matrix using known 3D points (X_i, Y_i, Z_i) and measured feature positions (u_i, v_i)
 - **Nonlinear** equations

$$u_i = \frac{m_{00}X_i + m_{01}Y_i + m_{02}Z_i + m_{03}}{m_{20}X_i + m_{21}Y_i + m_{22}Z_i + 1}$$
$$v_i = \frac{m_{10}X_i + m_{11}Y_i + m_{12}Z_i + m_{13}}{m_{20}X_i + m_{21}Y_i + m_{22}Z_i + 1}$$

Linear Regression

1. Bring denominator over to the LHS
2. Solve set of (over-determined) linear equations for m_{ij}
 - How? Least squares (pseudo-inverse)

$$\begin{aligned}u_i(m_{20}X_i + m_{21}Y_i + m_{22}Z_i + 1) &= m_{00}X_i + m_{01}Y_i + m_{02}Z_i + m_{03} \\v_i(m_{20}X_i + m_{21}Y_i + m_{22}Z_i + 1) &= m_{10}X_i + m_{11}Y_i + m_{12}Z_i + m_{13}\end{aligned}$$

Direct Linear Transformation (DLT) Algorithm for Camera Calibration

- Cross-product trick:
 - When an equation is only **known up to scale**, take the cross product of the LHS with both sides of the equation – no loss of information
 - $\mathbf{x}_i = k\mathbf{P}\mathbf{X}_i \Rightarrow \mathbf{x}_i \times \mathbf{P}\mathbf{X}_i = \mathbf{0}$ [Why?] (\mathbf{x}_i : image coords; \mathbf{X}_i : world coords)

$$[\mathbf{x}_i]_{\times} \mathbf{P} \mathbf{X}_i = \begin{pmatrix} y_i \mathbf{p}_3^{\top} \mathbf{X}_i - w_i \mathbf{p}_2^{\top} \mathbf{X}_i \\ w_i \mathbf{p}_1^{\top} \mathbf{X}_i - x_i \mathbf{p}_3^{\top} \mathbf{X}_i \\ x_i \mathbf{p}_2^{\top} \mathbf{X}_i - y_i \mathbf{p}_1^{\top} \mathbf{X}_i \end{pmatrix} = \begin{bmatrix} \mathbf{0}^{\top} & -w_i \mathbf{X}_i^{\top} & y_i \mathbf{X}_i^{\top} \\ w_i \mathbf{X}_i^{\top} & \mathbf{0}^{\top} & -x_i \mathbf{X}_i^{\top} \\ -y_i \mathbf{X}_i^{\top} & x_i \mathbf{X}_i^{\top} & \mathbf{0}^{\top} \end{bmatrix} \begin{pmatrix} \mathbf{p}_1 \\ \mathbf{p}_2 \\ \mathbf{p}_3 \end{pmatrix} = \mathbf{A}'_i \mathbf{p} = \mathbf{0}$$

$$\text{where } \mathbf{x}_i = (x_i, y_i, w_i)^{\top} \text{ and } \mathbf{P} = \begin{bmatrix} \mathbf{p}_1^{\top} \\ \mathbf{p}_2^{\top} \\ \mathbf{p}_3^{\top} \end{bmatrix}; \mathbf{p}_i \in \mathbb{R}^{4 \times 1}$$

Direct Linear Transformation (DLT) Algorithm for Camera Calibration

- Only 2 out of 3 equations are linearly independent, so pick two

$$\begin{bmatrix} \mathbf{0}^\top & -w_i \mathbf{X}_i^\top & y_i \mathbf{X}_i^\top \\ w_i \mathbf{X}_i^\top & \mathbf{0}^\top & -x_i \mathbf{X}_i^\top \end{bmatrix} \begin{pmatrix} \mathbf{p}_1 \\ \mathbf{p}_2 \\ \mathbf{p}_3 \end{pmatrix} = \mathbf{A}_i \mathbf{p} = \mathbf{0}$$

- Camera matrix has 11 DoF:
 - 12 parameters defined up to scale
 - Linear solution requires at least 6 points (in fact, 5½)

Direct Linear Transformation (DLT) Algorithm for Camera Calibration

- Using 6 points to solve for \mathbf{P} :

$$\mathbf{A}\mathbf{p} = \begin{bmatrix} \mathbf{A}_1 \\ \mathbf{A}_2 \\ \vdots \\ \mathbf{A}_6 \end{bmatrix} \mathbf{p} = \mathbf{0}$$

- $\mathbf{A} \in \mathbb{R}^{12 \times 12}$ but $\text{rank}(\mathbf{A}) = 11$
- So use the first 11 rows: $\mathbb{R}^{11 \times 12}$
- How to solve?
 - The trivial solution $\mathbf{p} = \mathbf{0}$ is not interesting
 - Compute the 1D null-space (e.g., via SVD)
 - Fix norm of \mathbf{p} afterwards (e.g., set $\|\mathbf{p}\| = 1$)

Direct Linear Transformation (DLT) Algorithm for Camera Calibration

- Using n points to solve for \mathbf{P} : the over-determined case

$$\mathbf{A}\mathbf{p} = \begin{bmatrix} \mathbf{A}_1 \\ \mathbf{A}_2 \\ \vdots \\ \mathbf{A}_n \end{bmatrix} \mathbf{p} = \mathbf{0}$$

- How to solve?
 - No exact non-trivial solution due to inexact measurements (e.g., noise)
 - $\mathbf{A}\mathbf{p} = \mathbf{0}$ is not possible, so minimise $\|\mathbf{A}\mathbf{p}\|$ subject to $\|\mathbf{p}\| = 1$
 - 1. Take the singular value decomposition (SVD) of \mathbf{A}
 - $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^\top$
 - 2. Take the rightmost column of \mathbf{V} [why?]
 - The right-singular vector of \mathbf{A} , corresponding to the smallest singular value (arranged in decreasing singular value order)

Direct Linear Transformation (DLT) Algorithm for Camera Calibration

- Objective:
 - Given $n \geq 6$ 2D–3D point correspondences $\{\mathbf{x}_i \leftrightarrow \mathbf{X}_i\}$, determine the 3×4 projection matrix \mathbf{P} such that $\mathbf{x}_i \approx \mathbf{P}\mathbf{X}_i$
- Algorithm:
 1. For each correspondence $\{\mathbf{x}_i \leftrightarrow \mathbf{X}_i\}$ compute \mathbf{A}_i , taking only the first two rows
 2. Assemble the n 2×12 \mathbf{A}_i matrices into a single $2n \times 12$ matrix \mathbf{A}
 3. Compute the SVD of \mathbf{A} : $\mathbf{U}\Sigma\mathbf{V}^\top$
 4. Take the last column of \mathbf{V} as the solution for \mathbf{p}
 5. Rearrange \mathbf{p} to obtain \mathbf{P}

Importance of Normalisation

$$\begin{bmatrix} 0 & 0 & 0 & 0 & -w_i X_1 & -w_i X_2 & -w_i X_3 & -w_i & y_i X_1 & y_i X_2 & y_i X_3 & y_i \\ w_i X_1 & w_i X_2 & w_i X_3 & w_i & 0 & 0 & 0 & 0 & -x_i X_1 & -x_i X_2 & -x_i X_3 & -x_i \end{bmatrix} \begin{pmatrix} \mathbf{p}_1 \\ \mathbf{p}_2 \\ \mathbf{p}_3 \end{pmatrix}$$

$\sim 10^2 \quad \sim 10^2 \quad \sim 10^2 \quad 1 \quad \sim 10^2 \quad \sim 10^2 \quad \sim 10^2 \quad 1 \quad \sim 10^4 \quad \sim 10^4 \quad \sim 10^4 \quad \sim 10^2$

- Orders of magnitude difference!

Normalised Direct Linear Transformation (DLT) Algorithm for Camera Calibration

- Objective:
 - Given $n \geq 6$ 2D–3D point correspondences $\{\mathbf{x}_i \leftrightarrow \mathbf{X}_i\}$, determine the 3×4 projection matrix \mathbf{P} such that $\mathbf{x}_i \approx \mathbf{P}\mathbf{X}_i$
- Algorithm:
 1. Normalise 2D and 3D points: $\tilde{\mathbf{x}}_i = \mathbf{T}\mathbf{x}_i$, $\tilde{\mathbf{X}}_i = \mathbf{S}\mathbf{X}_i$
 2. Apply the DLT algorithm to $\{\tilde{\mathbf{x}}_i \leftrightarrow \tilde{\mathbf{X}}_i\}$
 3. Denormalise the recovered solution $\tilde{\mathbf{P}}$ using $\mathbf{P} = \mathbf{T}^{-1}\tilde{\mathbf{P}}\mathbf{S}$
- Example normalisation matrices:

$$\mathbf{T} = \begin{bmatrix} w+h & 0 & w/2 \\ 0 & w+h & h/2 \\ 0 & 0 & 1 \end{bmatrix}^{-1}; \quad \mathbf{S} = \begin{bmatrix} \mathbf{V}\text{diag}(\lambda_1^{-1}, \lambda_2^{-1}, \lambda_3^{-1})\mathbf{V}^{-1} & -\mathbf{V}\text{diag}(\lambda_1^{-1}, \lambda_2^{-1}, \lambda_3^{-1})\mathbf{V}^{-1}\boldsymbol{\mu}_{\mathbf{x}_i} \\ 0 & 1 \end{bmatrix}$$

$$\mathbf{V}\text{diag}(\lambda_1, \lambda_2, \lambda_3)\mathbf{V}^{-1} = \text{eig}\left(\sum_i (\mathbf{X}_{i,\text{inhom}} - \boldsymbol{\mu}_{\mathbf{x}_i})(\mathbf{X}_{i,\text{inhom}} - \boldsymbol{\mu}_{\mathbf{x}_i})^T\right)$$

Camera Calibration

Recovering the Camera Intrinsics

Camera Matrix Decomposition: Computing the Camera Centre **C**

- $\mathbf{P} = \mathbf{K}[\mathbf{R} \mid -\mathbf{RC}] = [\mathbf{p}_1 \quad \mathbf{p}_2 \quad \mathbf{p}_3 \quad \mathbf{p}_4]$

- Careful! These \mathbf{p}_i are now column vectors of \mathbf{P}

1. It is the right null-space vector of \mathbf{P} (Hartley & Zisserman p. 163):

- Take last column of \mathbf{V} where $\mathbf{P} = \mathbf{U}\Sigma\mathbf{V}^\top$ is the SVD of \mathbf{P}

- Why? Observe that $\mathbf{PC} = \mathbf{KR} \begin{bmatrix} 1 & & & -X_C \\ & 1 & & -Y_C \\ & & 1 & -Z_C \\ & & & 1 \end{bmatrix} \begin{bmatrix} X_C \\ Y_C \\ Z_C \\ 1 \end{bmatrix} = \mathbf{0}$

2. Or, algebraic derivation (Hartley & Zisserman p. 163):

- $$\mathbf{C} = \begin{bmatrix} X \\ Y \\ Z \\ T \end{bmatrix} = \begin{bmatrix} \det([\mathbf{p}_2, \mathbf{p}_3, \mathbf{p}_4]) \\ -\det([\mathbf{p}_1, \mathbf{p}_3, \mathbf{p}_4]) \\ \det([\mathbf{p}_1, \mathbf{p}_2, \mathbf{p}_4]) \\ -\det([\mathbf{p}_1, \mathbf{p}_2, \mathbf{p}_3]) \end{bmatrix}$$

Camera Matrix Decomposition: Computing the Intrinsics **K** and Rotation **R**

- $\mathbf{P} = \mathbf{K}[\mathbf{R} \mid -\mathbf{RC}] = [\mathbf{M} \mid -\mathbf{MC}]$

1. RQ decomposition of **M**:

- $(\mathbf{R}_\Delta, \mathbf{Q}) = \text{RQ}(\mathbf{M})$

- $\mathbf{K} = \mathbf{R}_\Delta$: upper triangular matrix (Δ is just to distinguish it from rotation)

- $\mathbf{R} = \mathbf{Q}$: orthonormal matrix

2. Algebraic derivation:

- See next slide

Camera Matrix Decomposition: Computing the Intrinsics **K** and Rotation **R**

- $\mathbf{P} = \mathbf{K}[\mathbf{R} \mid -\mathbf{RC}] = [\mathbf{M} \mid -\mathbf{MC}]$

2. Algebraic derivation:

- Given rotations: $\mathbf{R}_x = \begin{bmatrix} 1 & 0 & 0 \\ 0 & c & -s \\ 0 & s & c \end{bmatrix}$ $\mathbf{R}_y = \begin{bmatrix} c' & 0 & s' \\ 0 & 1 & 0 \\ -s' & 0 & c' \end{bmatrix}$ $\mathbf{R}_z = \begin{bmatrix} c'' & -s'' & 0 \\ s'' & c'' & 0 \\ 0 & 0 & 1 \end{bmatrix}$

- $c = -\frac{m_{33}}{\sqrt{m_{32}^2 + m_{33}^2}} \quad s = \frac{m_{32}}{\sqrt{m_{32}^2 + m_{33}^2}}$

1. Multiply **M** by \mathbf{R}_x : the resulting term at (3,2) will be 0 because of the values selected for c and s
2. Multiply resulting matrix by \mathbf{R}_y such that resulting term at (3,1) is zero (select c' and s' accordingly)
3. Multiply resulting matrix by \mathbf{R}_z such that resulting term at (2,1) is zero (select c'' and s'' accordingly)

Camera Matrix Decomposition: Computing the Intrinsics **K** and Rotation **R**

- $\mathbf{P} = \mathbf{K}[\mathbf{R} \mid -\mathbf{RC}] = [\mathbf{M} \mid -\mathbf{MC}]$

2. Algebraic derivation:

- Given rotations: $\mathbf{R}_x = \begin{bmatrix} 1 & 0 & 0 \\ 0 & c & -s \\ 0 & s & c \end{bmatrix}$ $\mathbf{R}_y = \begin{bmatrix} c' & 0 & s' \\ 0 & 1 & 0 \\ -s' & 0 & c' \end{bmatrix}$ $\mathbf{R}_z = \begin{bmatrix} c'' & -s'' & 0 \\ s'' & c'' & 0 \\ 0 & 0 & 1 \end{bmatrix}$

- $c = -\frac{m_{33}}{\sqrt{m_{32}^2 + m_{33}^2}}$ $s = \frac{m_{32}}{\sqrt{m_{32}^2 + m_{33}^2}}$

- Then,

- $\mathbf{K} = \mathbf{M}\mathbf{R}_x\mathbf{R}_y\mathbf{R}_z$


- $\mathbf{M} = \mathbf{K}\mathbf{R}_z^\top\mathbf{R}_y^\top\mathbf{R}_x^\top \Rightarrow \mathbf{R} = \mathbf{R}_z^\top\mathbf{R}_y^\top\mathbf{R}_x^\top$

Camera Calibration

Geometric Solvers for Recovering the Projection Matrix

Camera Calibration with Geometric Solvers


- So far, minimisation of an **algebraic error** criterion
 - Advantage: **linear solutions**
 - Disadvantage: no explicit geometric meaning
- Refinement:
 - Nonlinear minimisation of a **geometric error**



Camera Calibration with Geometric Solvers

- Minimise objective function over \mathbf{P} :

$$\begin{aligned} & \min_{\mathbf{P}} \sum_i d(\mathbf{x}_i, \hat{\mathbf{x}}_i)^2 \\ &= \min_{\mathbf{P}} \sum_i d(\mathbf{x}_i, \mathbf{P}\mathbf{X}_i)^2 \\ &= \min_{\mathbf{P}} \sum_i \left(\frac{x_i}{w_i} - \frac{\mathbf{p}_1^\top \mathbf{X}_i}{\mathbf{p}_3^\top \mathbf{X}_i} \right)^2 + \left(\frac{y_i}{w_i} - \frac{\mathbf{p}_2^\top \mathbf{X}_i}{\mathbf{p}_3^\top \mathbf{X}_i} \right)^2 \\ & \text{for } \mathbf{x}_i = (x_i, y_i, w_i)^\top \text{ and } \mathbf{P} = \begin{bmatrix} \mathbf{p}_1^\top \\ \mathbf{p}_2^\top \\ \mathbf{p}_3^\top \end{bmatrix} \end{aligned}$$



- Nonlinear optimisation, so going to be slow

“Gold Standard” Algorithm for Camera Calibration

1. Compute an *initial* solution using the normalised DLT algorithm
2. Refine the normalised solution using iterative minimisation of a geometric error
 - Use a nonlinear solver, like lsqnonlin in Matlab
3. Denormalise solution


Summary: DLT Camera Calibration

- Advantages:
 - Very simple to formulate and solve
 - Can recover \mathbf{K} , \mathbf{R} , \mathbf{C} from \mathbf{P} using RQ decomposition [Golub & VanLoan 96]
- Disadvantages:
 - Does not compute internal parameters explicitly
 - Sometimes involves more unknowns than true degrees of freedom
 - Need a separate camera matrix for each new view


Practical / Popular Camera Calibration Algorithms

Practical Camera Calibration

1. Load images into a calibration toolbox
 2. Calibrate
- Example toolboxes:
 - C++: OpenCV
 - Matlab: Calibration toolbox by Jean-Yves Bouguet
http://www.vision.caltech.edu/bouguetj/calib_doc/



Multi-plane Calibration



Multi-plane Calibration: Zhang's Method



- Use several images of a planar target held at *unknown* poses [Zhang PAMI 99]

1. Compute plane homographies:

$$\begin{bmatrix} u_i \\ v_i \\ 1 \end{bmatrix} \sim \mathbf{K}[\mathbf{r}_1 \quad \mathbf{r}_2 \quad \mathbf{t}] \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix} \sim \mathbf{H}\mathbf{X}$$

2. Solve for $\mathbf{K}^{-\top}\mathbf{K}^{-1}$ from \mathbf{H}_k 's

- 1 plane if only f unknown
 - 2 planes if (f, u_c, v_c) unknown
 - 3+ planes for full \mathbf{K}
- Code available on OpenCV



Camera Resectioning

Absolute Camera Pose


Overview

- Objective:
 - Given 2D–3D point correspondences $\{\mathbf{x}_i \leftrightarrow \mathbf{X}_i\}$ & the camera intrinsics \mathbf{K} , estimate the position \mathbf{C} and orientation \mathbf{R} of the camera
- Synonyms:
 - Camera resectioning
 - Absolute camera pose estimation
 - Perspective-n-point problem

Motivation


- If we have an existing reconstruction and a calibrated camera, and we just want to know where our camera is right now
 - Offline: calibration, e.g., with the DLT algorithm
 - Online: absolute pose, relative pose, triangulation
 - Unless there is mechanical change, \mathbf{K} will remain constant [why might it change?]
- But didn't we just work out how to recover \mathbf{R} and \mathbf{C} ?
 - Recovering from the camera matrix involves estimating 11 DoFs
 - This is many more degrees-of-freedom than we need

Perspective-n-Point Problem




Perspective-n-Point Problem

- Degrees-of-freedom:
 - 6: 3 (translation) + 3 (rotation)
- Minimal solution:
 - 3 point correspondences
 - Perspective-3-point (P3P) problem



Perspective-n-Point Problem

- One solution
 - Haralick et al., IJCV 1994
- Variable elimination leads to a 4th order polynomial
 - 4 solutions
 - Use a fourth point correspondence to disambiguate



$$s_2^2 + s_3^2 - 2s_2s_3 \cos \alpha = a^2$$

$$s_1^2 + s_3^2 - 2s_1s_3 \cos \beta = b^2 \quad \text{Cosine rule}$$

$$s_1^2 + s_2^2 - 2s_1s_2 \cos \gamma = c^2$$

Next Lecture

- Two-view geometry:
 - Homographies
 - Homography estimation
 - Epipolar geometry