

# 3D Vision 3

Week 9

Two-view Geometry: Triangulation

Two-view Geometry: Stereo

# Announcements

- **Survey 2 – Week 8|9** is live on Wattle, please provide a response to improve the course for future years
- **Class Representatives:** another mechanism to provide feedback
  - Arjun Raj (u7526852@anu.edu.au) [COMP4528]
  - Ruxuan Li (u7763307@anu.edu.au) [ENGN4528]
  - Jiabao Han (u7765516@anu.edu.au) [COMP6528]
  - Xingchen Zhang (u7670173@anu.edu.au) [COMP6528]
- **Exam:** 9:00am–12:15pm Saturday 1 June 2024
  - 7-11 Barry Drive, First Floor Left Side

# Weekly Study Plan: Overview

Wk	Starting	Lecture	Lab	Assessment
1	19 Feb	Introduction	X	
2	26 Feb	Low-level Vision 1	1	
3	4 Mar	Low-level Vision 2	1	
		Mid-level Vision 1		
4	11 Mar	Mid-level Vision 2	1	CLab1 report due Friday
		High-level Vision 1		
5	18 Mar	High-level Vision 2	2	
6	25 Mar	High-level Vision 3 <sup>1</sup>	2	
	1 Apr	Teaching break	X	
	8 Apr	Teaching break	X	
7	15 Apr	3D Vision 1	2	CLab2 report due Friday
8	22 Apr	3D Vision 2	3	
9	29 Apr	3D Vision 3	3	
10	6 May	3D Vision 4	3	
		Mid-level Vision 3		
11	13 May	High-level Vision 4	X	CLab3 report due Friday
12	20 May	Course Review	X	

# Weekly Study Plan: Part B

Wk	Starting	Lecture	By
7	15 Apr	3D vision: introduction, camera model, single-view geometry	Dylan
8	22 Apr	3D vision: camera calibration, two-view geometry (homography)	Dylan
9	29 Apr	3D vision: two-view geometry (epipolar geometry, triangulation, stereo)	Dylan
10	6 May	3D vision: multiple-view geometry	Weijian
		Mid-level vision: optical flow, shape-from-X	Dylan
11	13 May	High-level vision: self-supervised learning, detection, segmentation	Dylan
12	20 May	Course review	Dylan

# Outline

1. Two-view Geometry: Triangulation
2. Two-view Geometry: Stereo
3. Single-view Geometry: Monocular Depth
4. Single-view Geometry: Optical flow

# Essential & Fundamental Matrices: Summary

- Algebraic representations of epipolar geometry:
  - Projection matrices (given intrinsics + extrinsics)
  - **Essential matrix** (given intrinsics):  $\mathbf{E} = [\mathbf{t}]_{\times} \mathbf{R}$
  - **Fundamental matrix**:  $\mathbf{F} = \mathbf{K}'^{-T} [\mathbf{t}]_{\times} \mathbf{R} \mathbf{K}^{-1} = [\mathbf{e}']_{\times} \mathbf{K}' \mathbf{R} \mathbf{K}^{-1}$  [HZ p.244]
    - Using identity  $[\mathbf{v}]_{\times} \mathbf{M} = \mathbf{M}^{-T} [\mathbf{M}^{-1} \mathbf{v}]_{\times}$  and  $\mathbf{e}' = \mathbf{K}' \mathbf{t}$  (image of camera centre 1)
- Epipolar constraint for corresponding points  $\{\mathbf{x}, \mathbf{x}'\}$ :
$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$$
  - $\mathbf{F}$  is rank 2 and is known only up to scale  $\rightarrow$  7 DoF
  - What is  $\mathbf{F} \mathbf{x}$ ?
- **Estimation**: DLT (again!); assemble a matrix  $\mathbf{A}$ , compute the SVD, enforce rank 2 (with another SVD); or use a nonlinear solver

# The Fundamental Matrix F: Recovering P & P'

- Since we can set the world coordinate system arbitrarily, set

$$\begin{aligned}P &= [I \mid 0] \\P' &= [[e']]_x F \mid e']\end{aligned}$$

where epipole can be found by solving

$$F^\top e' = 0$$

via SVD/eigendecomposition [HZ p. 256]

# Triangulation

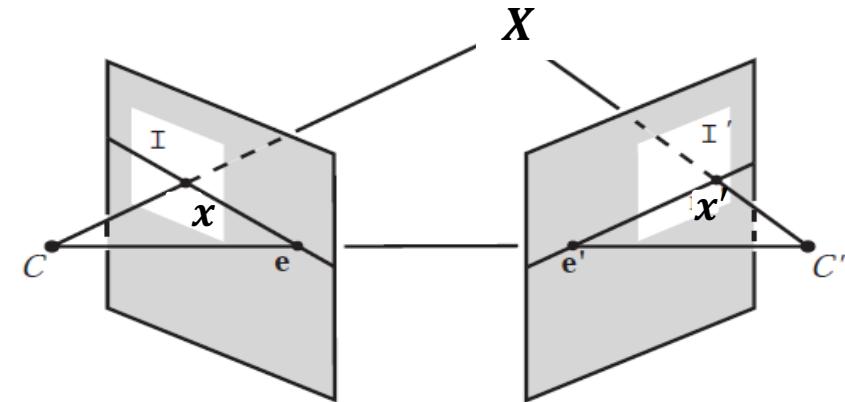
Two-view Geometry

# Triangulation

- **Problem:** Given image points  $x$  and  $x'$  in *correspondence* across two or more images (taken from calibrated cameras with known intrinsics and extrinsics), compute the 3D location  $X$

# Triangulation: DLT Linear Solution

- Generally, rays  $\overrightarrow{Cx}$  and  $\overrightarrow{C'x'}$  will not intersect exactly
- Can solve via the DLT algorithm (again!), finding a least squares solution to a system of equations
- Further reading:
  - HZ pp. 312–313



$$x \times (P X) = 0 \text{ and } x' \times (P' X) = 0$$

$$\therefore AX = \begin{bmatrix} u\mathbf{p}_3^T - \mathbf{p}_1^T \\ v\mathbf{p}_3^T - \mathbf{p}_2^T \\ u'\mathbf{p}'_3^T - \mathbf{p}'_1^T \\ v'\mathbf{p}'_3^T - \mathbf{p}'_2^T \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = 0$$

# Triangulation: DLT Linear Solution

Given  $\mathbf{P}, \mathbf{P}', \mathbf{x}, \mathbf{x}'$

1. Precondition points and projection matrices
2. Create matrix  $\mathbf{A}$
3. SVD:  $\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^T$
4.  $\mathbf{X} = \mathbf{V}_{:, -1}$  (last column of  $\mathbf{V}$ )
5. Then refine with respect to a geometric error

$$\mathbf{x} = \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \quad \mathbf{x}' = \begin{bmatrix} u' \\ v' \\ 1 \end{bmatrix}$$

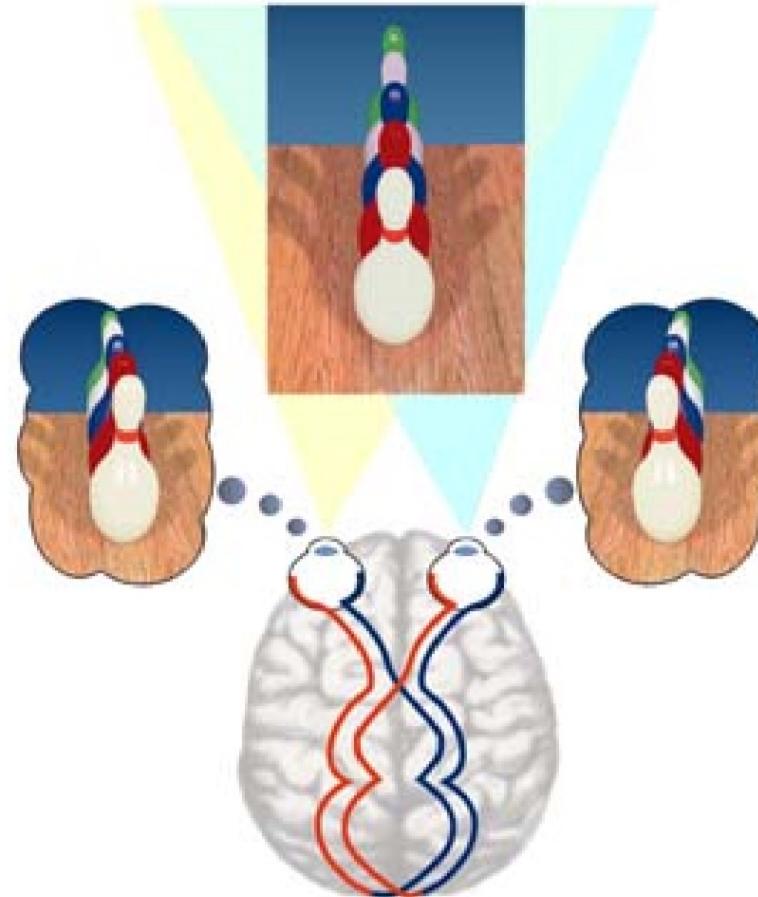
$$\mathbf{P} = \begin{bmatrix} \mathbf{p}_1^T \\ \mathbf{p}_2^T \\ \mathbf{p}_3^T \end{bmatrix} \quad \mathbf{P}' = \begin{bmatrix} \mathbf{p}'_1^T \\ \mathbf{p}'_2^T \\ \mathbf{p}'_3^T \end{bmatrix}$$

$$\mathbf{A} = \begin{bmatrix} u\mathbf{p}_3^T - \mathbf{p}_1^T \\ v\mathbf{p}_3^T - \mathbf{p}_2^T \\ u'\mathbf{p}'_3^T - \mathbf{p}'_1^T \\ v'\mathbf{p}'_3^T - \mathbf{p}'_2^T \end{bmatrix}$$

# Stereo Vision

Two-view Geometry

# Binocular Stereo Vision



# 3D Perception Using Stereo Vision

- Humans do it well from a *single* image and very well from *stereo* (binocular) images
- Mechanism is not well understood:
  - Biological components: some understanding
  - Algorithm: little understanding
- An intermediate goal of computer vision was to mimic the functionality of the biological system
  - *Not* to mimic its mechanics

# Stereo Vision

- Stereo matching computes depth information from two images
- Depth information can be used to...
  - Differentiate objects from the background
  - Differentiate objects from one another
  - Expose camouflaged objects
  - Navigate in environment avoiding obstacles

# Stereo Vision: Two Sub-Problems

## **Correspondence Problem**

- The problem of measuring the disparity of each point in the two eye (camera) projections

## **Interpretation Problem**

- The use of disparity information to recover the orientation and distance of surfaces in the scene

# Stereo Vision: Algorithmic Steps

- Basic steps to be performed in any stereo imaging system:
  - Image Acquisition
  - Camera Modelling
  - Feature Extraction
  - Image Matching
  - Depth Interpolation

# Stereo Image Acquisition

- Capturing two images with a very specific camera geometry

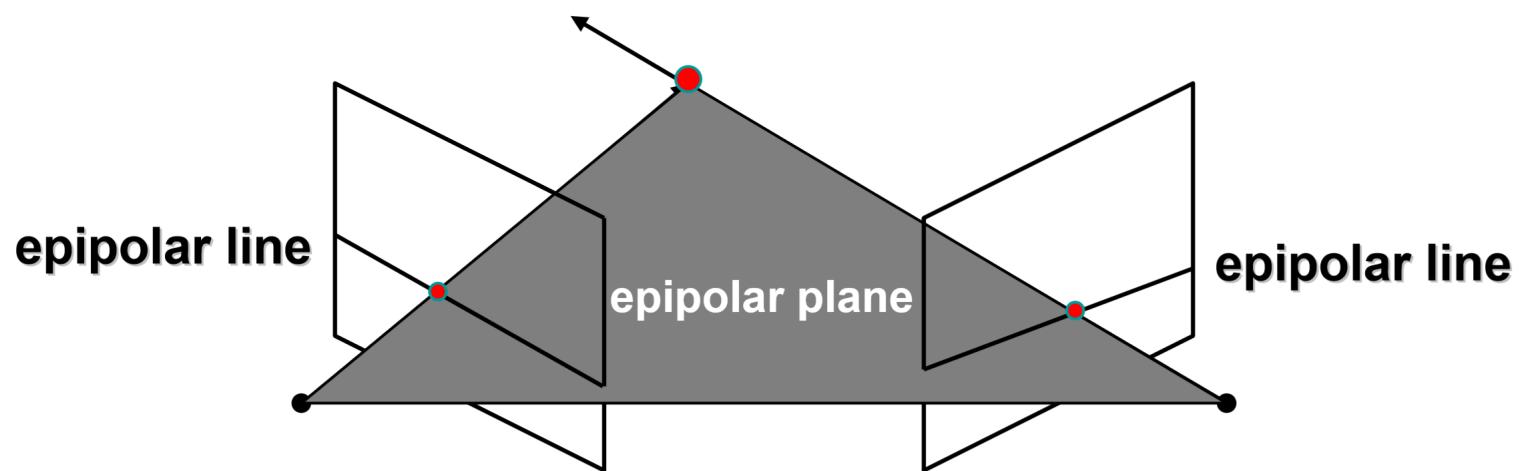


# Camera Modelling

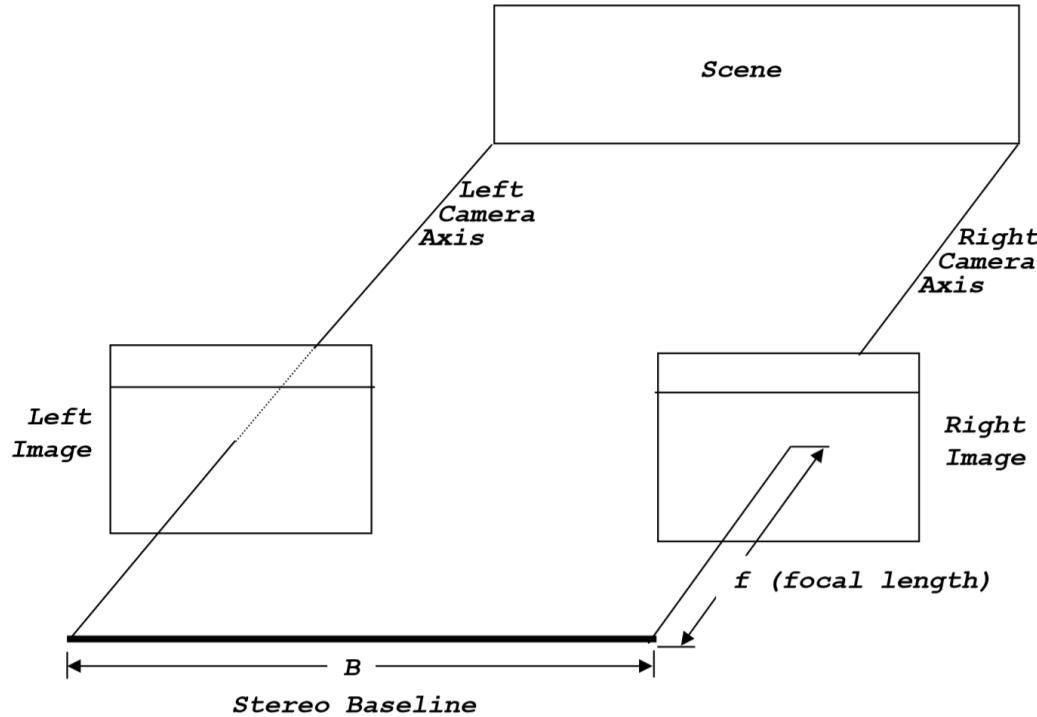
- For accurate depth results the camera parameters must be known
  - Intrinsic parameters: focal length, pixel skew/shape, distortion, etc.
  - Extrinsic parameters: relative rotation and translation



# General Epipolar Constraint



# Simple (Ideal) Stereo Geometry



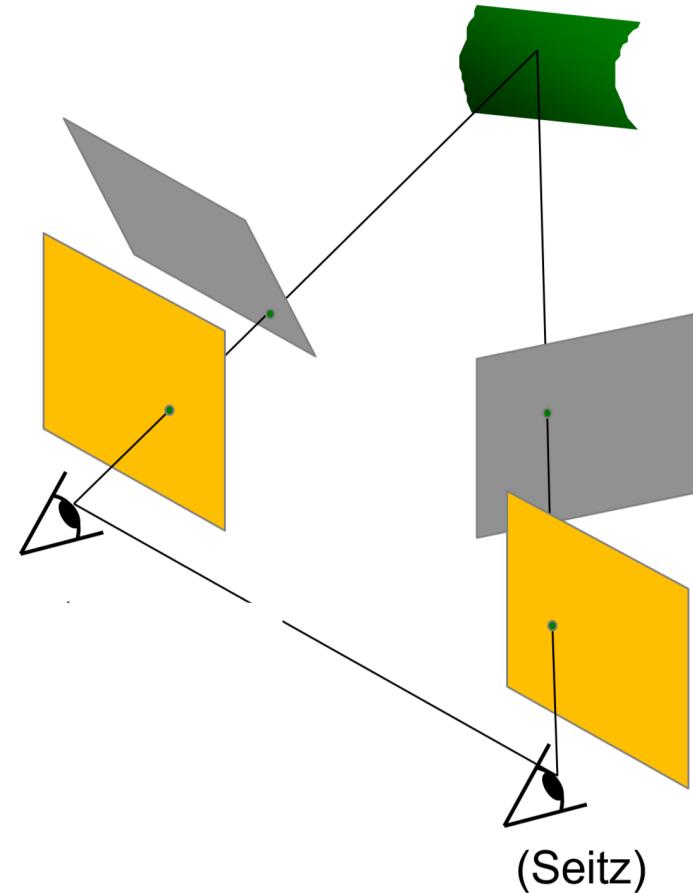
- Two slightly different images
- Extrinsics: identical rotations, translation has an offset in the X direction

# Simplest Case: Rectified Stereo Images

- Assumptions:
  - Image planes of cameras are parallel
  - Principal axes are at same height
  - Identical focal lengths
- Then the epipolar lines fall along the horizontal scan lines of the images
- We will assume images have been rectified so that epipolar lines correspond to scan lines:
  - Simplifies algorithms
  - Improves efficiency

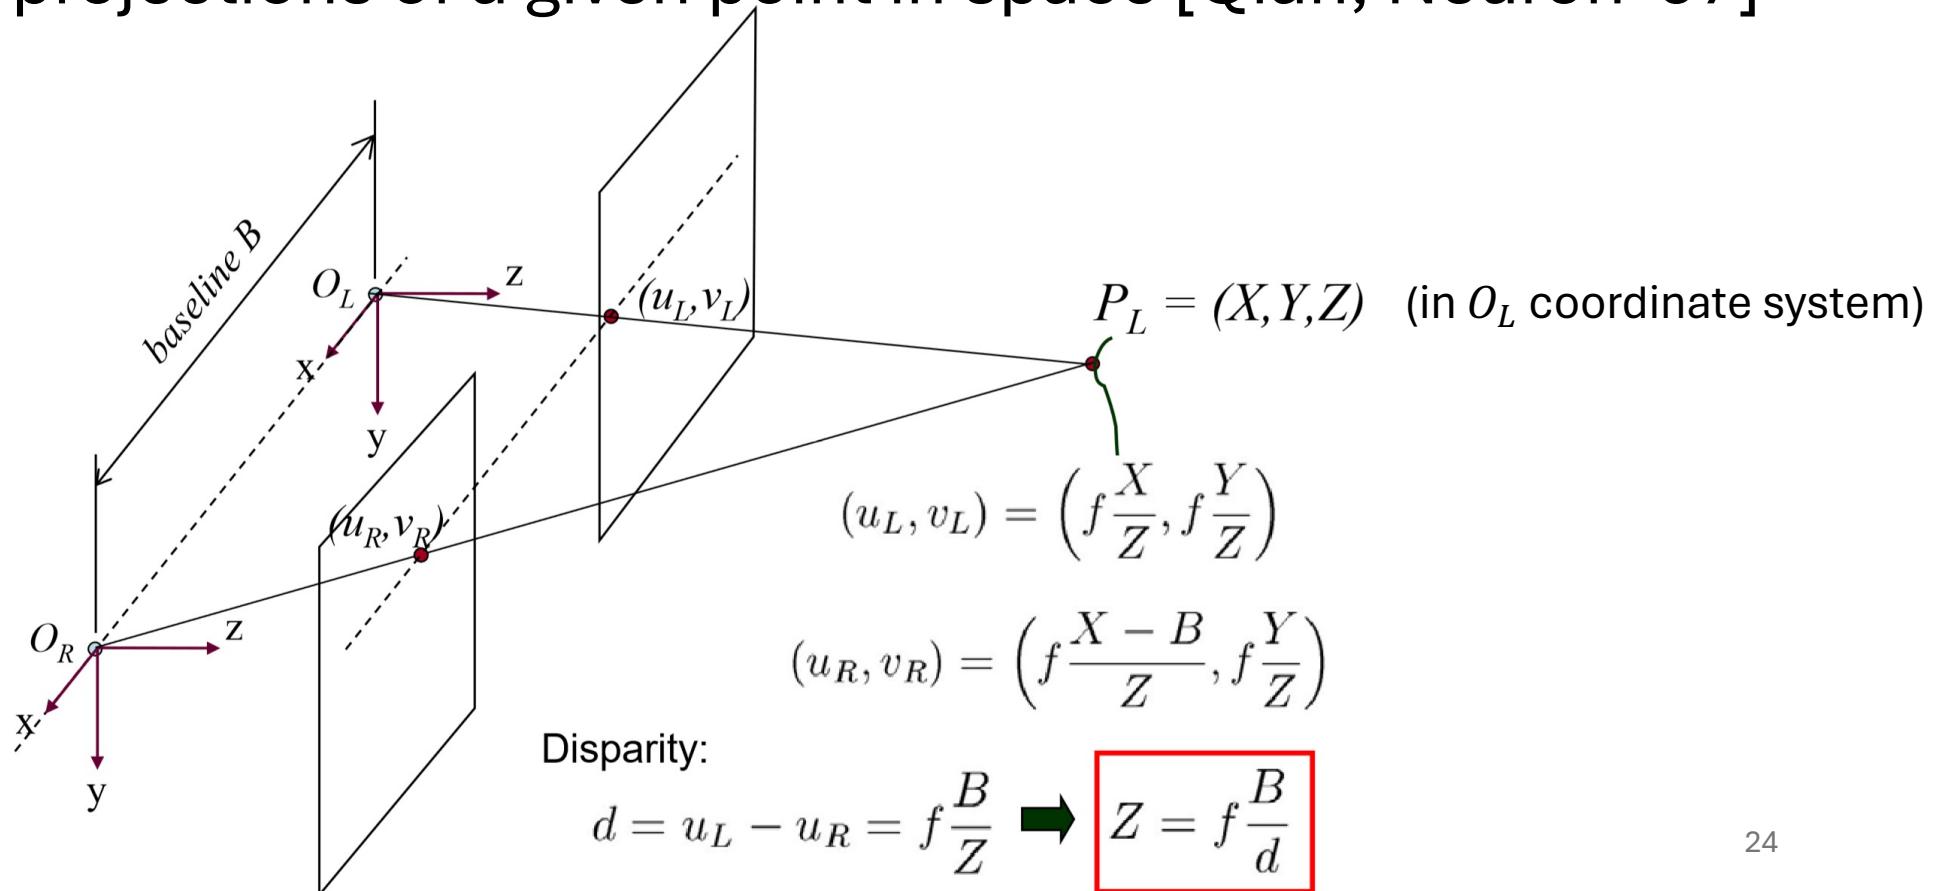
# Simplest Case: Rectified Stereo Images

- We can always achieve this ideal geometry with stereo-rectification
- Image reprojection:
  - Reproject image planes onto a common plane parallel to the line between optical centres using suitable planar homographies
  - E.g., using homography estimation
  - Not covered in detail here



# Disparity-to-Depth Computation

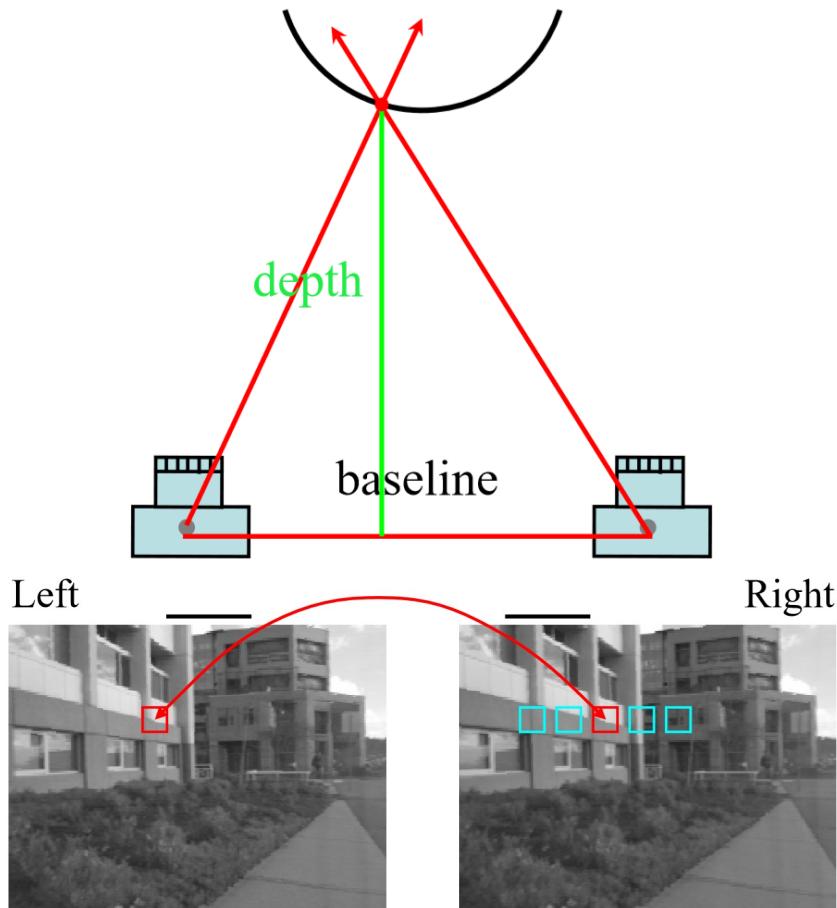
- Disparity (human vision): the positional difference between the two retinal projections of a given point in space [Qian, Neuron '97]



# The Correspondence Problem

- What should we match?
  - Objects?
  - Edges?
  - Pixels?
  - Super-pixels (set of pixels)?
  - ...
  - DINO features?

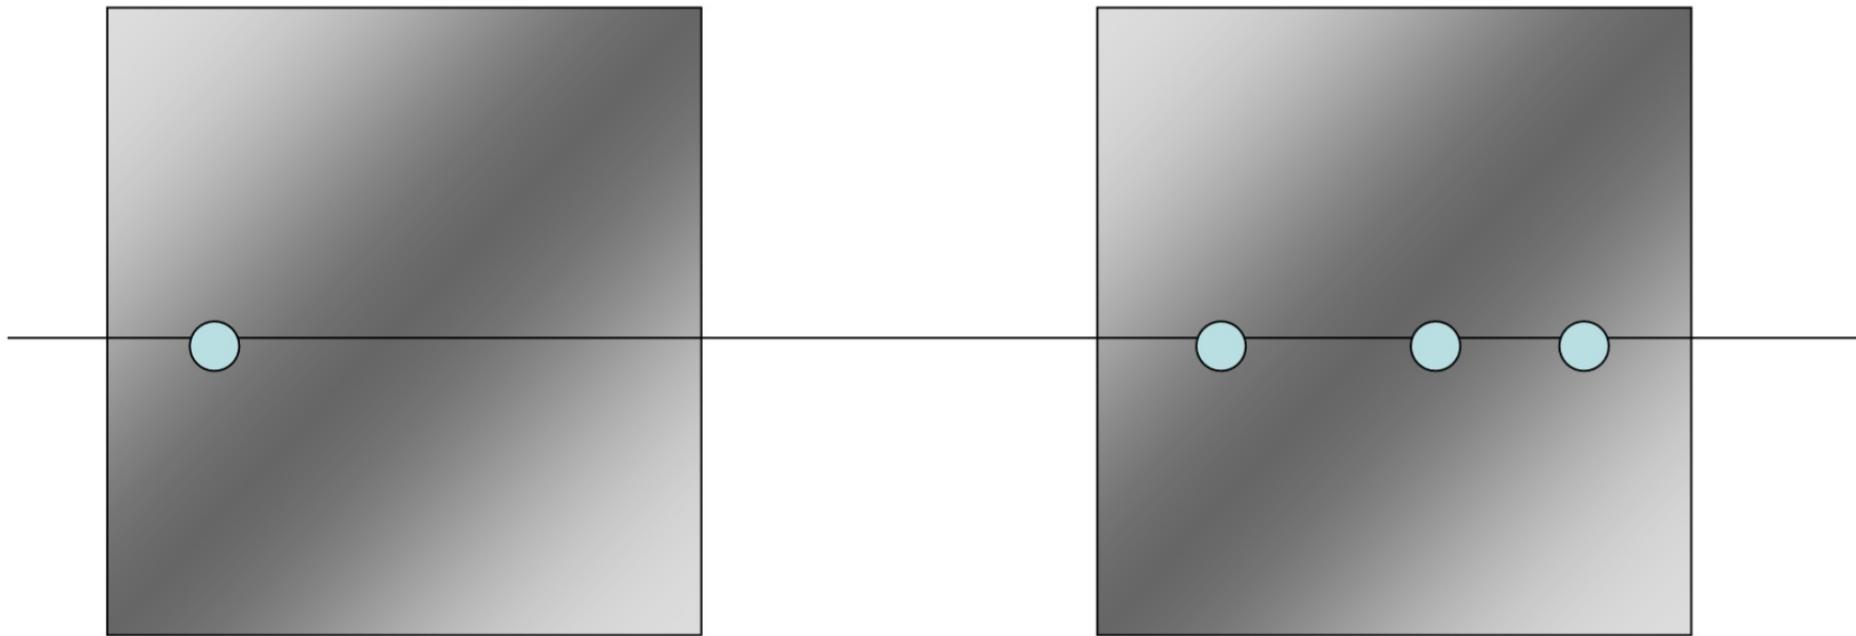
# The Correspondence Problem



- Triangulate on two images of the same point to recover depth
  - Feature matching across views
  - Calibrated cameras
- Matching windows across horizontal scan lines

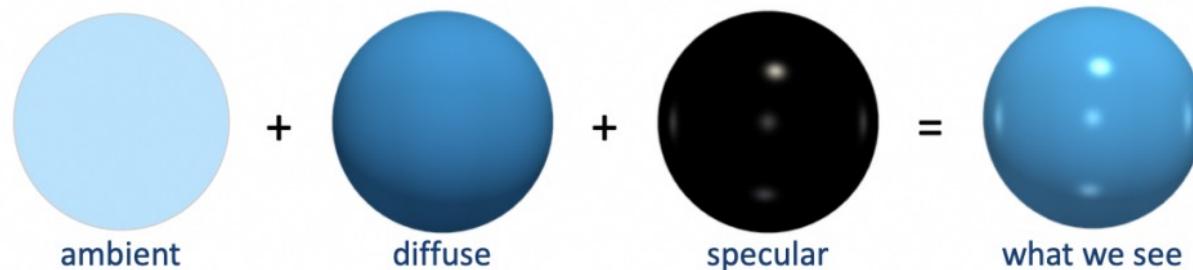
# The Correspondence Problem: The Epipolar Constraint

- The epipolar constraint removes some ambiguity

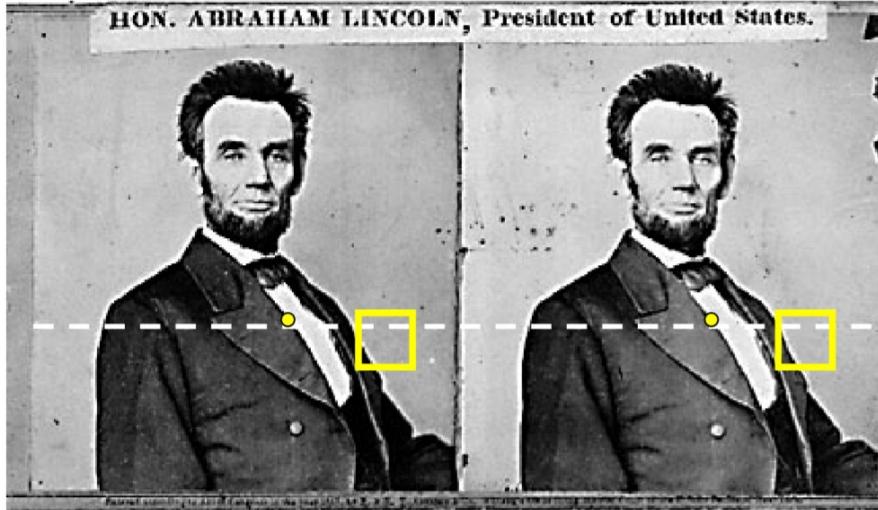


# The Correspondence Problem: The Colour Constancy “Constraint”

- Same world point has same intensity (or colour) in both images
  - True for Lambertian surfaces
    - A Lambertian surface has a brightness that is independent of viewing angle
  - Violations:
    - Noise
    - Specularity
    - Non-Lambertian materials
    - Pixels that contain multiple surfaces
    - Occlusions

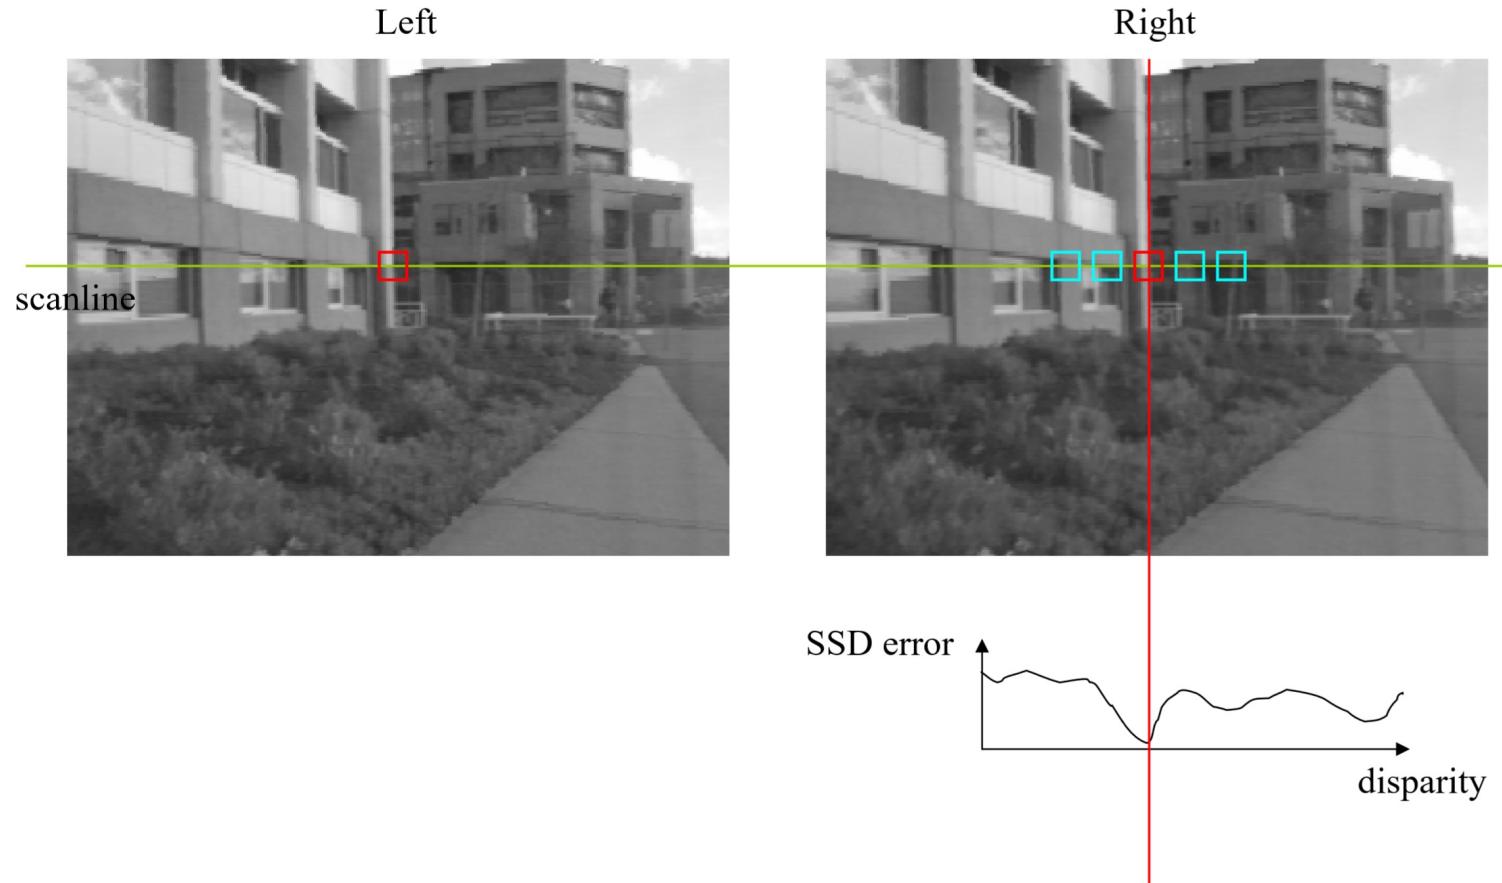


# The Correspondence Problem: Pixel Matching

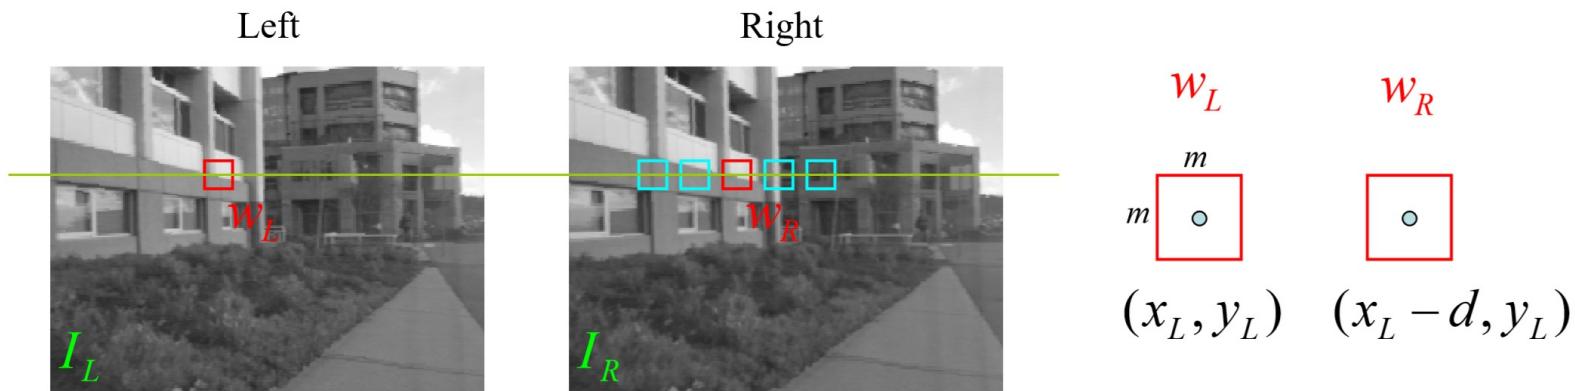


- For each epipolar line:
  - For each pixel in the left image:
    - Compare with every pixel on same epipolar line in right image
    - Pick pixel with minimum matching cost
- Still too much ambiguity: match windows instead

# The Correspondence Problem: Correspondence Using Correlation



# Sum of Squared (Pixel) Differences



$w_L$  and  $w_R$  are corresponding  $m$  by  $m$  windows of pixels.

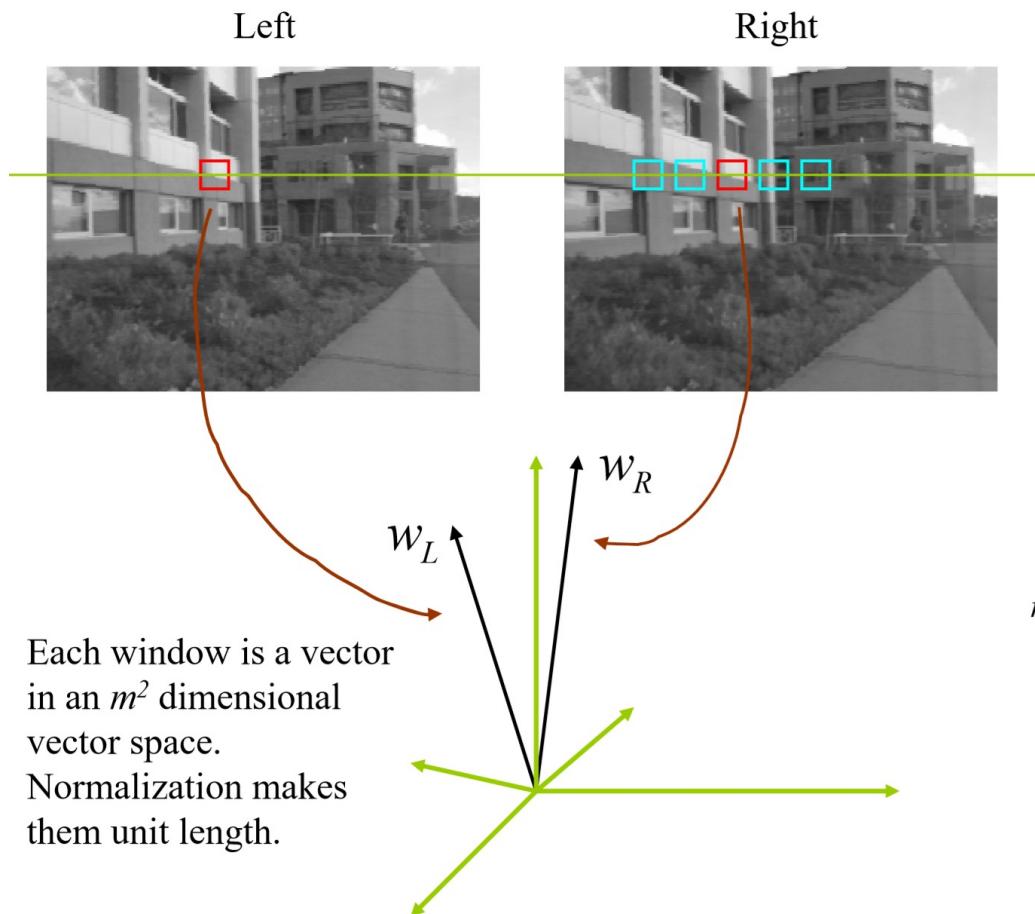
We define the window function :

$$W_m(x, y) = \{u, v \mid x - \frac{m}{2} \leq u \leq x + \frac{m}{2}, y - \frac{m}{2} \leq v \leq y + \frac{m}{2}\}$$

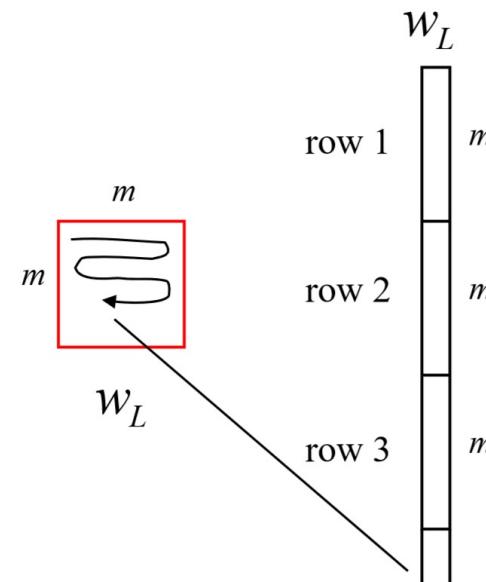
The SSD cost measures the intensity difference as a function of disparity :

$$C_r(x, y, d) = \sum_{(u, v) \in W_m(x, y)} [I_L(u, v) - I_R(u - d, v)]^2$$

# Image Patches as Vectors



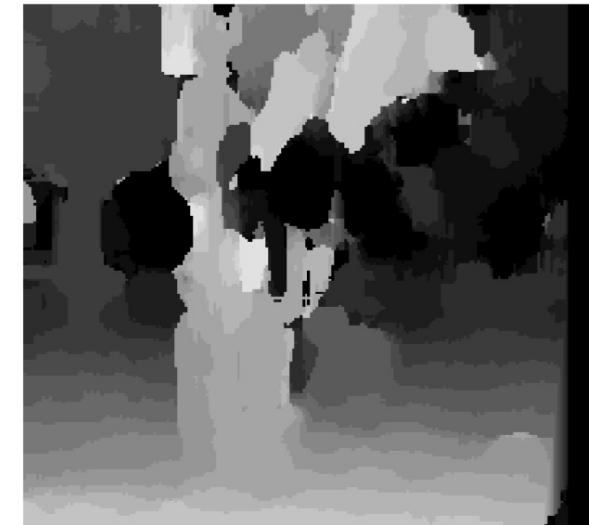
“Unwrap”  
image to form  
vector, using  
raster scan order



# Effect of Window Size



$W = 3$



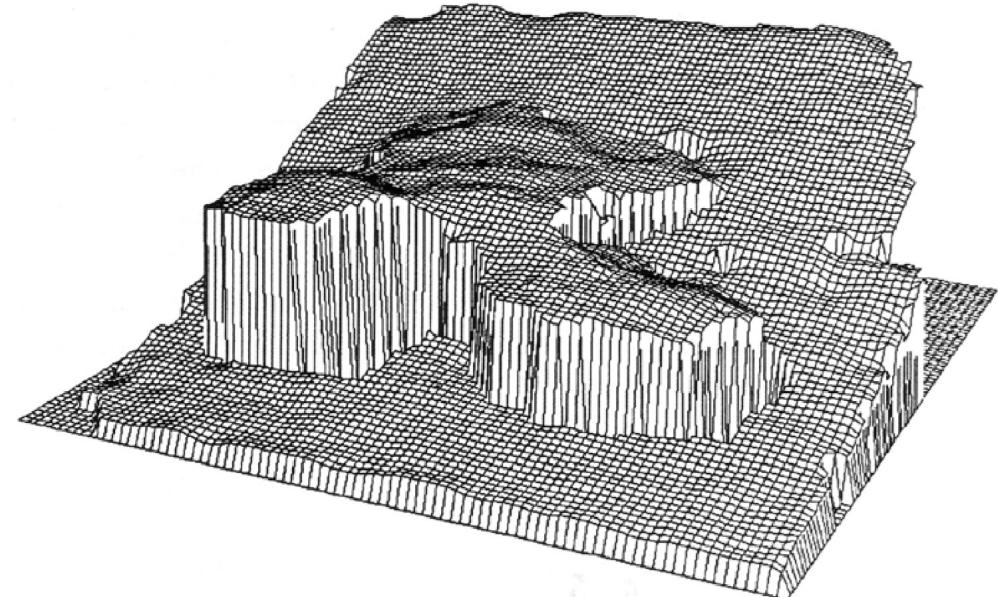
$W = 20$

- Improvement: use an adaptive window size (try multiple sizes and select best match)

# Example Stereo Pair

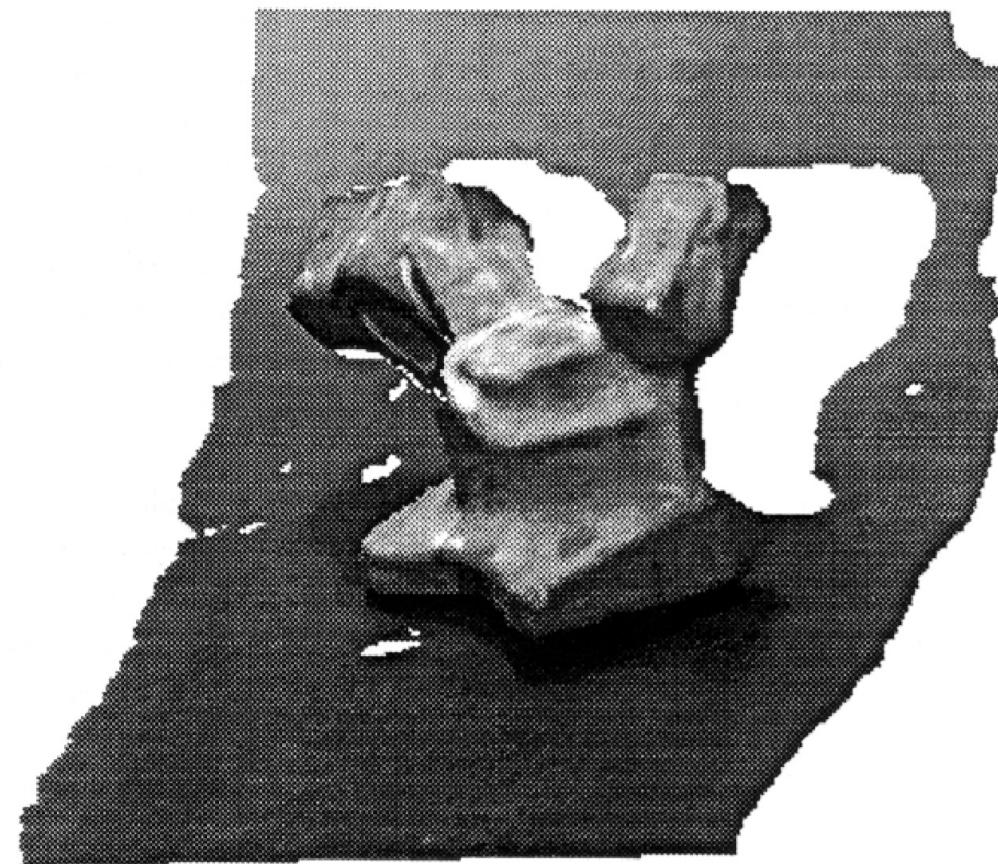


**Left Camera**



- Obtained depth map in 3D

# Novel View Synthesis

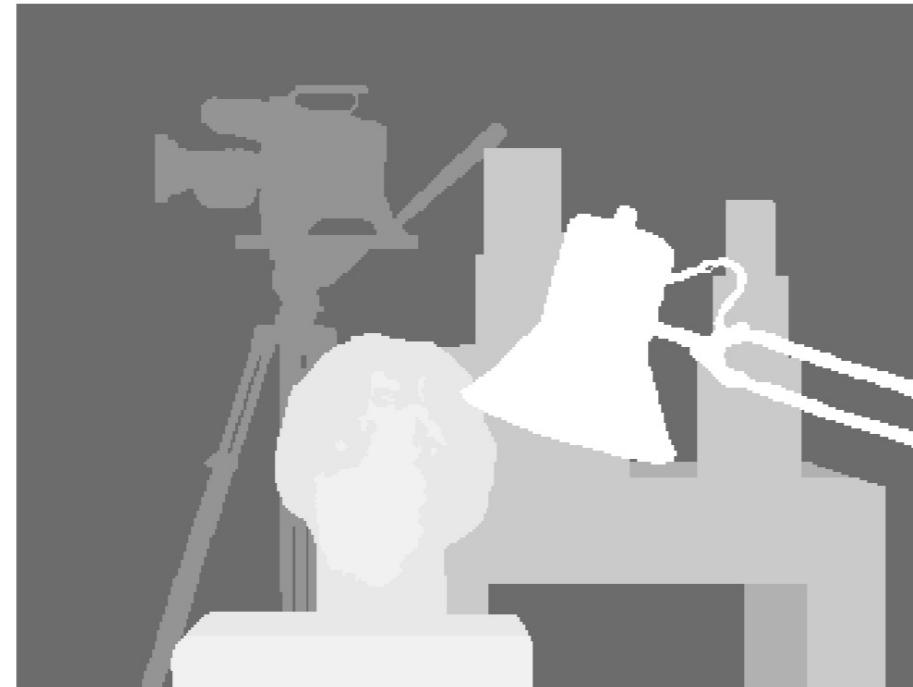


# A Taxonomy of Stereo Algorithms

- D. Scharstein and R. Szeliski, “A Taxonomy and Evaluation of Dense Two-Frame Stereo Correspondence Algorithms”, International Journal of Computer Vision, 47 (2002), pp. 7-42



Scene



Ground truth

# A Taxonomy of Stereo Algorithms



True disparities



19 – Belief propagation



11 – GC + occlusions



20 – Layered stereo



10 – Graph cuts



\*4 – Graph cuts



13 – Genetic algorithm



6 – Max flow



12 – Compact windows



9 – Cooperative alg.



15 – Stochastic diffusion



\*2 – Dynamic prgr.



14 – Realtime SAD



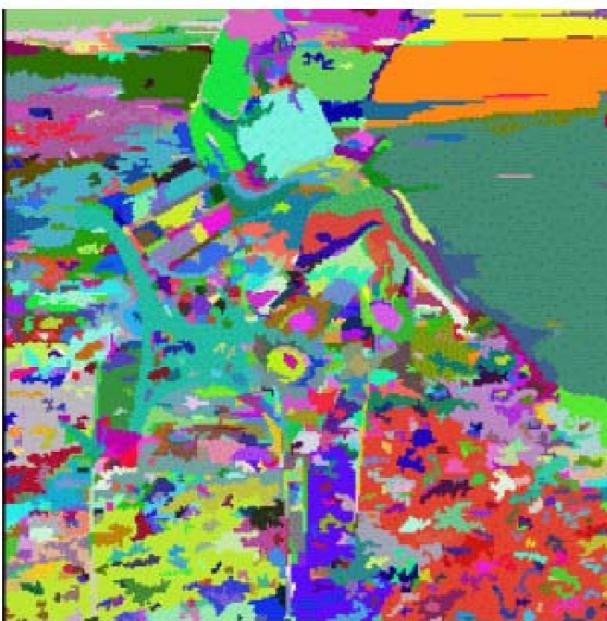
\*3 – Scanline opt.



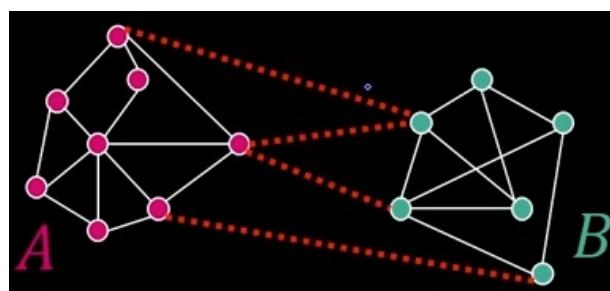
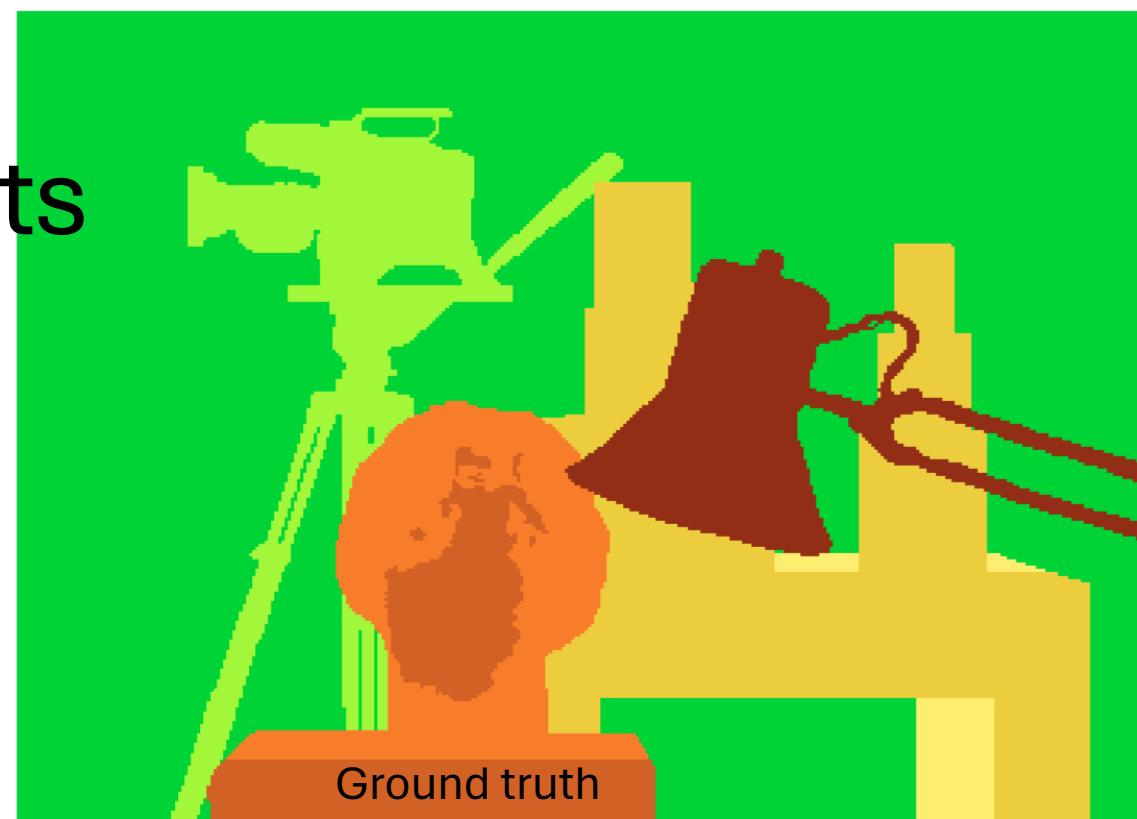
7 – Pixel-to-pixel stereo

\*1 – SSD+MF  
Scharstein and Szeliski

# Stereo from Segmentation



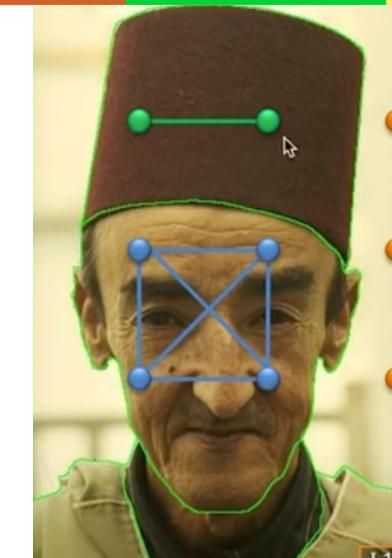
# Stereo from Graph-Cuts



An image is represented by a graph

$$Cut(A, B) = \sum_{p \in A, q \in B} w_{pq}$$

Minimise the cost of a cut,  
also called min-cut  
Each subgraph is an image  
segment



# Summary of Stereo Vision

- Stereo matching + triangulation
  - The correspondence problem and the interpretation problem
- Constraints:
  - Geometry: epipolar constraint
  - Photometric: brightness constancy “constraint”

# Next Week

- Multiple-view geometry: Structure-from-Motion
- Multiple-view geometry: Radiance fields (fancy triangulation)
- Multiple-view geometry: Learning-based pose-free reconstruction
- Mid-level vision: Optical flow
- Mid-level vision / single-view geometry: Shape-from-X