



Australian
National
University



COMP4650/6490 Document Analysis

Deep Neural Networks — Part I

ANU School of Computing

Administrative matters

- Assignment 2
 - Minor update: Fixed 2 small issues in `clustering.py`
- Assignment 1
 - Expected return date: Tuesday 29 August

- Neural networks & why NN
- Feedforward neural network
 - From logistic regression to Feedforward NN
 - Non-linear activation functions
- Back-propagation
 - Gradient descent
 - Back-propagation essentials
 - Computation graph
- Optimisers
 - Stochastic gradient descent
 - SGD with momentum & Adam

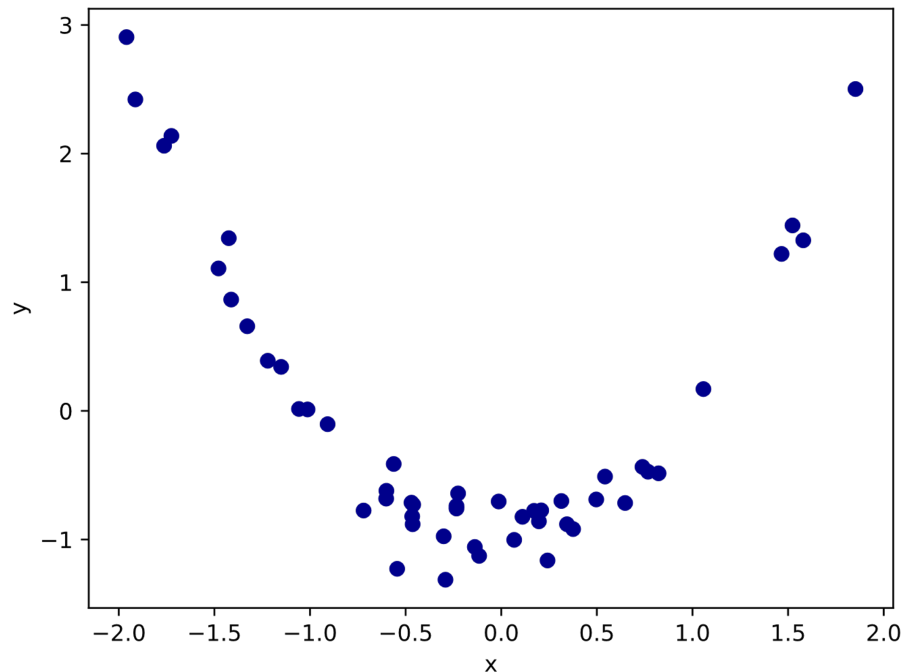
- **Neural networks & why NN**
- Feedforward neural network
 - From logistic regression to Feedforward NN
 - Non-linear activation functions
- Back-propagation
 - Gradient descent
 - Back-propagation essentials
 - Computation graph
- Optimisers
 - Stochastic gradient descent
 - SGD with momentum & Adam

Neural networks

- Computing systems initially inspired by *simplified* models of the brain
- We keep some of the ideas and terminology, but it is *NOT* a current model of the brain
- A better way to think of modern neural networks
 - Differentiable vector cascades
 - Complex functions: A neural network is a stack of linear and non-linear functions

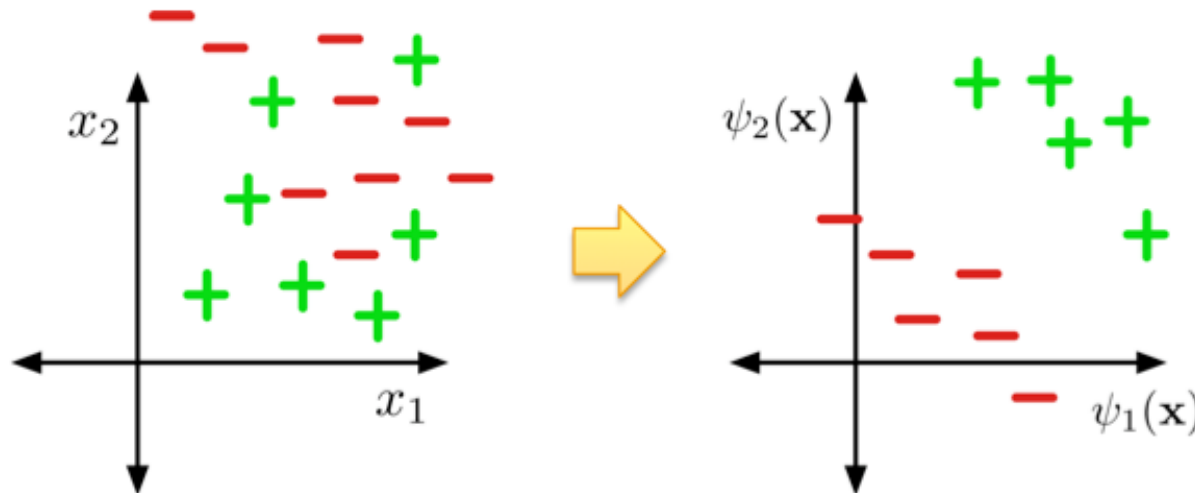
Why neural networks

- Non-linear relationships
 - What do we do if the target does not have a linear relationship with the input?
 - We need to fit a *non-linear* function to our data.



Why neural networks

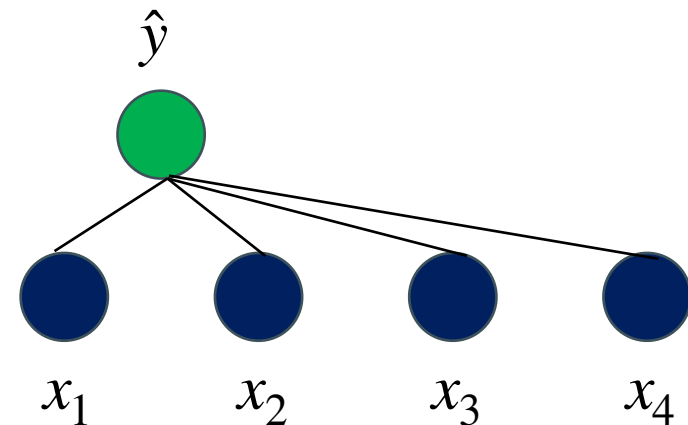
- Non-linear relationships
 - What do we do if the target does not have a linear relationship with the input?
 - We need to fit a *non-linear* function to our data.
- Feature learning
 - Neural nets can be viewed as a way of learning features



- Neural networks & why NN
- **Feedforward neural network**
 - From logistic regression to Feedforward NN
 - Non-linear activation functions
- Back-propagation
 - Gradient descent
 - Back-propagation essentials
 - Computation graph
- Optimisers
 - Stochastic gradient descent
 - SGD with momentum & Adam

From logistic regression to feedforward NN

- Binary logistic regression
 - A 1-layer NN (single output)
- Multinomial logistic regression
 - A 1-layer NN (multiple outputs)
- Multi-layer perceptron (MLP)
 - Stack multiple (non-)linear models
 - Fully connected feedforward NN
 - Composition of functions

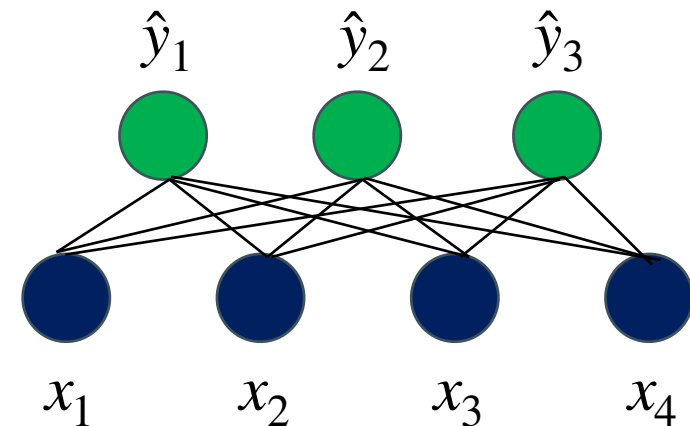


$$\hat{y} = \sigma(\mathbf{w}^T \mathbf{x} + b)$$

Feedforward NN

From logistic regression to feedforward NN

- Binary logistic regression
 - A 1-layer NN (single output)
- Multinomial logistic regression
 - A 1-layer NN (multiple outputs)
- Multi-layer perceptron (MLP)
 - Stack multiple (non-)linear models
 - Fully connected feedforward NN
 - Composition of functions

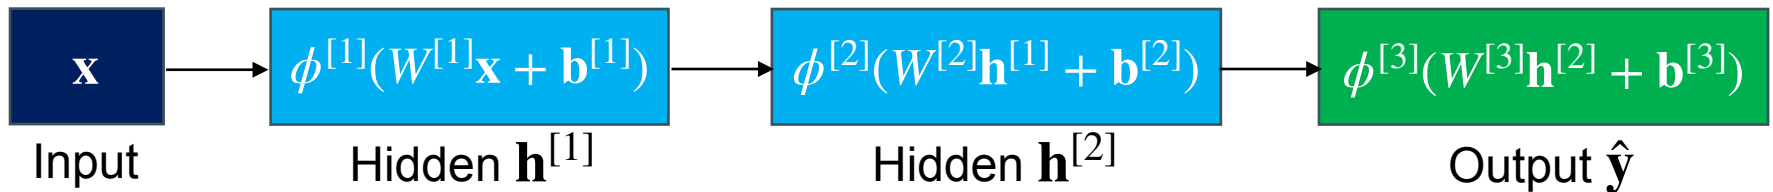
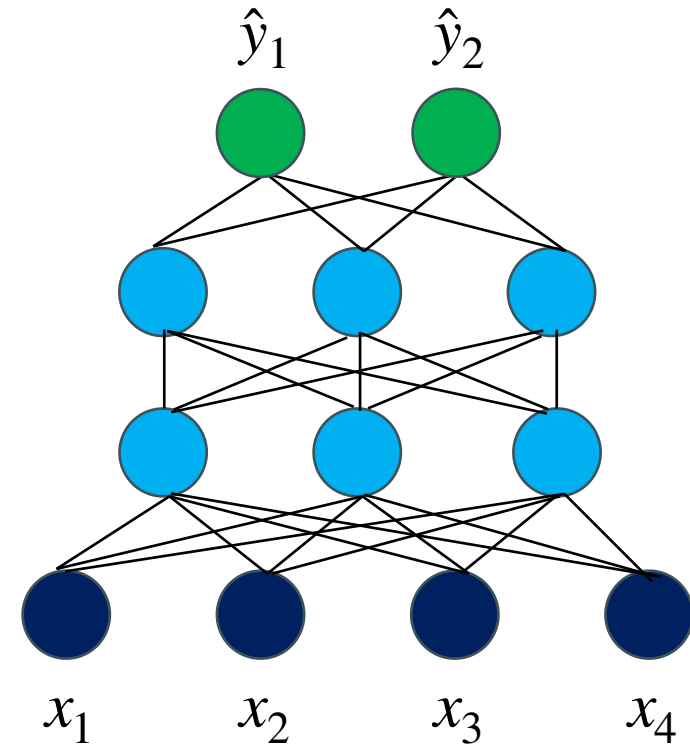


$$\hat{\mathbf{y}} = \text{softmax}(W\mathbf{x} + \mathbf{b})$$

Feedforward NN

From logistic regression to feedforward NN

- Binary logistic regression
 - A 1-layer NN (single output)
- Multinomial logistic regression
 - A 1-layer NN (multiple outputs)
- Multi-layer perceptron (MLP)
 - Stack multiple (non-)linear models
 - Fully connected feedforward NN
 - Composition of functions



Notation

Computation in layer l (superscript $[l]$ notation denotes layer l)

$$\mathbf{h}^{[l]} = \phi^{[l]}(W^{[l]}\mathbf{h}^{[l-1]} + \mathbf{b}^{[l]})$$

- $\mathbf{h}^{[l]}$: output of layer l
- $\phi^{[l]}$: activation function of layer l (usually a non-linear function)
- $W^{[l]}$: weight matrix of layer l
- $\mathbf{h}^{[l-1]}$: input of layer l , i.e. output of layer $l-1$, and $\mathbf{h}^{[0]} = \mathbf{x}$
- $\mathbf{b}^{[l]}$: biases of layer l , note that sometimes we write without the bias since it can be absorbed into $W^{[l]}$ by adding an extra dimension of value 1 to the input vector $\mathbf{h}^{[l-1]}$

Feedforward NN

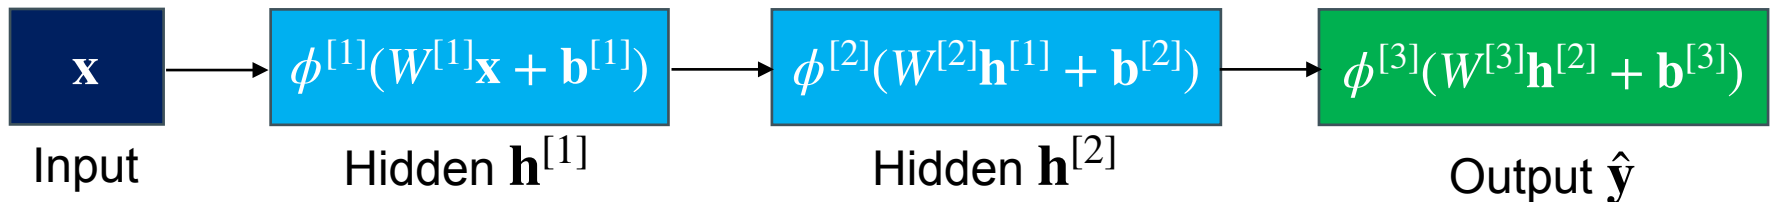
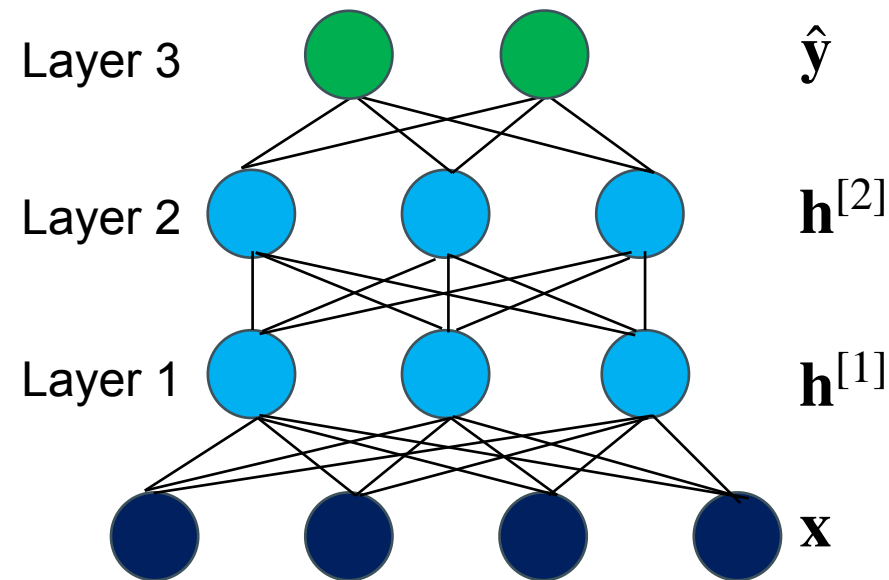
Notation

Computation in layer l (superscript $[l]$ notation denotes layer l)

$$\mathbf{h}^{[l]} = \phi^{[l]}(W^{[l]}\mathbf{h}^{[l-1]} + \mathbf{b}^{[l]})$$

Example:

- A 3-layer fully connected Feedforward NN (i.e. MLP) and its forward computation
- We usually don't count the input layer



Non-linear activation functions

Why use a non-linearity?

$$\begin{aligned} NN(\mathbf{x}) &= W^{[2]}(W^{[1]}\mathbf{x} + \mathbf{b}^{[1]}) + \mathbf{b}^{[2]} \\ &= W^{[2]}W^{[1]}\mathbf{x} + (W^{[2]}\mathbf{b}^{[1]} + \mathbf{b}^{[2]}) \end{aligned}$$

Non-linear activation functions

Why use a non-linearity?

$$\begin{aligned} NN(\mathbf{x}) &= W^{[2]}(W^{[1]}\mathbf{x} + \mathbf{b}^{[1]}) + \mathbf{b}^{[2]} \\ &= W^{[2]}W^{[1]}\mathbf{x} + (W^{[2]}\mathbf{b}^{[1]} + \mathbf{b}^{[2]}) \\ &= W'\mathbf{x} + \mathbf{b}' \end{aligned}$$

where $W' = W^{[2]}W^{[1]}$ and $\mathbf{b}' = W^{[2]}\mathbf{b}^{[1]} + \mathbf{b}^{[2]}$.

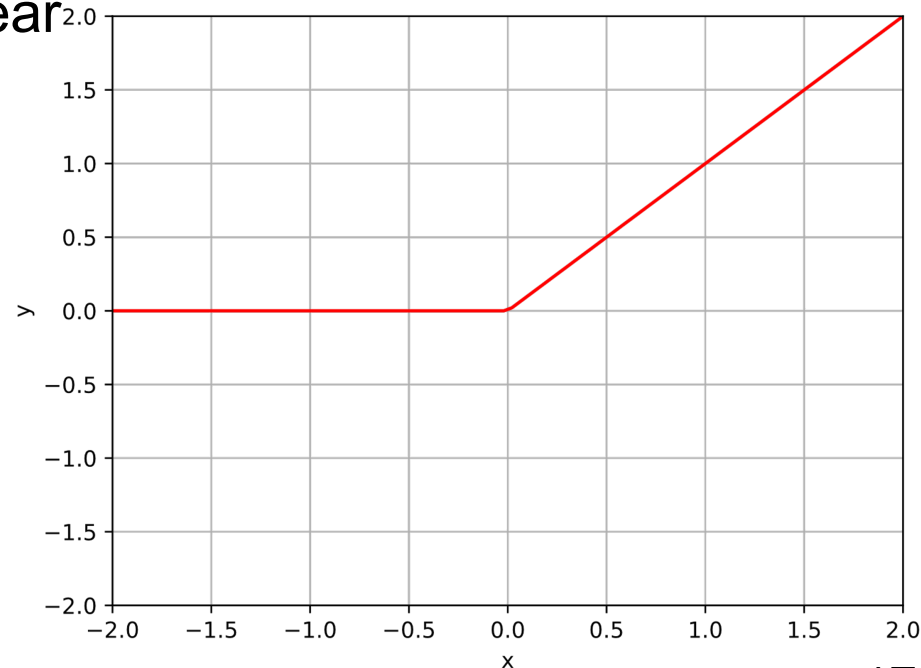
This is still a linear model!

Non-linear activation functions

Why use a non-linearity?

$$NN(\mathbf{x}) = W^{[2]}\phi(W^{[1]}\mathbf{x} + \mathbf{b}^{[1]}) + \mathbf{b}^{[2]}$$

- ϕ is *some* function that is non-linear and (mostly) differentiable
- ϕ is called the *activation function*
- In principle we can use any non-linear (mostly) differentiable function as the activation function
- Commonly used in practice is $\text{ReLU}(x) = \max(x, 0)$ which is fast to compute and works well



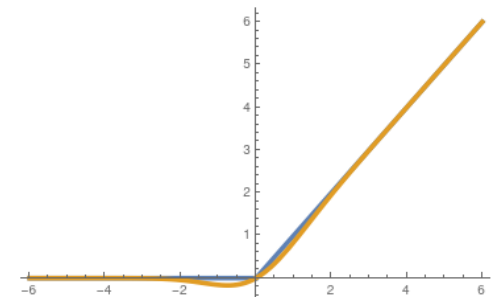
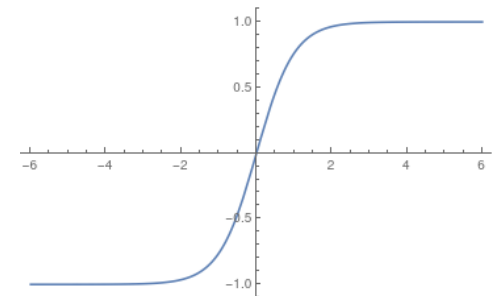
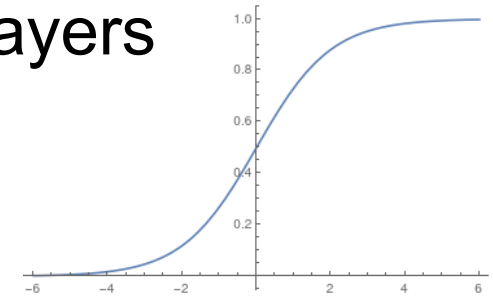
Non-linear activation functions

Typical activation functions for *hidden* layers

- Sigmoid
 - $\sigma(z) = (1 + \exp(-z))^{-1}$
 - It was popular for a long period of time
- Tanh:

$$\begin{aligned}\tanh(z) &= \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)} \\ &= 2\sigma(2z) - 1\end{aligned}$$

- Rectified linear unit (ReLU) and its smooth approximations
 - $\text{ReLU}(z) = \max(z, 0)$
 - e.g. GELU, see



[https://en.wikipedia.org/wiki/Rectifier_\(neural_networks\)](https://en.wikipedia.org/wiki/Rectifier_(neural_networks))

Non-linear activation functions

Typical activation functions for the *output* layer

- For classification
 - Use *softmax* (multi-class) or *sigmoid* (multi-label)
- For regression
 - Use a linear last layer
i.e. $\phi^{[L]}$ is the identity function in an L-layer NN

Terminology

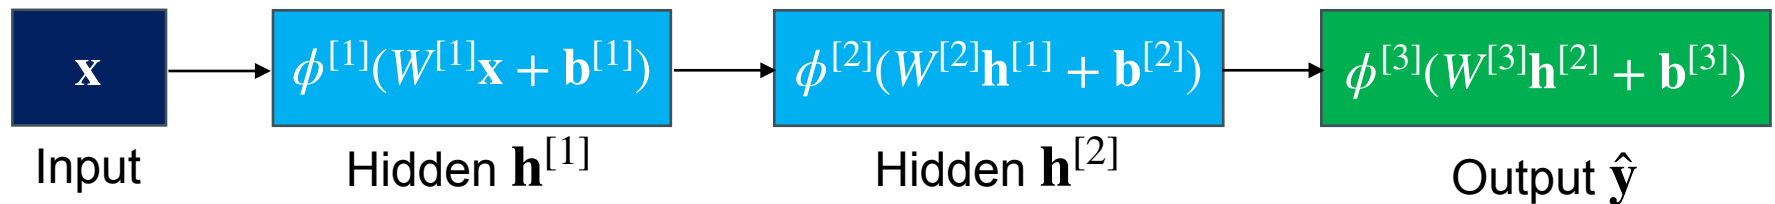
- Layers
 - For a Feedforward NN, we typically think of each composition of an activation function with a linear function as a *layer*
 - Layers except the input and output layers are *hidden layers*
 - In general, *almost anything*, regardless of complexity, can be called a layer, much like almost anything can be called a function.
 - The term is used to conceptually distinguish between different components

Terminology

- Hidden dimensions
 - The size of the *input* layer is defined by the number of input features
 - The size of the *output* layer is defined by the number of output targets
 - The size of the *hidden* layers can be anything, this is a hyper-parameter for you to choose
 - The dimension of the output of a *hidden* layer is called the *number of hidden units* in that layer

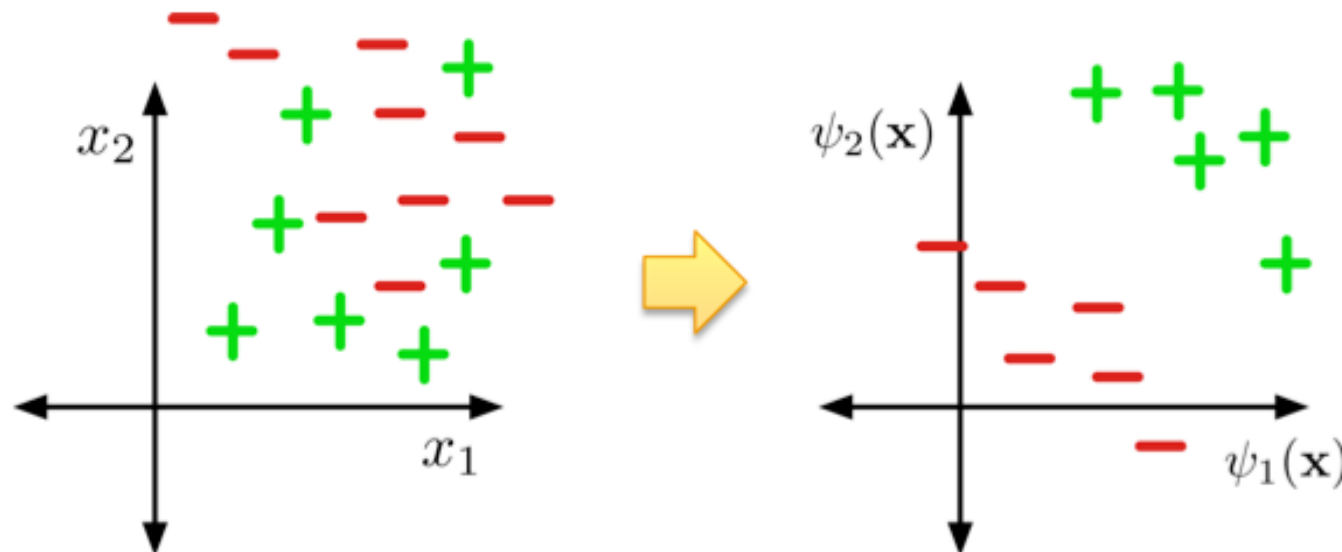
Terminology

- Feature / representation learning
 - The intermediate models learn to output a *useful representation* of the input, e.g. $\mathbf{h}^{[1]}$ and $\mathbf{h}^{[2]}$ below.
 - The final model still learns to predict the desired output
 - Neural nets can be viewed as a way of learning features



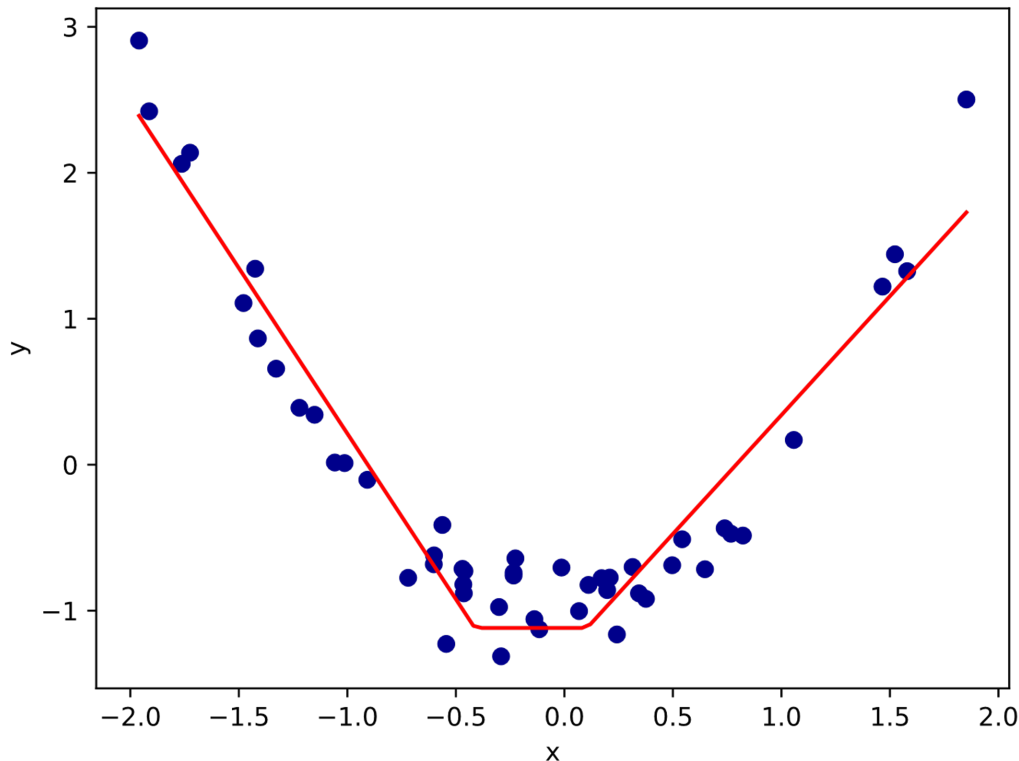
Terminology

- Feature / representation learning
 - The intermediate models learns to output a *useful representation* of the input, e.g. $\mathbf{h}^{[1]}$ and $\mathbf{h}^{[2]}$ below.
 - The final model still learns to predict the desired output
 - Neural nets can be viewed as a way of learning features



Example

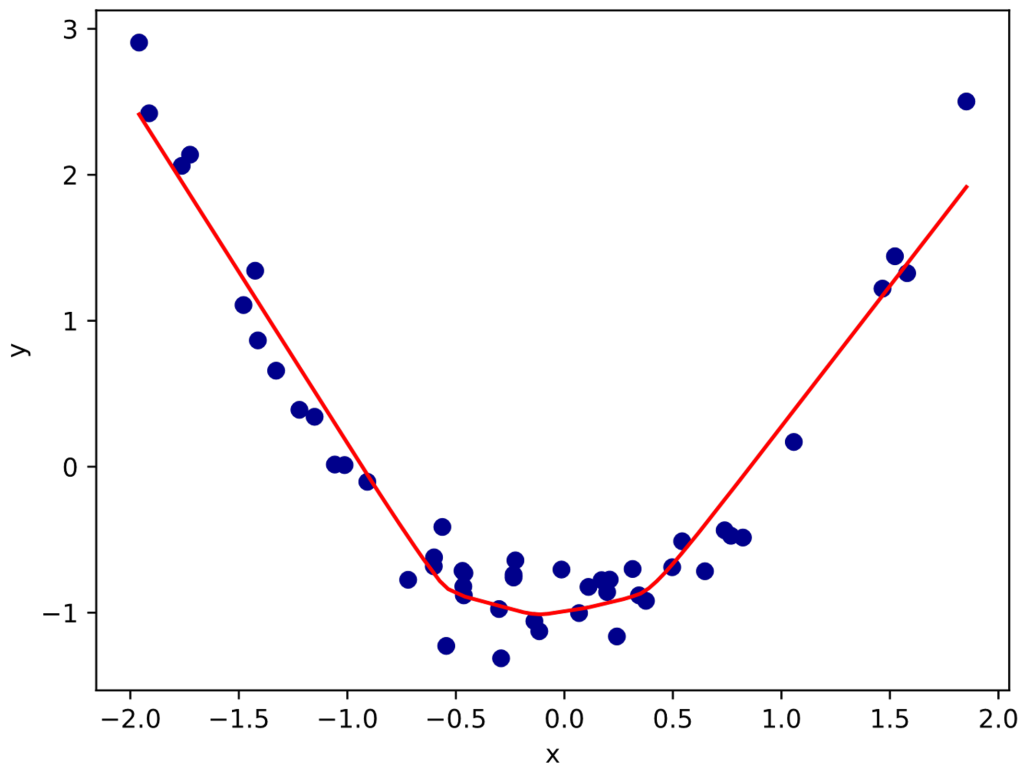
Fitting a Feedforward NN with 1 hidden layer of 2 hidden units with ReLU activations



- Each hidden unit learns a linear function then applies ReLU.
- Output is then a linear function of the hidden outputs.
- Resulting graph is equivalent to scaling, shifting and flipping two ReLU graphs and then adding them.

Example

Fitting a Feedforward NN with 1 hidden layer of 50 hidden units with ReLU activations



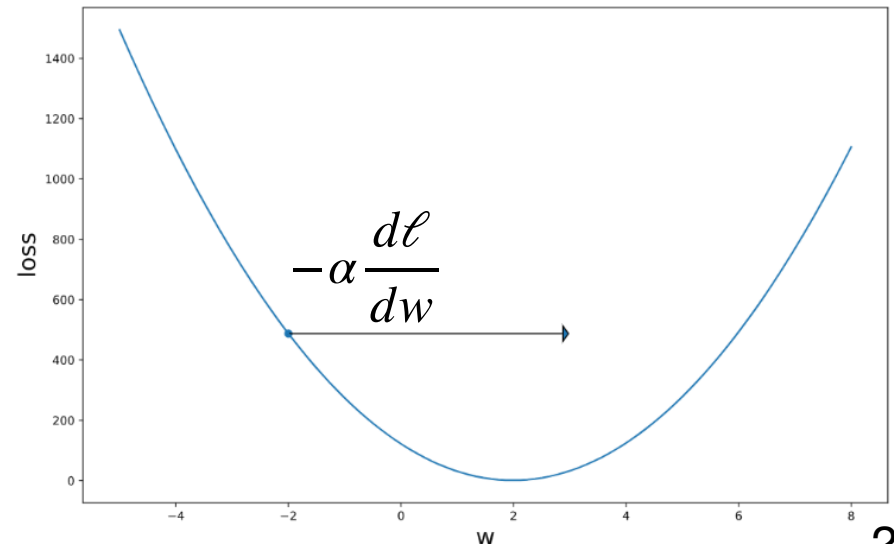
- Increasing the number of hidden units means we add together more (scaled) versions of the activation function.
- As we add more units we can approximate more complicated functions.
- As we add more units, our model becomes more capable of fitting noise in the training dataset — overfitting.

- Neural networks & why NN
- Feedforward neural network
 - From logistic regression to Feedforward NN
 - Non-linear activation functions
- **Back-propagation**
 - Gradient descent
 - Back-propagation essentials
 - Computation graph
- Optimisers
 - Stochastic gradient descent
 - SGD with momentum & Adam

Gradient descent

How can we learn the parameters $W^{[l]}$ and $\mathbf{b}^{[l]}$?

- Use gradient descent, and we need
 - Loss function $\ell(\mathbf{y}, \hat{\mathbf{y}})$
 - Partial derivatives of the loss function ℓ w.r.t. the parameters $W^{[l]}$ and $\mathbf{b}^{[l]}$
- To efficiently compute $\frac{\partial \ell}{\partial W^{[l]}}$ and $\frac{\partial \ell}{\partial \mathbf{b}^{[l]}}$
 - Use back-propagation



Back-propagation essentials

To efficiently compute the partial derivatives of the loss with respect to the parameters

- Forward pass
 - Compute the loss and the output of each layer
- Backward pass
 - Compute the partial derivatives layer by layer, starting from the last layer, using the *chain rule* and the results computed in the forward pass
 - A (bottom-up) *dynamic programming* (or memoisation) algorithm that avoids redundant calculations of intermediate terms

Chain rule

- Given $z = f(x(t))$, the derivative

$$\frac{dz}{dt} = \frac{dz}{dx} \frac{dx}{dt}$$

- Given $z = g(x(t), y(t))$, the derivative

$$\frac{dz}{dt} = \frac{\partial z}{\partial x} \frac{dx}{dt} + \frac{\partial z}{\partial y} \frac{dy}{dt}$$

Example: Forward pass

Given the model with loss function for data point (x, y) :

$$\ell(y, \hat{y}) = \frac{1}{2} (\phi(wx + b) - y)^2$$

We can introduce intermediate variables z, \hat{y} to give:

$$z = wx + b$$

$$\hat{y} = \phi(z)$$

$$\ell = \frac{1}{2}(\hat{y} - y)^2$$

Compute and store z, \hat{y}, ℓ in the forward pass.

Example: Backward pass

Compute the derivative of the loss ℓ w.r.t. \hat{y}

$$\frac{d\ell}{d\hat{y}} = \frac{d \frac{1}{2}(\hat{y} - y)^2}{d\hat{y}} = \hat{y} - y$$

Note that both z , \hat{y} , y are known:

- z , \hat{y} are computed and stored during the forward pass
- y is the ground truth

Forward pass:

$$z = wx + b$$

$$\hat{y} = \phi(z)$$

$$\ell = \frac{1}{2}(\hat{y} - y)^2$$

Back-propagation

Example: Backward pass

Compute the derivative of the loss ℓ w.r.t. \hat{y}

$$\frac{d\ell}{d\hat{y}} = \frac{d \frac{1}{2}(\hat{y} - y)^2}{d\hat{y}} = \hat{y} - y$$

Note that both z, \hat{y}, y are known:

- z, \hat{y} are computed and stored during the forward pass
- y is the ground truth

The chain rule states $\frac{d\ell}{dz} = \frac{d\ell}{d\hat{y}} \frac{d\hat{y}}{dz} = \frac{d\ell}{d\hat{y}} \phi'(z)$, and

$$\frac{d\ell}{dw} = \frac{d\ell}{dz} \frac{dz}{dw} = \frac{d\ell}{dz} x, \text{ and } \frac{d\ell}{db} = \frac{d\ell}{dz} \frac{dz}{db} = \frac{d\ell}{dz} \times 1$$

Forward pass:

$$z = wx + b$$

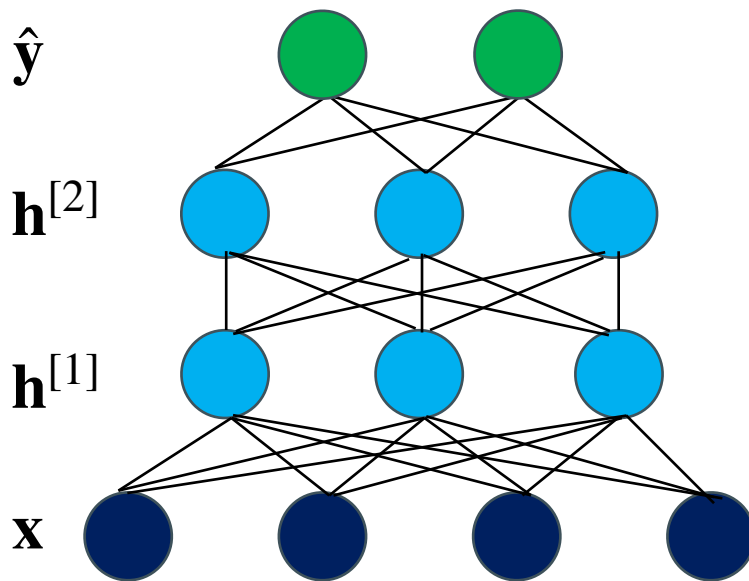
$$\hat{y} = \phi(z)$$

$$\ell = \frac{1}{2}(\hat{y} - y)^2$$

Back-propagation

Another example: Forward pass

Given data point (\mathbf{x}, \mathbf{y}) and loss function $\ell(\mathbf{y}, \hat{\mathbf{y}})$



$$\mathbf{a}^{[1]} = \mathbf{W}^{[1]}\mathbf{x} + \mathbf{b}^{[1]}$$

$$\mathbf{h}^{[1]} = \phi^{[1]}(\mathbf{a}^{[1]})$$

$$\mathbf{a}^{[2]} = \mathbf{W}^{[2]}\mathbf{h}^{[1]} + \mathbf{b}^{[2]}$$

$$\mathbf{h}^{[2]} = \phi^{[2]}(\mathbf{a}^{[2]})$$

$$\mathbf{a}^{[3]} = \mathbf{W}^{[3]}\mathbf{h}^{[2]} + \mathbf{b}^{[3]}$$

$$\hat{\mathbf{y}} = \phi^{[3]}(\mathbf{a}^{[3]})$$

$$\ell(\mathbf{y}, \hat{\mathbf{y}})$$

Back-propagation

Another example: Backward pass

Given data point (\mathbf{x}, \mathbf{y}) and loss function $\ell(\mathbf{y}, \hat{\mathbf{y}})$

Forward pass:

$$\mathbf{a}^{[1]} = \mathbf{W}^{[1]}\mathbf{x} + \mathbf{b}^{[1]}$$

$$\mathbf{h}^{[1]} = \phi^{[1]}(\mathbf{a}^{[1]})$$

$$\mathbf{a}^{[2]} = \mathbf{W}^{[2]}\mathbf{h}^{[1]} + \mathbf{b}^{[2]}$$

$$\mathbf{h}^{[2]} = \phi^{[2]}(\mathbf{a}^{[2]})$$

$$\mathbf{a}^{[3]} = \mathbf{W}^{[3]}\mathbf{h}^{[2]} + \mathbf{b}^{[3]}$$

$$\hat{\mathbf{y}} = \phi^{[3]}(\mathbf{a}^{[3]})$$

$$\ell(\mathbf{y}, \hat{\mathbf{y}})$$

Compute $\frac{\partial \ell}{\partial \mathbf{W}^{[1]}}$ using vector calculus¹

(more in this week's lab):

$$\frac{\partial \ell}{\partial \hat{\mathbf{y}}}, \quad \frac{\partial \ell}{\partial \mathbf{a}^{[3]}} = \frac{\partial \ell}{\partial \hat{\mathbf{y}}} \frac{\partial \hat{\mathbf{y}}}{\partial \mathbf{a}^{[3]}}$$

$$\frac{\partial \ell}{\partial \mathbf{h}^{[2]}} = \frac{\partial \ell}{\partial \mathbf{a}^{[3]}} \frac{\partial \mathbf{a}^{[3]}}{\partial \mathbf{h}^{[2]}}, \quad \frac{\partial \ell}{\partial \mathbf{a}^{[2]}} = \frac{\partial \ell}{\partial \mathbf{h}^{[2]}} \frac{\partial \mathbf{h}^{[2]}}{\partial \mathbf{a}^{[2]}}$$

$$\frac{\partial \ell}{\partial \mathbf{h}^{[1]}} = \frac{\partial \ell}{\partial \mathbf{a}^{[2]}} \frac{\partial \mathbf{a}^{[2]}}{\partial \mathbf{h}^{[1]}}, \quad \frac{\partial \ell}{\partial \mathbf{a}^{[1]}} = \frac{\partial \ell}{\partial \mathbf{h}^{[1]}} \frac{\partial \mathbf{h}^{[1]}}{\partial \mathbf{a}^{[1]}}$$

$$\frac{\partial \ell}{\partial \mathbf{W}^{[1]}} = \frac{\partial \ell}{\partial \mathbf{a}^{[1]}} \frac{\partial \mathbf{a}^{[1]}}{\partial \mathbf{W}^{[1]}}$$

¹ Using numerator layout, see https://en.wikipedia.org/wiki/Matrix_calculus

Computation graph

- Computing partial derivatives by hand is tedious and error-prone
- Luckily this can be *automated* by representing a network as a *computation graph* and providing the forward/backward procedures of operators involved
- We can diagram out the computations using a computation graph
 - Nodes represent all the inputs and computed quantities
 - Edges represent which nodes are computed directly as a function of which other nodes

Back-propagation

Computation graph

Example

$$\ell(y, \hat{y}) = \frac{1}{2} (\phi(wx + b) - y)^2$$

$$\begin{aligned} z &= wx + b \\ \hat{y} &= \phi(z) \\ \ell &= \frac{1}{2}(\hat{y} - y)^2 \end{aligned}$$

Back-propagation

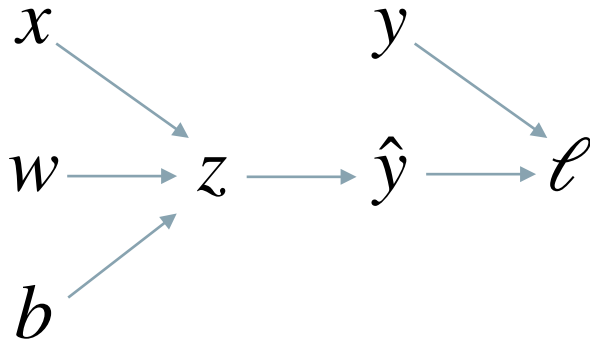
Computation graph

Example

$$\ell(y, \hat{y}) = \frac{1}{2} (\phi(wx + b) - y)^2$$

$$\begin{aligned} z &= wx + b \\ \hat{y} &= \phi(z) \\ \ell &= \frac{1}{2}(\hat{y} - y)^2 \end{aligned}$$

Forward pass



Back-propagation

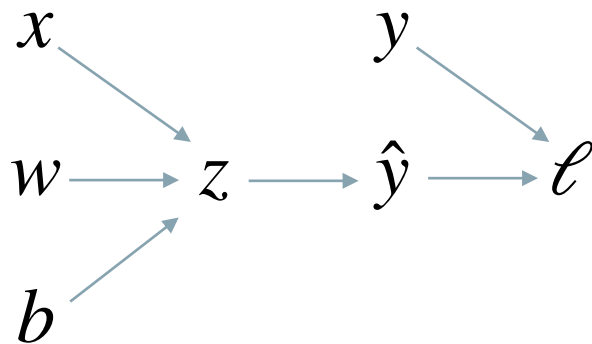
Computation graph

Example

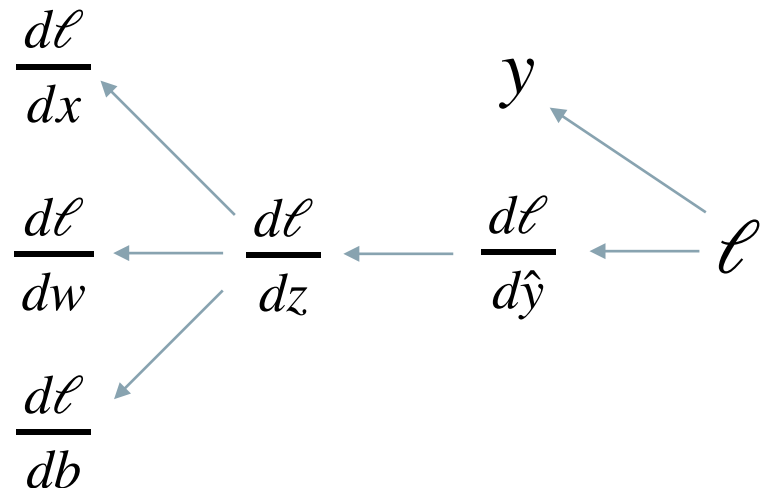
$$\ell(y, \hat{y}) = \frac{1}{2} (\phi(wx + b) - y)^2$$

$$\begin{aligned} z &= wx + b \\ \hat{y} &= \phi(z) \\ \ell &= \frac{1}{2}(\hat{y} - y)^2 \end{aligned}$$

Forward pass



Backward pass

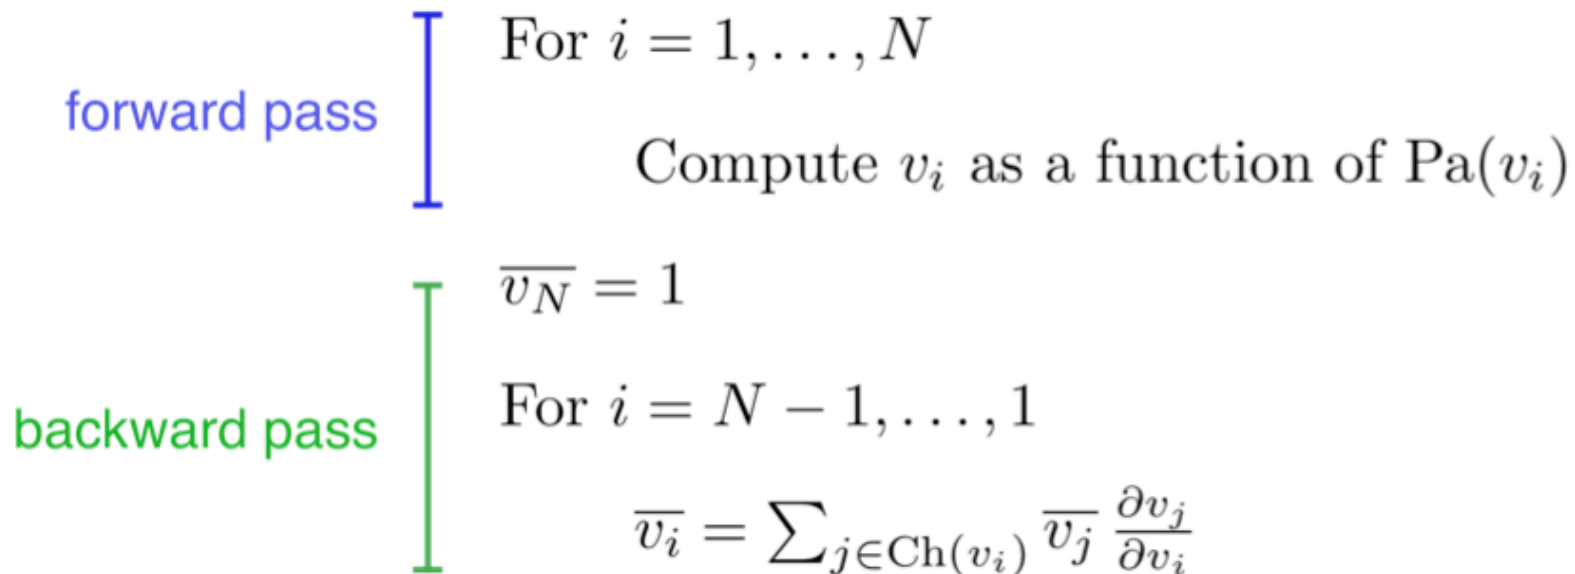


Note: updating x if learning embeddings

Back-propagation

Full back-propagation algorithm

Let v_1, \dots, v_N be a topological ordering of the computation graph (i.e. parents come before children), and v_N denotes the variable we're trying to compute derivatives (e.g. loss)



Back-propagation

Automatic differentiation with PyTorch

Calculus is hard! Use Autograd tools such as PyTorch (more in this week's lab), TensorFlow, JAX

```
import torch
x = torch.tensor([3., 2., 1.], requires_grad=False)
y = torch.tensor([0.7], requires_grad=False)

w = torch.tensor([1., 2., 3.], requires_grad=True)
b = torch.tensor([1.], requires_grad=True)
y_hat = torch.dot(w, x) + b

L = 1/2 * (y_hat - y)**2
L.backward()

print(w.grad)
print(b.grad)

tensor([30.9000, 20.6000, 10.3000])
tensor([10.3000])
```

PyTorch example

```
import torch
import torch.nn as nn
from sklearn.datasets import load_iris

# Get classification dataset
X_raw, y_raw = load_iris(return_X_y=True)
X = torch.tensor(X_raw, dtype=torch.float32)
y = torch.tensor(y_raw, dtype=torch.long)

loss_fn = nn.CrossEntropyLoss() # Setup loss function

# Define layers in our MLP
linear1 = nn.Linear(in_features=X.shape[1], out_features=10)
linear2 = nn.Linear(in_features=10, out_features=int(max(y_raw))+1)

# Define optimiser and give it layer parameters
all_params = list(linear1.parameters()) + list(linear2.parameters())
optimiser = torch.optim.SGD(all_params, lr=0.001)

# Run the MLP forwards
h = nn.functional.relu(linear1(X))
out_logits = linear2(h)

loss = loss_fn(out_logits, y) # Calculate loss

optimiser.zero_grad() # Zero out the stored gradients: important don't forget this
loss.backward()       # Compute gradients (Backpropagation)
optimiser.step()      # Do SGD step
```


- Neural networks & why NN
- Feedforward neural network
 - From logistic regression to Feedforward NN
 - Non-linear activation functions
- Back-propagation
 - Gradient descent
 - Back-propagation essentials
 - Computation graph
- **Optimisers**
 - Stochastic gradient descent
 - SGD with momentum & Adam

Stochastic gradient descent (SGD)

- Standard gradient descent
 - Gradients are computed on the loss of the entire training dataset
 - Slow if the dataset is very large
- Mini-batch SGD
 - At each update we randomly sample a batch of B data points and compute gradients only of the loss on these data points
 - The batch size B is a hyper-parameter
 - Usually B is one of 32, 64, ..., 512
 - SGD is much faster to compute per step
 - SGD acts as a regulariser: models trained with SGD usually generalise better

Popular optimisers

- Several improvements can be made to the basic stochastic gradient descent optimisation algorithm
- Two popular optimisers are *SGD with momentum* and *Adam*

SGD with momentum

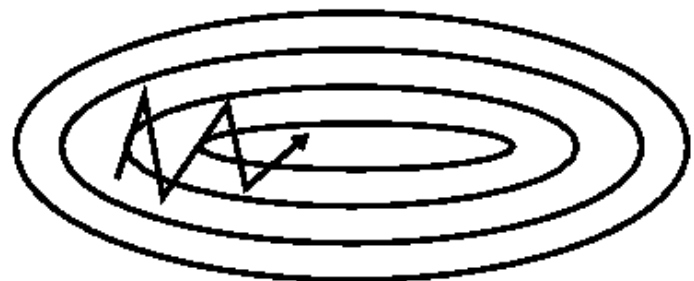
- Momentum is an exponentially decaying moving average of gradients from previous batches
- Momentum helps traverse gullies quicker, and can drastically reduce the number of steps needed for training.
- Update rule: $\boldsymbol{\nu}_t \leftarrow \alpha \boldsymbol{\nu}_{t-1} - \eta \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_t)$

where $\boldsymbol{\theta}_{t+1} \leftarrow \boldsymbol{\theta}_t + \boldsymbol{\nu}_t$

- $\boldsymbol{\nu}$ is the momentum, initialised to $\mathbf{0}$.
- $\alpha \in [0, 1]$ is the momentum coefficient, a hyper-parameter determining how quickly the previous gradients terms in $\boldsymbol{\nu}$ decay. Typically values: 0.5, 0.9, 0.99.
- η is the learning rate, $\boldsymbol{\theta}$ the parameters, and $J(\boldsymbol{\theta})$ the cost function.



SGD without momentum



SGD with momentum

Adam (Adaptive moment estimation)

- Adam keeps a running average of gradients and variances of gradients and uses them to normalise changes
- Adam is less sensitive to the learning rate than SGD with momentum
- In practice, Adam often converges in fewer iterations than SGD with momentum (Though this is not always the case).

Learning rate schedules

- Change the learning rate throughout training
- Exponentially decaying schedule multiplies learning rate by some $\lambda \in (0, 1)$ after every step
- Allows for larger learning rate to be used initially, while still converging.

Summary

- Neural networks & why NN
- Feedforward neural network
 - From logistic regression to Feedforward NN
 - Non-linear activation functions
- Back-propagation
 - Gradient descent
 - Back-propagation essentials
 - Computation graph
- Optimisers
 - Stochastic gradient descent
 - SGD with momentum & Adam

Reference

- Chapter 7, Speech and Language Processing