



Australian
National
University



COMP4650/6490 Document Analysis

Language Modelling & Smoothing

ANU School of Computing



Outline

- Language Models
- Smoothing
- Evaluation of Language Models



Outline

- Language Models
- Smoothing
- Evaluation of Language Models

Goal: assign a probability to a word sequence

- Speech recognition:
 - $P(\text{I ate a cherry}) > P(\text{Eye eight Jerry})$
- Spelling correction:
 - $P(\text{Australian National University}) > P(\text{Australian National Univerisity})$
- Collocation error correction:
 - $P(\text{high wind}) > P(\text{large wind})$
- Machine Translation:
 - $P(\text{The magic of Usain Bolt on show...}) > P(\text{The magic of Usain Bolt at the show ...})$
- Question-answering, summarisation, etc

- A language model *computes the probability of a sequence of words*
 - A vocabulary \mathcal{V}
 - $P(x_1, x_2, \dots, x_l) \geq 0$
 - $\sum_l \sum_{(x_1, x_2, \dots, x_l)} P(x_1, x_2, \dots, x_l) = 1$
- Related task: probability of an upcoming word
 - $P(x_4 \mid x_1, x_2, x_3)$
- Language model (LM):
Either $P(x_1, x_2, \dots, x_l)$ or $P(x_l \mid x_1, x_2, \dots, x_{l-1})$

How to Compute $P(x_1, x_2, \dots, x_l)$

- Apply the chain rule of probability

$$P(x_1, x_2, \dots, x_l) = P(x_1) P(x_2 | x_1) P(x_3 | x_1, x_2) \cdots P(x_l | x_1, x_2, \dots, x_{l-1})$$

- Example

$P(\text{John Smith's hotel room bugged})$

$= P(\text{John}) P(\text{Smith's} | \text{John}) P(\text{hotel} | \text{John Smith's})$

$\dots P(\text{bugged} | \text{John Smith's hotel room})$

Estimate the Probabilities

- Maximum likelihood estimation (MLE):

$P(\text{bugged} \mid \text{John Smith's hotel room}) =$

$\frac{\text{count}(\text{John Smith's hotel room bugged})}{\text{count}(\text{John Smith's hotel room})}$

- $P(\text{bugged} \mid \text{John Smith's hotel room})$ will most likely result in 0, as such specific long sequences are not likely to appear in the training set
- So, we need to simplify the calculations

Markov Assumption

- Simplification:

$$P(\text{bugged} \mid \text{John Smith's hotel room}) \approx P(\text{bugged} \mid \text{room})$$

OR

$$P(\text{bugged} \mid \text{John Smith's hotel room}) \approx P(\text{bugged} \mid \text{hotel room})$$

- First-order Markov assumption:

$$\begin{aligned} P(x_1, x_2, \dots, x_l) &= P(x_1) \prod_{i=2}^l P(x_i \mid x_{1:i-1}) \\ &= P(x_1) \prod_{i=2}^l P(x_i \mid x_{i-1}) \end{aligned}$$

Unigram Model

- Zero-order Markov assumption

$$P(x_1, x_2, \dots, x_l) = \prod_{i=1}^l P(x_i)$$

- Examples generated from a unigram model

*Months the my and issue of year foreign new exchange's
September were recession exchange new endorsed a acquire
to six executives*

- Not very good!

Bigram Model

- First-order Markov assumption

$$P(x_1, x_2, \dots, x_l) = P(x_1) \prod_{i=2}^l P(x_i | x_{i-1})$$

- $P(\textit{I want to eat Chinese food}) = ?$
- Estimate bigram probabilities from a training corpus

Bigram Counts from Berkeley Restaurant Corpus

x_i

x_{i-1}

| | i | want | to | eat | chinese | food | lunch | spend |
|----------------|----------|-------------|-----------|------------|----------------|-------------|--------------|--------------|
| i | 5 | 827 | 0 | 9 | 0 | 0 | 0 | 2 |
| want | 2 | 0 | 608 | 1 | 6 | 6 | 5 | 1 |
| to | 2 | 0 | 4 | 686 | 2 | 0 | 6 | 211 |
| eat | 0 | 0 | 2 | 0 | 16 | 2 | 42 | 0 |
| chinese | 1 | 0 | 0 | 0 | 0 | 82 | 1 | 0 |
| food | 15 | 0 | 15 | 0 | 1 | 4 | 0 | 0 |
| lunch | 2 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| spend | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 |

Table contains $count(x_{i-1}, x_i)$, only a subset of the full table of counts shown

Compute Bigram Probabilities

| | i | want | to | eat | chinese | food | lunch | spend |
|------------------|----------|-------------|-----------|------------|----------------|-------------|--------------|--------------|
| $count(x_{i-1})$ | 2533 | 927 | 2417 | 746 | 158 | 1093 | 341 | 278 |

- Maximum likelihood estimation:

$$P(x_i | x_{i-1}) = \frac{count(x_{i-1}, x_i)}{\sum_x count(x_{i-1}, x)} = \frac{count(x_{i-1}, x_i)}{count(x_{i-1})}$$

where special symbols $\langle s \rangle$ and $\langle /s \rangle$ are added at the beginning/end of a sentence in the corpus (for a bigram LM).

- Bigram probabilities, e.g. $P(\text{want} | \text{I}) = 827 / 2533 \approx 0.33$
- Log probabilities (logarithm helps with numerical stability)
 $\log P(\langle s \rangle \text{ I want to eat Chinese food } \langle /s \rangle)$
 $= \log P(\text{I} | \langle s \rangle) + \log P(\text{want} | \text{I}) + \log P(\text{to} | \text{want}) + \log P(\text{eat} | \text{to}) +$
 $\log P(\text{Chinese} | \text{eat}) + \log P(\text{food} | \text{Chinese}) + \log P(\langle /s \rangle | \text{food})$

Trigram Models

- Second order Markov assumption

$$P(x_1, x_2, \dots, x_l) = P(x_1)P(x_2 | x_1) \prod_{i=3}^l P(x_i | x_{i-2}, x_{i-1})$$

- Long-distance dependencies of language

“The iPhone which I bought one week ago does not stand the cold.”

- We can extend to 4-grams, 5-grams ...

Sequence Generation

- Given conditional probabilities:
 - $P(\text{want} \mid I) = 0.32$
 - $P(\text{to} \mid I) = 0$
 - $P(\text{eat} \mid I) = 0.004$
 - $P(\text{Chinese} \mid I) = 0$
 - $P(I \mid I) = 0.002$
 - ...
- Sampling:
 - Ensures you don't just get the same sentence all the time!
 - Can be done by generating a random number in $[0,1]$ based on a uniform distribution
 - Words that are a better fit (in the n-gram context) are more likely to be selected

Approximating Shakespeare

- Generate sentences from a **unigram** model:
 - Every enter now severally so, let
 - Hill he late speaks; or! a more to leg less first you enter
- From a **bigram** model:
 - What means, sir. I confess she? then all sorts, he is trim, captain
 - Why dost stand forth thy canopy, forsooth; he is this palpable hit the King Henry
- From a **trigram** model:
 - Sweet prince, Falstaff shall die
 - This shall forbid it should be branded, if renown made it empty

The Perils of Overfitting

- $P(I \text{ want to eat Chinese food}) = ?$
when $\text{count}(\text{Chinese food}) = 0$
- In practice, the test corpus is often different than the training corpus
- Unknown words
- Unseen n-grams



Outline

- Language Models
- **Smoothing**
- Evaluation of Language Models

Smoothing

- To prevent a language model from assigning zero probability to unseen events (e.g. unknown words, unseen contexts), i.e. not to overfit the training corpus
- Smoothing: adjusting low probabilities (such as zero probabilities) upwards, and high probabilities downwards
- Typically by adjusting the counts in MLE to produce more accurate probabilities
- Works very well, also used to improve NN approaches

Interpolation

- Key idea: mix lower order n-gram probabilities
- λ s are hyper-parameters

- For **bigram** models:

$$\hat{P}(x_i | x_{i-1}) = \lambda P(x_i | x_{i-1}) + (1 - \lambda)P(x_i), \quad \lambda \in [0,1]$$

- For **trigram** models:

$$\hat{P}(x_i | x_{i-2}, x_{i-1}) = \lambda_3 P(x_i | x_{i-2}, x_{i-1}) + \lambda_2 P(x_i | x_{i-1}) + \lambda_1 P(x_i)$$

$$\sum_{i=1}^3 \lambda_i = 1, \quad \lambda_i \in [0,1]$$

How to Set the λ s?

- Choose λ_i that maximise the probability of held-out data

Training data

Held-out
data

Test data

- Typically, Expectation Maximisation (EM) is used
- One crude approach (Collins et al. Course notes 2013)

$$\lambda_3 = \frac{\text{count}(x_{i-2}, x_{i-1})}{\text{count}(x_{i-2}, x_{i-1}) + \gamma}, \quad \lambda_2 = (1 - \lambda_3) \frac{\text{count}(x_{i-1})}{\text{count}(x_{i-1}) + \gamma},$$

$$\lambda_1 = 1 - \lambda_2 - \lambda_3$$

- Ensures λ_i is larger when count is larger
- Different λ s for each n-gram
- Only one parameter to estimate (i.e. γ)

Absolute Discounting

- Aims to deal with sequences that occur infrequently
- The discount d reserves some probability mass for the unseen n-grams

$$P_{\text{AbsDiscount}}(x_i | x_{i-1}) = \frac{\max(\text{count}(x_{i-1}, x_i) - d, 0)}{\text{count}(x_{i-1})} + \lambda(x_{i-1}) P(x_i)$$

$$\lambda(x_{i-1}) = \frac{d}{\text{count}(x_{i-1})} | \{x : \text{count}(x_{i-1}, x) > 0\} |$$

- How much to discount? Estimate using a held-out corpus. Typically, $d = 0.75$ is used.
- Is $P_{\text{AbsDiscount}}(x_i | x_{i-1})$ a true probability distribution?

Absolute Discounting

For all n-grams that had count c in the training corpus, calculating the average count \hat{c} of all those n-grams in the held-out corpus, and this can be used to estimate d

| Bigram count in training set | Bigram count in heldout set |
|------------------------------|-----------------------------|
| 0 | 0.0000270 |
| 1 | 0.448 |
| 2 | 1.25 |
| 3 | 2.24 |
| 4 | 3.23 |
| 5 | 4.21 |
| 6 | 5.23 |
| 7 | 6.21 |
| 8 | 7.21 |
| 9 | 8.26 |

AP newswire texts.
22 million words for training
22 million words for held-out

We tend to systematically overestimate n-gram counts

Kneser-Ney Smoothing

- Augments absolute discounting
- Key idea (assumption):
 - If a **word** appears after a small number of contexts, it should be less likely to appear in a **novel context**
 - If there are only a few words that come after a **context**, then a **novel word** in that context should be less likely
- Consider the word: *Francisco*
 - Might be common because: *San Francisco* is a common term
 - But *Francisco* occurs in very few contexts

Kneser-Ney Smoothing

- For task: I can't see without my reading _____?
 - The word *glasses* seems more likely here than *Francisco*
 - However, a standard unigram model will assign *Francisco* a higher probability than *glasses* because *Francisco* is more common (since *San Francisco* is a very frequent term)
 - We want our unigram model to prefer *glasses*, because *Francisco* is frequent mainly because *San Francisco* is frequent
- $P_{\text{continuation}}(x_i)$ instead of $P(x_i)$
 - Words that have appeared in more contexts are more likely to appear in new context, e.g. *glasses*

Kneser-Ney Smoothing

- $P_{\text{continuation}}(x_i)$
 - How likely is word x_i to continue a new context?
 - Proportional to the number of different words it follows:
$$P_{\text{continuation}}(x_i) \propto |\{x : \text{count}(x, x_i) > 0\}|$$
 - And
$$|\{x_{i-1} : \text{count}(x_{i-1}, \text{Francisco}) > 0\}| \ll |\{x_{i-1} : \text{count}(x_{i-1}, \text{glasses}) > 0\}|$$
- Normalising
$$P_{\text{continuation}}(x_i) = \frac{|\{x : \text{count}(x, x_i) > 0\}|}{\sum_{x'} |\{x : \text{count}(x, x') > 0\}|}$$

Kneser-Ney Smoothing

- Absolute discounting

$$P_{\text{AbsDiscount}}(x_i | x_{i-1}) = \frac{\max(\text{count}(x_{i-1}, x_i) - d, 0)}{\text{count}(x_{i-1})} + \lambda(x_{i-1}) P(x_i)$$

- Kneser-Ney smoothing

$$P_{\text{KN}}(x_i | x_{i-1}) = \frac{\max(\text{count}(x_{i-1}, x_i) - d, 0)}{\text{count}(x_{i-1})} + \lambda(x_{i-1}) P_{\text{continuation}}(x_i)$$

$\lambda(x_{i-1})$ is defined the same as in absolute discounting

- Is $P_{\text{KN}}(x_i | x_{i-1})$ a true probability distribution?

Smoothing for Web-scale N-grams

- *Stupid Backoff* (Brants *et al.* 2007): No discounting, simply backoff to a lower order n-gram (with fixed weight λ) if a higher order n-gram has 0 count
- Does not give a probability distribution

$$S(x_i | x_{i-n+1:i-1}) = \begin{cases} \frac{\text{count}(x_{i-n+1:i})}{\text{count}(x_{i-n+1:i-1})} & \text{if } \text{count}(w_{i-n+1:i}) > 0 \\ \lambda S(x_i | x_{i-n+2:i-1}) & \text{otherwise} \end{cases}$$

$$S(x) = \frac{\text{count}(x)}{N}$$

- N is the size of the training corpus in words
- $\lambda = 0.4$ worked well

Smoothing for Web-scale N-grams

- *Stupid Backoff*: A simplistic type of smoothing
- Inexpensive to train on large data sets, and approaches the quality of Kneser-Ney smoothing as the amount of training data increases
- Try to use higher-order n-gram's, otherwise drop to shorter sequence, hence backoff!
 - Every time you 'backoff', e.g. reduce from trigram to bigram because you've encountered a 0 probability, you multiply by 0.4
 - i.e. use trigram if good evidence available, otherwise bigram, otherwise unigram!
- S = scores, not probabilities!



Outline

- Language Models
- Smoothing
- **Evaluation of Language Models**

- Extrinsic evaluation:
 - Put each model in a task, e.g. spelling correction, machine translation etc.
 - Usually time consuming
- Intrinsic evaluation:
 - Measures the quality of a language model independent of any application
 - Needs a held-out dataset $x_{1:L}$
 - Perplexity

$$PP(x_{1:L}) = P(x_{1:L})^{-\frac{1}{L}}$$
$$= \sqrt[L]{\frac{1}{P(x_{1:L})}}$$

Perplexity

- The inverse probability of held-out data, normalised by the number of words
- The lower the better
- Works well when test and training data are similar
- Pre-processing and vocabulary matter (only comparable if two LMs use identical vocabulary)
- Improvement in perplexity does not guarantee increased performance of an NLP task

| | Unigrams | Bigram | Trigram |
|------------|----------|--------|---------|
| Perplexity | 962 | 170 | 109 |

Perplexity of a 1.5M-word WSJ test set

- SRILM
 - <http://www.speech.sri.com/projects/srilm/>
- Berkeley LM
 - <https://code.google.com/archive/p/berkeleylm/>
- KenLM
 - <https://kheafield.com/code/kenlm/>
- Available Language Models (and training recipe)
 - <http://www.keithv.com/software/csr/>

References

- Chapter 3. Speech and Language Processing (3rd ed. draft)
- Brants et al. Large language models in machine translation. EMNLP/CoNLL. 2007.
- Course notes for NLP. Michael Collins.
<http://www.cs.columbia.edu/~mcollins/lm-spring2013.pdf>
- NLP Lunch Tutorial: Smoothing. Bill MacCartney.
<https://nlp.stanford.edu/~wcmac/papers/20050421-smoothing-tutorial.pdf>