

Lab 2

SQL Basics

In this lab, we will begin to use the database management system (DBMS) PostgreSQL. The documentation for PostgreSQL, including a great tutorial, can be found at <https://www.postgresql.org/docs/14/index.html>.

1 Command-Line Interface `psql`

The command-line interface for PostgreSQL works similarly to the UNIX shell, and is entered from the shell by typing the command `psql`. **Students may mix up commands for the shell and DBMS, try to always remember which one you are in. The prompts are different:**

- The shell on your lab computer shows you

```
u1234567@n11X1tYZ:$
```

- The shell on the partch server shows you

```
u1234567@partch:$
```

- The `psql` on partch or on the local computer shows you

```
u1234567=>
```

Note: In case you see `u1234567->` instead of `u1234567=>`, this indicates an unfinished SQL command. Type a semicolon `;` to finish the command and press Enter to execute.

(1) *Start the PostgreSQL command-line interface by entering `psql` on the partch server.*

Open a command shell and log in to `partch` by typing `ssh u123456@partch.anu.edu.au` at the command line in the terminal window

```
u1234567@n11X1tYZ:$ ssh u1234567@partch.anu.edu.au
```

Start the PostgreSQL interface by entering `psql` in your terminal.

```
u1234567@partch:$ psql
```

Please note that `psql` has also been installed locally on each lab computer, but you are recommended to use the `psql` installed on the `partch` server. The information (e.g., tables) stored through the local `psql` on each lab computer is not linked to your personal account and you can not retrieve it next time you login. The information stored through the `psql` on the `partch` server is linked to your personal account and you can retrieve that next time you login to the `partch` server (from the same or different lab computer, or access the `partch` server remotely from your personal computer).

There are two kinds of commands that the PostgreSQL command-line interface can handle: SQL commands and meta-commands. The meta-commands always begin with `\`. Below are some of the most useful meta-commands in PostgreSQL.

<code>\?</code>	Help on psql commands
<code>\q</code>	Quit psql and return to the Unix shell
<code>\dt</code>	List the currently defined tables
<code>\d <table-name></code>	Describe a table, listing its columns and datatypes
<code>\i <file-name></code>	Input commands from a file

(2) *Type `\dt` to list the tables in the current database.*

At the moment, your database has no table.

(3) *Quit out of `psql` using `\q` and start it up again.*

You can also quit by pressing CTRL-d (hold the CTRL key down and press 'd' at the same time).

2 Data Definition Language

(4) *Enter the following SQL command, which will create a new table in your database.*

```
CREATE TABLE student (  
    name varchar(20),  
    email varchar(50)  
);
```

Note: Don't forget the semicolon ';' at the end of each SQL command. In case that happens, you will see that the command prompt is changed from

“=>” to “->” to indicate an unfinished SQL command. Type a semicolon to finish the command and press Enter to execute.

(5) *Now, add one record to the new table using the following command.*

```
INSERT INTO student (name, email) VALUES
    ('Peter', 'peter@gmail.com');
```

Rather than typing everything through the command-line interface, it is often more convenient to prepare code in a file and run it from that.

(6) *Open a text editor on your lab computer, and save the following code into a file called `insertStudents.sql` and copy this file from your own computer to your directory for this week’s lab on partch.*

```
INSERT INTO student (name, email) VALUES
    ('Aiden', 'aiden@hotmail.com');
INSERT INTO student (name, email) VALUES
    ('Emilia', 'emilia@yahoo.com');
INSERT INTO student (name, email) VALUES
    ('Ian', 'ian@github.com');
```

To run this .sql code, type `\i insertStudents.sql` within `psql`. If you get an error like

```
insertStudents.sql: No such file or directory
```

then you are probably not in the directory where you saved the file. Exit `postgreSQL` using `\q`, then do a `pwd` to see where you are, and use `ls` to see whether the file `insertStudents.sql` is there. Find your way to the correct directory, or save the file where you should have saved it and try again.

You may also use the following query command to show all records in the table `student`.

(7) *Run this query.*

```
SELECT * FROM student;
```

(8) *Download the following two files from the Wattle course site (in the folder Lab 2: SQL Basics), and copy the files to the directory for this week’s lab on partch.*

- `employeeCreate.sql`,
- `employeeQueries.sql`

(9) *Open `employeeCreate.sql` in an editor, and look through the code.*

Notice some language features such as `NOT NULL`.

(10) *Go and look at the PostgreSQL manual in your browser (<https://www.postgresql.org/docs/14/index.html>)*

Section II, “The SQL Language” is a great reference that you can use to find neat ways of doing things, and learn the meaning of pieces of code you do not understand. The `CREATE TABLE` command is covered in Section II.5, “Data Definition”. Click on the link to Section II.5.4.2, “Not-Null Constraints”, and read a little of the entry.

(11) *Use the PostgreSQL manual to find out the meaning of some other not-so-obvious constraints in `employeeCreate.sql`, such as: check constraints, primary keys and foreign keys.*

The file `employeeCreate.sql`, as you have seen, is written to create the database tables and add records into them. There are several `DROP TABLE` commands included at the beginning of `employeeCreate.sql`, which are used to remove the existing tables (if any) in our example database before creating the new tables.

(12) *Run `employeeCreate.sql` now by typing*

```
\i employeeCreate.sql
```

You should see a lot of messages saying `CREATE TABLE` and `INSERT 0 1`. Note that `INSERT 0 1` indicates that exactly one row is added to the table.

Now we have a database to play with.

PostgreSQL provides the `ALTER TABLE` command to change the definition of a table.

(13) *Consider the table `project` in our example database, how can you add an attribute `StartDate` to `project`?*

```
ALTER TABLE project ADD StartDate Date;
```

(14) *Consider the table `department` in the example database, how can you ensure that the values of the attribute `mgrssn` are valid, i.e., each of them must be a `ssn` in the table `employee`? (Hint: add a foreign key constraint on `department`)*

```
ALTER TABLE department ADD FOREIGN KEY (mgrssn) REFERENCES employee(ssn);
```

3 Simple Queries

The above example database state (created by `employeeCreate.sql`) is shown in the following figure. Here you may try to run some simple queries and please don't worry if you don't understand them right now.

EMPLOYEE	fname	minit	lname	ssn	bdate	address	salary	superssn	dno
	Michio		Morishima	20118	1973-07-18	79 Macpherson St, Turner	52107.00	21286	1000
	John		Backus	20766	1984-12-03	25 Burns St, Yarralumla	46789.00	21287	1007
	Gramsci		Antonio	20876	1991-01-22	27 Garibaldi St, Ashfield, NSW	71569.00	20915	1001
	Ada		Lovelace	21286	1985-12-10	17 Ainslie Ave, Reid, ACT	62107.00	21286	1000
	Milton		Friedman	29057	1972-07-31	75 Wakefield Ave, Ainslie	37764.00	21287	1007
	Edsger	W	Dijkstra	20765	1980-05-11	192 Wattle St, O'Connor ACT	73567.00	20766	1000
	Grace	M	Hopper	20864	1976-12-09	45 Cobol St, Parramatta, NSW	78563.00	21286	1000
	Frederick	W	Taylor	20915	1986-03-20	14 Blackett St, Downer, ACT	56098.00	20915	1001
	John	M	Keynes	21287	1983-06-05	94 Earle St, Lyneham, ACT	73567.00	21287	1007
(9 rows)									

DEPARTMENT	dname	dnumber	mgrssn	mgrstartdate
	Information Technology	1000	20765	2007-01-01
	Administration	1001	20915	2004-02-29
	Finance	1007	21287	2005-06-07
(3 rows)				

DEPT_LOCATION	dnumber	dlocation
	1000	Canberra
	1000	Sydney
	1001	Canberra
	1007	Canberra
	1007	Sydney
(5 rows)		

PROJECT	pname	pnumber	plocation	dnum
	Difference Engine	9000	Canberra	1000
	Red tape is Fun	9001	Canberra	1001
	Object Oriented COBOL	9002	Sydney	1000
(3 rows)				

WORKS_ON	ssn	pno	hours
	20765	9000	100
	20765	9001	500
	20864	9002	50
	20915	9000	250
(4 rows)			

(15) Open `employeeQueries.sql`, and take a look at the queries there.

(16) For each query: work out what result you would expect it to yield against our example database. Write and save this result in a comment-block below the query. Then run it in `psql` by copying and pasting into the terminal. Compare actual and expected results, and get an idea of simple SQL queries.

(17) Refer to the online resources to learn PostgreSQL <https://www.postgresql.org/docs/online-resources/>