



## SQL – Part 4

### Data Manipulation Language (Advanced SQL Queries)



## Advanced SQL Queries – Set Operations

- SQL incorporates several set operations: **UNION** (set union) and **INTERSECT** (set intersection), and sometimes **EXCEPT** (set difference / minus).
- Set operations result in return of a relation of tuples (no duplicates).
- Set operations apply to relations that have the same attribute types appearing in the same order, e.g., list all students who have either a gmail or hotmail email account.

```
(SELECT * FROM STUDENT WHERE Email like '%@gmail.com')  
UNION  
(SELECT * FROM STUDENT WHERE Email like '%@hotmail.com');
```

- For example, the following query will not work

```
(SELECT StudentID, Name FROM STUDENT)  
UNION  
(SELECT Email FROM STUDENT);
```

## Advanced SQL Queries – Join Operations

- When we want to retrieve data from *more than one relations*, we often need to use **join** operations.
- Consider the following queries, which both need a join operation between two relations:
  - List the names of all courses which have been enrolled by at least one student.
  - List all students, and their enrolled courses if any.

STUDENT			
StudentID	Name	DoB	Email

COURSE		
No	Cname	Unit

ENROL				
StudentID	CourseNo	Semester	Status	EnrolDate



## Advanced SQL Queries – Inner Join

- **Inner Join**: tuples are included in the result only if there is at least one matching in both relations.
- For the query “list the names of all courses which have been enrolled by at least one student”, we use:

```
SELECT DISTINCT c.Cname  
FROM COURSE c INNER JOIN ENROL e ON c.No=e.CourseNo;
```

COURSE		
No	Cname	Unit
COMP2400	Relational Databases	6
COMP3900	Advanced Database Concepts	6

ENROL				
StudentID	CourseNo	Semester	Status	EnrolDate
456	COMP1130	2016 S1	active	25/02/2016
458	COMP1130	2016 S1	active	25/02/2016
456	COMP2400	2016 S2	active	09/03/2016


- Result:

Cname
Relational Databases


## Advanced SQL Queries – Outer Join

- **Outer Join** includes **Left Join** and **Right Join**.
- **Left/Right Join**: all tuples of the left/right table are included in the result, even if there are no matches in the relations.


Inner Join



Left Join



Right Join





## Advanced SQL Queries – Outer Join

- **Left Join:** A left join retains all rows of the left table regardless of whether there is a row that matches on the right table.

ENROL1		
<u>StudentID</u>	<u>CourseNo</u>	<u>Semester</u>
456	COMP1130	2016 S1
457	COMP1130	2016 S1
456	COMP2400	2016 S2

STUDENT			
<u>StudentID</u>	Name	DoB	Email
456	Tom	25/01/1988	tom@gmail.com
458	Peter	20/02/1991	peter@hotmail.com

```
SELECT *  
FROM STUDENT s LEFT JOIN ENROL1 e  
ON s.StudentID=e.StudentID;
```

StudentID	Name	DoB	Email	StudentID	CourseNo	Semester
456	Tom	25/01/1988	tom@gmail.com	456	COMP1130	2016 S1
456	Tom	25/01/1988	tom@gmail.com	456	COMP2400	2016 S2
458	Peter	20/02/1991	peter@hotmail.com	null	null	null



## Advanced SQL Queries – Outer Join

- **Right Join:** A right join retains all rows of the right table regardless of whether there is a row that matches on the left table.

ENROL1		
<u>StudentID</u>	<u>CourseNo</u>	<u>Semester</u>
456	COMP1130	2016 S1
457	COMP1130	2016 S1
456	COMP2400	2016 S2

STUDENT			
<u>StudentID</u>	Name	DoB	Email
456	Tom	25/01/1988	tom@gmail.com
458	Peter	20/02/1991	peter@hotmail.com

```
SELECT *  
FROM STUDENT s RIGHT JOIN ENROL1 e  
ON s.StudentID=e.StudentID;
```

StudentID	Name	DoB	Email	StudentID	CourseNo	Semester
456	Tom	25/01/1988	tom@gmail.com	456	COMP1130	2016 S1
null	null	null	null	457	COMP1130	2016 S1
456	Tom	25/01/1988	tom@gmail.com	456	COMP2400	2016 S2



## Advanced SQL Queries – Outer Join

- For the query “list all students, and their enrolled courses if any”, we can use either of the following statements:

```
SELECT s.*, e.CourseNo, e.Semester  
FROM STUDENT s LEFT JOIN ENROL1 e  
ON s.StudentID=e.StudentID;
```

```
SELECT s.*, e.CourseNo, e.Semester  
FROM ENROL1 e RIGHT JOIN STUDENT s  
ON e.StudentID=s.StudentID;
```

- If we have 1000 tuples in STUDENT, then the query result should contain at least 1000 tuples (one tuple in STUDENT may occur multiple times) with the following attributes:

StudentID	Name	DoB	Email	CourseNo	Semester
...	...	...	...	...	...



## Advanced SQL Queries – Natural Join

- **Motivation:** An inner join retains all the data of the two tables for , with duplication

- ```
SELECT *  
FROM STUDENT s INNER JOIN ENROL1 e  
On s.StudentID=e.StudentID;
```

| ENROL1           |                 |                 |
|------------------|-----------------|-----------------|
| <u>StudentID</u> | <u>CourseNo</u> | <u>Semester</u> |
| 456              | COMP1130        | 2016 S1         |
| 457              | COMP1130        | 2016 S1         |
| 456              | COMP2400        | 2016 S2         |

| STUDENT          |       |            |                   |
|------------------|-------|------------|-------------------|
| <u>StudentID</u> | Name  | DoB        | Email             |
| 456              | Tom   | 25/01/1988 | tom@gmail.com     |
| 458              | Peter | 20/02/1991 | peter@hotmail.com |

- **Result:**

| StudentID | Name | DoB        | Email         | StudentID | CourseNo | Semester |
|-----------|------|------------|---------------|-----------|----------|----------|
| 456       | Tom  | 25/01/1988 | tom@gmail.com | 456       | COMP1130 | 2016 S1  |
| 456       | Tom  | 25/01/1988 | tom@gmail.com | 456       | COMP2400 | 2016 S2  |

## Advanced SQL Queries – Natural Join

- **Natural Join:** A natural join retains all the data of the two tables for only the matched rows, without duplication

- ```
SELECT *  
FROM STUDENT s NATURAL JOIN ENROL1 e;
```

ENROL1		
<u>StudentID</u>	<u>CourseNo</u>	<u>Semester</u>
456	COMP1130	2016 S1
457	COMP1130	2016 S1
456	COMP2400	2016 S2

STUDENT			
<u>StudentID</u>	<u>Name</u>	<u>DoB</u>	<u>Email</u>
456	Tom	25/01/1988	tom@gmail.com
458	Peter	20/02/1991	peter@hotmail.com

- **Result:**

<u>StudentID</u>	<u>Name</u>	<u>DoB</u>	<u>Email</u>	<u>CourseNo</u>	<u>Semester</u>
456	Tom	25/01/1988	tom@gmail.com	COMP1130	2016 S1
456	Tom	25/01/1988	tom@gmail.com	COMP2400	2016 S2



## Advanced SQL Queries – Natural Join

- **Natural Join:** One kind of inner join, in which two relations are joined implicitly by comparing all attributes of the same names in both relations.
- For the query “list all students who have enrolled and their courses”, use:

```
SELECT * FROM STUDENT NATURAL JOIN ENROL;
```

ENROL				
<u>StudentID</u>	CourseNo	Semester	Status	EnrolDate
456	COMP1130	2016 S1	active	25/02/2016
457	COMP1130	2016 S1	active	25/02/2016

STUDENT			
<u>StudentID</u>	Name	DoB	Email
456	Tom	25/01/1988	tom@gmail.com
458	Peter	20/02/1991	peter@hotmail.com

- Result: (STUDENT.StudentID=ENROL.StudentID is used in the query)

StudentID	Name	DoB	Email	CourseNo	Semester	Status	EnrolDate
456	Tom	25/01/1988	tom@gmail.com	COMP1130	2016 S1	active	25/02/2016



## Advanced SQL Queries – Subqueries

- **Subqueries** are just queries that are used where a relation is required.
- Subqueries can be specified within the FROM-clause (usually in conjunction with aliases and renaming) to create *inline view* (exist only for the query)
- Subqueries can also be specified within the WHERE-clause, e.g.,
  - **IN** *subquery* tests if tuple occurs in the result of the subquery
  - **EXISTS** *subquery* tests whether the subquery results in non-empty relation
  - using **ALL**, **SOME** or **ANY** before a subquery makes subqueries usable in comparison formulae
  - in all these cases the condition involving the subquery can be negated using a preceding **NOT**



## Subqueries – In

- Recall that, for the query “list all students who have enrolled and their courses”, we have:

```
SELECT *  
FROM STUDENT NATURAL JOIN ENROL;
```

- Now if we want to query: “list all students who have enrolled in a course *that has less than 10 students enrolled* and the CourseNo of these courses”, we have

```
SELECT s.*,e1.CourseNo  
FROM STUDENT s NATURAL JOIN ENROL e1  
WHERE e1.CourseNo IN  
      (SELECT e2.CourseNo  
       FROM ENROL e2  
       GROUP BY e2.CourseNo  
       HAVING COUNT(*)<10);
```



## Subqueries – Exists

- For the query: “list all students who have enrolled in at least one course”, we have

```
SELECT s.*
FROM STUDENT s
WHERE EXISTS (SELECT *
              FROM ENROL e
              WHERE s.StudentID=e.StudentID);
```

- For the query: “list all students who have *not* enrolled in any course”, we have

```
SELECT s.*
FROM STUDENT s
WHERE NOT EXISTS (SELECT *
                  FROM ENROL e
                  WHERE s.StudentID=e.StudentID);
```



## Subqueries – More Complicated

- For the query: “list the courses that have the largest number of students enrolled in Semester 2 2016”, we have

```
SELECT e.CourseNo
FROM (SELECT e1.CourseNo, COUNT(*) AS NoOfStudents
      FROM ENROL e1
      WHERE e1.Semester = '2016 S2'
      GROUP BY e1.CourseNo) e
WHERE e.NoOfStudents =
      (SELECT MAX(e2.NoOfStudents)
       FROM (SELECT e1.CourseNo, COUNT(*) AS NoOfStudents
             FROM ENROL e1
             WHERE e1.Semester = '2016 S2'
             GROUP BY e1.CourseNo) e2);
```



## Subqueries – More Complicated

- For the query: “list all the courses that have more students enrolled than at least one other course in Semester 2 2016”, we have

```
SELECT e.CourseNo
FROM (SELECT e1.CourseNo, COUNT(*) AS NoOfStudents
      FROM ENROL e1
      WHERE e1.Semester = '2016 S2'
      GROUP BY e1.CourseNo) e
WHERE e.NoOfStudents > ANY
      (SELECT e2.NoOfStudents
       FROM (SELECT e1.CourseNo, COUNT(*) AS NoOfStudents
             FROM ENROL e1
             WHERE e1.Semester = '2016 S2'
             GROUP BY e1.CourseNo) e2);
```





## Views in SQL

- A view in SQL is a virtual table that is derived from other tables in the same database or previously defined views.
- How to Create Views?
  - Suppose we already have tables STUDENT(StudentID, Name, DoB, Email) and ENROL(StudentID, CourseNo, Semester, Status, EnrolDate). Then we can create a view ENROL1 as follows:

```
CREATE VIEW ENROL1
AS SELECT s.StudentID, s.Name, e.CourseNo, e.EnrolDate
FROM STUDENT s, ENROL e
WHERE s.StudentID=e.StudentID;
```