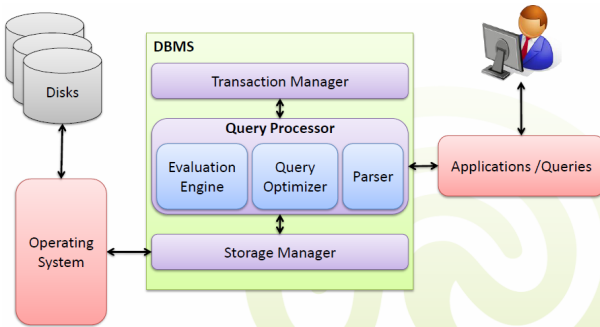




Query Processing

Query Processing – Overview

- 1 Users **submit** SQL queries to a DBMS.
- 2 The DBMS **processes and executes** them in a database.



- **Note:** SQL is a declarative language, so it is the task of DBMSs to decide how SQL queries should be executed.



Query Processing – Example

- **From:**

```
SELECT name FROM Person WHERE age<21;
```

- **To:**

name
Rickon Bran

- **Questions:**

- How does a relational DBMS process this?
- How can a relational DBMS process this efficiently?




Query Processing – Example

```
SELECT name FROM Person WHERE age<21;
```

High-level language
(SQL)

\Downarrow
 $\pi_{name}(\sigma_{age<21}(\text{Person}))$

Low-level language
(Relational Algebra)



Execution plan
(Query tree)

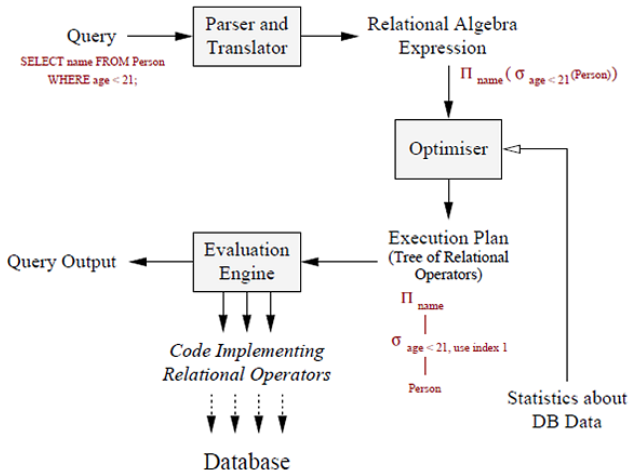
\Downarrow

name
Rickon
Bran

Query result



Query Processing – Example



Query Processing Steps

- **Query parser and translator**

- 1 Check the syntax of SQL queries
- 2 Verify that the relations do exist
- 3 Transform into relational algebra expressions

- **Query optimiser**

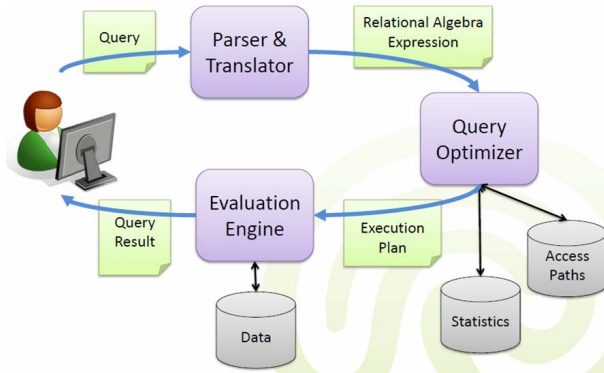
- 1 Transform into the best possible execution plan
- 2 Specify the implementation of each operator in the execution plan

- **Evaluation engine**

- 1 Evaluate the query execution plan
- 2 Return the result to the user

Query Processing – Parser

- The **parser** checks the syntax of the query:
 - Validation of table names, attributes, data types, access permission ...;
 - Either the query is executable or an error message is generated.





Query Processing – Parser

- Consider the relation schema:

Person(id:integer, name:string, age:integer, address:string)

- Note:** **System catalog** (also called **data dictionary**) is used at this stage, which contains the information about data managed by the DBMS.

Example:

attr_name	rel_name	type	position
id	Person	integer	1
name	Person	string	2
age	Person	integer	3
address	Person	string	4
...

- Question:** Can the following query be accepted by the parser?

```
SELECT fname, lname FROM Person WHERE address<21;
```




Query Processing – Parser

- Consider the relation schema:

`Person(id:integer, name:string, age:integer, address:string)`

- Question:** Can the following query be accepted by the parser?

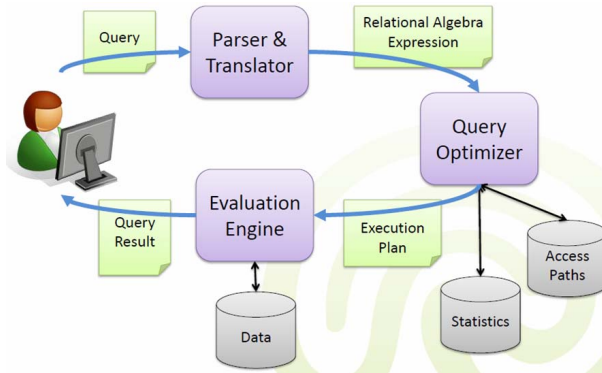
```
SELECT fname, lname FROM Person WHERE address<21;
```

- Answer:** The query **would be rejected** because

- The attributes `fname` and `lname` are not defined;
- The attribute `address` is not comparable with 21.

Query Processing – Translator

- The **translator** translates queries into RA expressions (not necessarily equivalent due to duplicates):
 - A query is first decomposed into **query blocks**.
 - Each query block is translated into an RA expression.





Recall: RA and SQL Queries

● RA operators

- **selection** σ_{φ}
- **projection** π_{A_1, \dots, A_n}
- **Cartesian product** $R_1 \times R_2$
- **join** $R_1 \bowtie_{\varphi} R_2$ and $R_1 \bowtie R_2$
- **renaming** $\rho_{R(A_1, \dots, A_n)}$
- **union** $R_1 \cup R_2$
- **intersection** $R_1 \cap R_2$
- **difference** $R_1 - R_2$

● SQL statement

```
SELECT attribute_list
  FROM table_list
  [WHERE condition]
  [GROUP BY attribute_list
  [HAVING group_condition]]
  [ORDER BY attribute_list];
```

$\sigma_{\varphi}(R) \Leftrightarrow \text{SELECT } * \text{ FROM } R \text{ WHERE } \varphi;$

$\pi_{A_1, \dots, A_n}(R) \Leftrightarrow \text{SELECT DISTINCT } A_1, \dots, A_n \text{ FROM } R;$

$R_1 \times R_2 \Leftrightarrow \text{SELECT DISTINCT } * \text{ FROM } R_1, R_2;$

...

- Aggregate operations in SQL require extended RA expressions.



Recall: RA and SQL Queries

- Nested subqueries are decomposed into separate query blocks.
- **Example:**

```
SELECT Lname, Fname
FROM EMPLOYEE
WHERE Salary > (SELECT Salary
                FROM EMPLOYEE
                WHERE ssn=5);
```

Outer query block

```
SELECT Lname, Fname FROM EMPLOYEE WHERE
Salary > c
```

⇓ translated

$\pi_{Lname, Fname}(\sigma_{Salary > c}(EMPLOYEE))$

Inner query block

```
(SELECT Salary FROM EMPLOYEE WHERE
ssn=5)
```

⇓ translated

$\pi_{Salary}(\sigma_{ssn=5}(EMPLOYEE))$



Query Processing – Query Optimiser

- 1 Transform into the best possible execution plan

There are different possible relational algebra expressions for a single query!

(will be covered in this course)

- 2 Specify the implementation of each operator in the execution plan

There are different possible implementations for a relational algebra operator!

(will not be covered in this course)



Query Processing – Query Optimiser

- SQL queries only specify **what data to be retrieved** and **not how to retrieve data**.
- There are **many possible execution plans** for a SQL query.
- Query optimiser is responsible for identifying **an efficient execution plan**:
 - 1 enumerating alternative plans (typically, a subset of all possible plans);
 - 2 choosing the one with the least estimated cost.
- Query optimisation is one of the most important tasks of a relational DBMS.
A good DBMS must have a good query optimiser!



Equivalent RA Expressions

- Suppose that we have:

Students(matNr, firstName, lastName, email)

Exams(matNr, crsNr, result, semester)

Courses(crsNr, title, unit)

```
SELECT lastName, result, title
FROM STUDENTS, EXAMS, COURSES
WHERE STUDENTS.matNr=EXAMS.matNr AND
      EXAMS.crsNr=COURSES.crsNr AND result≤1.3;
```

- **Question:**

How many equivalent RA expressions for this SQL query can you find?



Equivalent RA Expressions

Students(matNr, firstName, lastName, email)

Exams(matNr, crsNr, result, semester)

Courses(crsNr, title, unit)

```
SELECT lastName, result, title
```

```
FROM STUDENTS, EXAMS, COURSES
```

```
WHERE STUDENTS.matNr=EXAMS.matNr AND
```


```
EXAMS.crsNr=COURSES.crsNr AND result≤1.3;
```

● Answer:

- 1 $\pi_{lastName, result, title}(\sigma_{result \leq 1.3}((\text{Students} \bowtie_{\text{Students.matNr}=\text{Exams.matNr}} \text{Exams}) \bowtie_{\text{Exams.crsNr}=\text{Courses.crsNr}} \text{Courses}))$
- 2 $\pi_{lastName, result, title}(\sigma_{result \leq 1.3}(\sigma_{\text{EXAMS.crsNr}=\text{Courses.crsNr}}(\sigma_{\text{STUDENTS.matNr}=\text{Exams.matNr}}(\text{Students} \times \text{Exams} \times \text{Courses}))))$
- 3 $\pi_{lastName, result, title}((\text{Students} \bowtie_{\text{Students.matNr}=\text{Exams.matNr}} (\sigma_{result \leq 1.3}(\text{Exams}))) \bowtie_{\text{Exams.crsNr}=\text{Courses.crsNr}} \text{Courses})$

Query Trees

- Each RA expression can be represented as a **query tree**:
 - leaf nodes** represent the input relations;
 - internal nodes** represent the intermediate result;
 - the root node** represents the resulting relation.
- Example:**

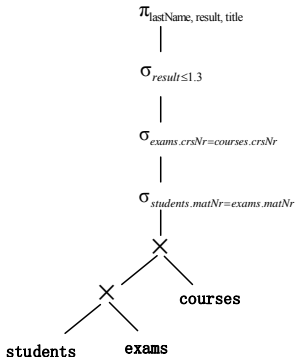
$$\pi_{lastName, result, title}(\sigma_{result \leq 1.3}((Students \bowtie_{Students.matNr=Exams.matNr} Exams) \bowtie_{\sigma_{Exams.crsNr=Courses.crsNr}} Courses))$$


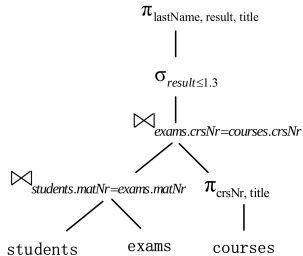


Query Trees

- Exercise:** Can you draw the query tree for the following RA expression?

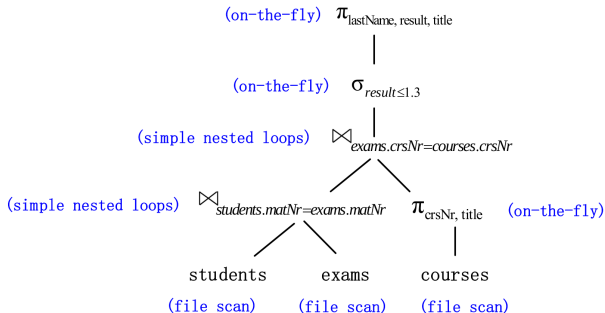
$\pi_{lastName, result, title}(\sigma_{result \leq 1.3}(\sigma_{Exams.crsNr=Courses.crsNr}(\sigma_{Students.matNr=Exams.matNr}(Students \times Exams \times Courses))))$





Execution Plan

- A **query execution plan** consists of an (extended) query tree with additional annotation at each node indicating:
 - (1) the *access method* to use for each table, and
 - (2) the *implementation method* for each RA operator.



Query Processing – Evaluation Engine

- The **evaluation engine** executes an execution plan, and returns the query answer to the user.

