



Database Transactions – Part 1

Introduction



Transaction – A Classical Example

- **Scenario:** Suppose that Steve's account balance is **\$1000** and Bob's balance is **\$200**. Now Steve wants to transfer **\$500** into Bob's account.
- There are several steps involved in transferring the money:
 - 1 **Check** Steve's balance;
 - 2 **Update** Steve's balance;
 - 3 **Check** Bob's balance;
 - 4 **Update** Bob's balance.
- Steve later checked his balance (it was **\$500**), which looked good to Steve. However, Bob told Steve that he hadn't received his money yet (still **\$200** in Bob's account instead of **\$700**).

Question: What did happen?



Transaction – A Classical Example

- **Reason:** Due to power outage, the system **stopped working just after updating** Steve's balance.
- **Task:** Transfer **\$500** from Steve's account to Bob's account

1 SELECT balance FROM ACCOUNT
WHERE name = 'Steve';

2 UPDATE ACCOUNT
SET balance = balance-500
WHERE name='Steve';

3 SELECT balance FROM ACCOUNT
WHERE name = 'Bob';

4 UPDATE ACCOUNT
SET balance = balance+500
WHERE name = 'Bob';

Operations	Steve	Bob
before 1	\$1000	\$200
after 1	\$1000	\$200
after 2	\$500	\$200
after 3	\$500	\$200
after 4	\$500	\$700

Transaction – A Classical Example

- We need an approach to ensure that
 - either the balances of Steve and Bob remain unchanged **if the money transfer fails**
 - or Steve's balance is **\$500** and Bob's is **\$700** **if the money transfer succeeds.**

1 SELECT balance FROM ACCOUNT
WHERE name = 'Steve';

2 UPDATE ACCOUNT
SET balance = balance-500
WHERE name='Steve';


3 SELECT balance FROM ACCOUNT
WHERE name = 'Bob';

4 UPDATE ACCOUNT
SET balance = balance+500

Operations	Steve	Bob
before 1	\$1000	\$200
after 1	\$1000	\$200
after 2	\$500	\$200
after 3	\$500	\$200
after 4	\$500	\$700

What is a Transaction?

- DBMSs provide **transaction support** for solving this kind of problem.
- A **transaction** is a **sequence of database operations grouped together** for execution as a **logic unit** in a DBMS.
 - Different from an execution of a program outside the DBMS (e.g., a C program) in many ways!





What is a Transaction?

- Database applications often access a database **by transactions** rather than individual operations.
 - e.g., large databases and hundreds of concurrent users: banking, supermarket checkout, airline reservation, online purchasing, etc.
- **Why transactions?** They can **enforce data integrity** in the following situations:
 - multiple users may modify and share data at the same time;
 - transaction, system, and media failures may happen from time to time.
- **What does a transaction look like?**
 - `INSERT`, `SELECT`, `UPDATE`, `DELETE`, `BEGIN`, `COMMIT`, `ABORT` (`ROLLBACK`), etc. from a high-level language perspective;
 - `read`, `write`, `begin`, `commit`, `abort` at the internal process level.



Transaction – Language Level

- **Database operations** of a transaction (at the SQL language level) may include: **SELECT**, **INSERT**, **UPDATE**, **DELETE**.
- **Other operations**: **BEGIN**, **COMMIT**, **ABORT** (**ROLLBACK**)

BEGIN TRANSACTION

- 1 `SELECT balance FROM ACCOUNT WHERE name = 'Steve';`
- 2 `UPDATE ACCOUNT`
`SET balance = balance-500 WHERE name='Steve';`
- 3 `SELECT balance FROM ACCOUNT WHERE name = 'Bob';`
- 4 `UPDATE ACCOUNT`
`SET balance = balance+500 WHERE name = 'Bob';`

COMMIT



Transactions - Internal Process Level

- **Basic operations** of a transaction (at the internal process level) are
 - **read(*X*)**: loads object *X* into main memory;
 - **write(*X*)**: modifies in-memory copy of object *X* (and writes it to disk later on);
- **Granularity of objects**: tables, rows, cells, or memory pages,
- **Other operations**:
 - **begin**: marks the beginning of a transaction;
 - **commit**: signals a successful end of the transaction - all changes can safely be applied to the database permanently;
 - **abort**: signals the transaction has ended unsuccessfully - undo all operations of the transaction.



Transactions - Internal Process Level

```
T: BEGIN TRANSACTION
T: SELECT balance FROM ACCOUNT WHERE name = 'Steve';
T: UPDATE ACCOUNT SET balance = balance-500 WHERE name='Steve';
T: SELECT balance FROM ACCOUNT WHERE name = 'Bob';
T: UPDATE ACCOUNT SET balance = balance+500 WHERE name = 'Bob';
T: COMMIT;
```

Objects:

- A - Steve's account balance;
- B - Bob's account balance.

Steps	<i>T</i>
1	read(A)
2	write(A) ($A := A - 500$)
3	read(B)
4	write(B) ($B := B + 500$)
5	commit