




Database Transactions – Part 3

Concurrent Transactions


Concurrent Transactions

- **Interleaved processing:** transactions are interleaved in a single CPU.



Concurrent Transactions

- **Parallel processing:** transactions are executed in parallel in multiple CPUs.





Concurrent Transactions

- Executing transactions concurrently will **improve database performance**
 - ↪ **Increase throughput** (*average number of completed transactions*)
 - For example, while one transaction is waiting for an object to be read from disk, the CPU can process another transaction (because I/O activity can be done in parallel with CPU activity).
 - ↪ **Reduce latency** (*average time to complete a transaction*)
 - For example, interleave execution of a short transaction with a long transaction usually allows the short one to be completed more quickly.
- But the DBMS has to guarantee that the interleaving of transactions **does not lead to inconsistencies**, i.e., **concurrency control**.



Why is Concurrency Control Needed?

- Concurrency control is needed for preventing the following problems:
 - 1 The **lost update** problem
 - 2 The **dirty read** problem
 - 3 The **unrepeated read** problem
 - 4 The **phantom read** problem



(1) - The Lost Update Problem

- Example:** Bob withdraws **\$100** from his account (T_1) while Alice deposits **\$500** into Bob's account (T_2).

```
 $T_1$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_2$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_1$ : UPDATE ACCOUNT SET balance=balance-100 WHERE name='Bob';  
 $T_1$ : COMMIT;  
 $T_2$ : UPDATE ACCOUNT SET balance=balance+500 WHERE name='Bob';  
 $T_2$ : COMMIT;
```

Steps	T_1	T_2
1	read(B)	read(B)
2		
3	write(B) ($B := B - 100$)	
4	commit	write(B) ($B := B + 500$)
5		
6		

Steps	B(Bob)
before 1	\$200
after 2	\$200
after 4	\$100
after 6	\$700

- Question:** What is the problem?



(1) - The Lost Update Problem

- Example:** Bob withdraws **\$100** from his account (T_1) while Alice deposits **\$500** into Bob's account (T_2).

```
 $T_1$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_2$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_1$ : UPDATE ACCOUNT SET balance=balance-100 WHERE name='Bob';  
 $T_1$ : COMMIT;  
 $T_2$ : UPDATE ACCOUNT SET balance=balance+500 WHERE name='Bob';  
 $T_2$ : COMMIT;
```


Steps	T_1	T_2
1	read(B)	read(B)
2		
3	write(B) ($B := B - 100$)	
4	commit	write(B) ($B := B + 500$)
5		
6		

Steps	B(Bob)
before 1	\$200
after 2	\$200
after 4	\$100
after 6	\$700

- Answer:** Bob's balance should be **\$600**. The update by T_1 is lost!

(1) - The Lost Update Problem

- Occurs when two transactions update the same object, and one transaction could overwrite the value of the object which has already been updated by another transaction (**write-write conflicts**).
- Example:**



- $write(B)$ by T_2 overwrites B , and the update by T_1 is *lost*.



(2) - The Dirty Read Problem

- Example:** Bob withdraws **\$100** from his account (T_1) while Alice deposits **\$500** into Bob's account (T_2).

```
 $T_1$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_1$ : UPDATE ACCOUNT SET balance=balance-100 WHERE name='Bob';  
 $T_2$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_1$ : ABORT;  
 $T_2$ : UPDATE ACCOUNT SET balance=balance+500 WHERE name='Bob';  
 $T_2$ : COMMIT;
```

Steps	T_1	T_2
1	read(B)	read(B) write(B) (B:=B+500) commit
2	write(B) (B:=B-100)	
3	abort	
4		
5		
6		

Steps	B(Bob)
before 1	\$200
after 1	\$200
after 2	\$100
after 4	\$200
after 6	\$600

- Question:** What is the problem?

(2) - The Dirty Read Problem

- Example:** Bob withdraws **\$100** from his account (T_1) while Alice deposits **\$500** into Bob's account (T_2).

```

 $T_1$ : SELECT balance FROM ACCOUNT WHERE name='Bob';
 $T_1$ : UPDATE ACCOUNT SET balance=balance-100 WHERE name='Bob';
 $T_2$ : SELECT balance FROM ACCOUNT WHERE name='Bob';
 $T_1$ : ABORT;
 $T_2$ : UPDATE ACCOUNT SET balance=balance+500 WHERE name='Bob';
 $T_2$ : COMMIT;
  
```


Steps	T_1	T_2
1	read(B)	read(B)
2	write(B) ($B := B - 100$)	
3		
4	abort	
5		write(B) ($B := B + 500$)
6		commit

Steps	B(Bob)
before 1	\$200
after 1	\$200
after 2	\$100
after 4	\$200
after 6	\$600

- Answer:** Bob's balance should be **\$700** since T_1 was not completed.

(2) - The Dirty Read Problem

- Occurs when one transaction could read the value of an object that has been updated by another transaction but has not yet committed (**write-read conflicts**).
- Example:**



- T_1 fails and must change the value of B back to **\$200**; but T_2 has read the uncommitted (\cong *dirty*) value of B (**\$100**).



(3) - The Unrepeatable Read Problem

- Example:** Bob checks his account (T_1) twice (takes time to decide whether to withdraw **\$200**) while Alice withdraws **\$500** from Bob's account (T_2).

```
T1: SELECT balance FROM ACCOUNT WHERE name='Bob';  
T2: SELECT balance FROM ACCOUNT WHERE name='Bob';  
T2: UPDATE ACCOUNT SET balance=balance-500 WHERE name='Bob';  
T2: COMMIT;  
T1: SELECT balance FROM ACCOUNT WHERE name='Bob';
```

Steps	T_1	T_2
1	read(B)	
2		read(B)
3		write(B) (B:=B-500)
4		commit
5	read(B)	

Steps	B(Bob)
before 1	\$500
after 2	\$500
after 3	\$0
after 4	\$0
after 5	\$0

- Question:** What is the problem?



(3) - The Unrepeatable Read Problem

- Example:** Bob checks his account (T_1) twice (takes time to decide whether to withdraw **\$200**) while Alice withdraws **\$500** from Bob's account (T_2).

```
 $T_1$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_2$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_2$ : UPDATE ACCOUNT SET balance=balance-500 WHERE name='Bob';  
 $T_2$ : COMMIT;  
 $T_1$ : SELECT balance FROM ACCOUNT WHERE name='Bob';
```

Steps	T_1	T_2
1	read(B)	
2		read(B)
3		write(B) (B:=B-500)
4		commit
5	read(B)	

Steps	B(Bob)
before 1	\$500
after 2	\$500
after 3	\$0
after 4	\$0
after 5	\$0


- Answer:** Bob received two different account balances **\$500** and **\$0**, even though he hasn't withdrawn any money yet.



(3) - The Unrepeatable Read Problem

- A transaction could change the value of an object that has been read by another transaction but is still in progress (could issue two read for the object, or a write after reading the object) (**read-write conflicts**).

- **Example:**





(4) - The Phantom Read Problem

- **Example:** A query is submitted for finding all customers whose account balances are less than **\$300** (T_1) while Alice is opening a new account with the balance **\$200** (T_2).
- Assume that only Bob (B) has an account whose balance is less than **\$300** before Alice (A) opens his new account.

T_1 : SELECT name FROM ACCOUNT WHERE balance<300;

T_2 : INSERT INTO ACCOUNT(id, name, balance) VALUES(99, 'Alice', 250);

T_2 : COMMIT;

T_1 : SELECT name FROM ACCOUNT WHERE balance<300;

Steps	T_1	T_2
1	read(R)	
2		write(R)
3		commit
4	read(R)	

Steps	Query result
before 1	$R = \{B\}$
after 1	$R = \{B\}$
after 2	$R = \{A, B\}$
after 4	$R = \{A, B\}$

- **Question:** What is the problem?

(4) - The Phantom Read Problem

- **Example:** A query is submitted for finding all customers whose account balances are less than **\$300** (T_1) while Alice is opening a new account with the balance **\$200** (T_2).
- Assume that only Bob (B) has an account whose balance is less than **\$300** before Alice (A) opens his new account.

```

T1: SELECT name FROM ACCOUNT WHERE balance<300;
T2: INSERT INTO ACCOUNT(id, name, balance) VALUES(99, 'Alice', 250);
T2: COMMIT;
T1: SELECT name FROM ACCOUNT WHERE balance<300;

```

Steps	T_1	T_2
1	read(R)	
2		write(R)
3		commit
4	read(R)	

Steps	Query result
before 1	$R = \{B\}$
after 1	$R = \{B\}$
after 2	$R = \{A, B\}$
after 4	$R = \{A, B\}$

- **Answer:** T_1 reads Account based on the condition $\text{balance} < 300$ twice but gets two different results $\{B\}$ and $\{A, B\}$.

(4) - The Phantom Read Problem

- Occurs when tuples updated by a transaction T_1 satisfy the search conditions of another transaction so that, by the same search condition, the transaction obtains different results at different times.
- Example:**

