

Week 5 Announcements

- Reminder from last week (W4):
 - Survey 1 (weeks 4 and 5).
 - Pointing to W3 reminders:
 - Video Assignment #1 → no late submission (hard deadline).
 - Group Project.
 - UML Exercise.
 - Control Flow Graph Exercise.
- Industry talks: 18 & 25 September.
- Android content is now available (in prep for W6 labs).
- COMP2100 are you okay? (zero e-mails) 😬 😞
- How are you doing? Is everything okay?
 - If you consider yourself at risk, please reach out!

COMP2100/6442

Software Design Methodologies / Software Construction

Tree Data Structures



Continuation...

Bernardo Pereira Nunes and
Sergio J. Rodríguez Méndez

Outline

- ~~Data Structures~~
- ~~Tree Data Structures~~
 - B-Tree
- Exercises



B-Tree

B-Tree

Bayer and McCreight never explained what, if anything, the *B* stands for: *Boeing*, *balanced*, *between*, *broad*, *bushy*, and *Bayer* have been suggested.^{[4][5][6]} McCreight has said that "the more you think about what the B in B-trees means, the better you understand B-trees."^[5]

The **B** stands for Bayer*

> Generalization of Binary Search Trees

- Not binary trees
- Overcome BST Limitation: each node is read individually (and access to secondary memory is slow)


> Designed for searching data stored on block-oriented devices

- Map tree nodes into blocks

Finally, a balanced tree!

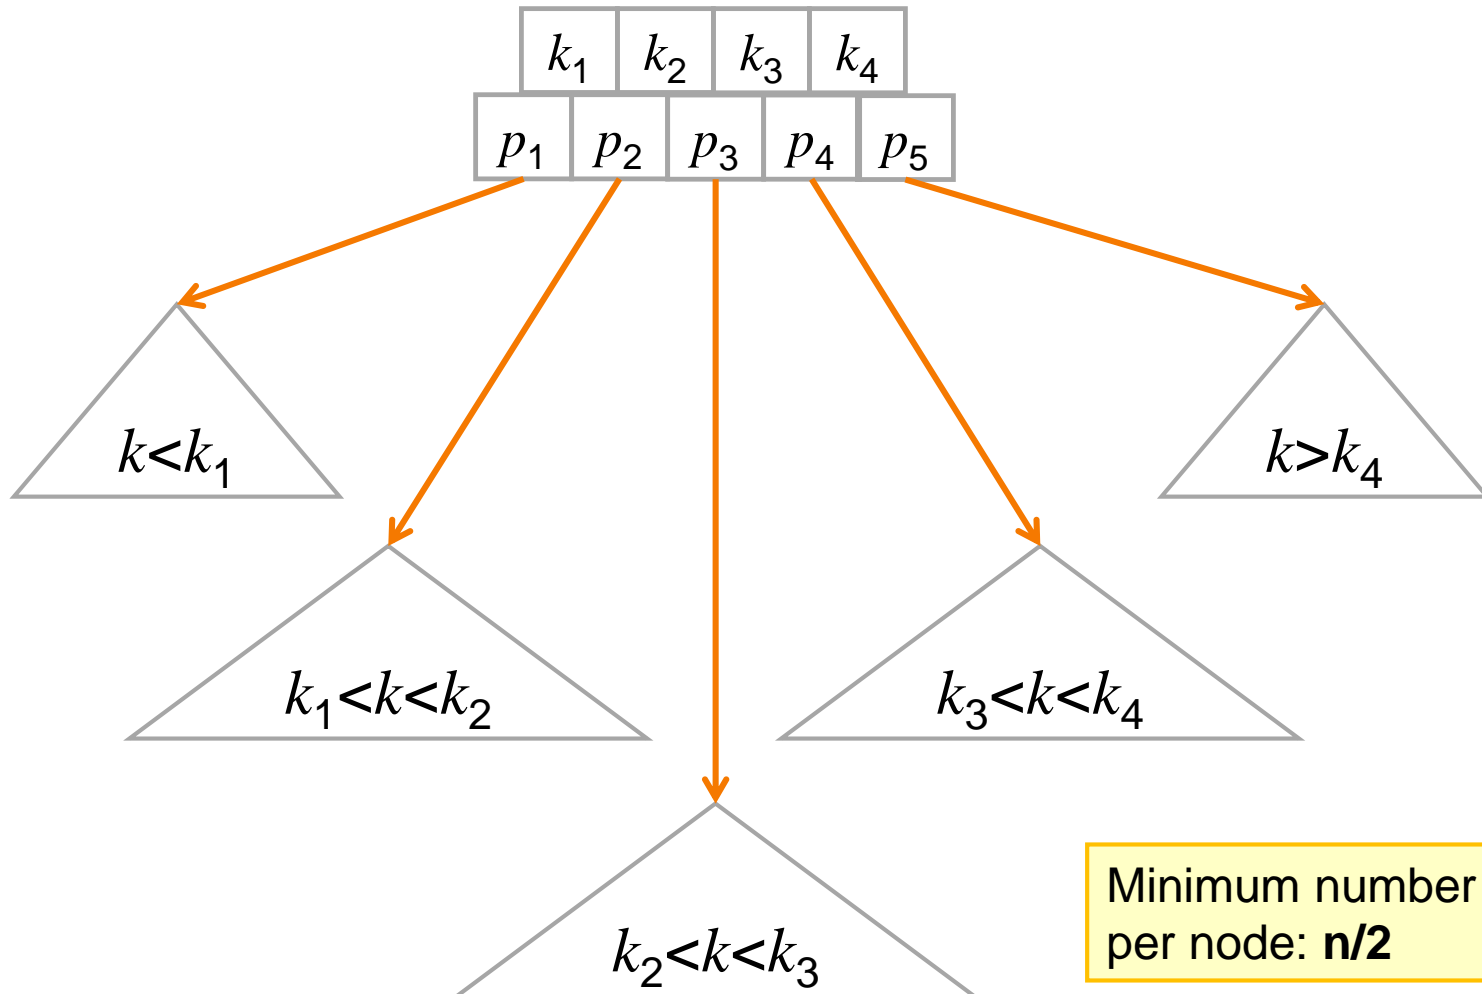
B-Tree

- > Each access to secondary memory brings a group of elements
- > Subtrees are divided into blocks (**pages**)



B-Tree: General Idea

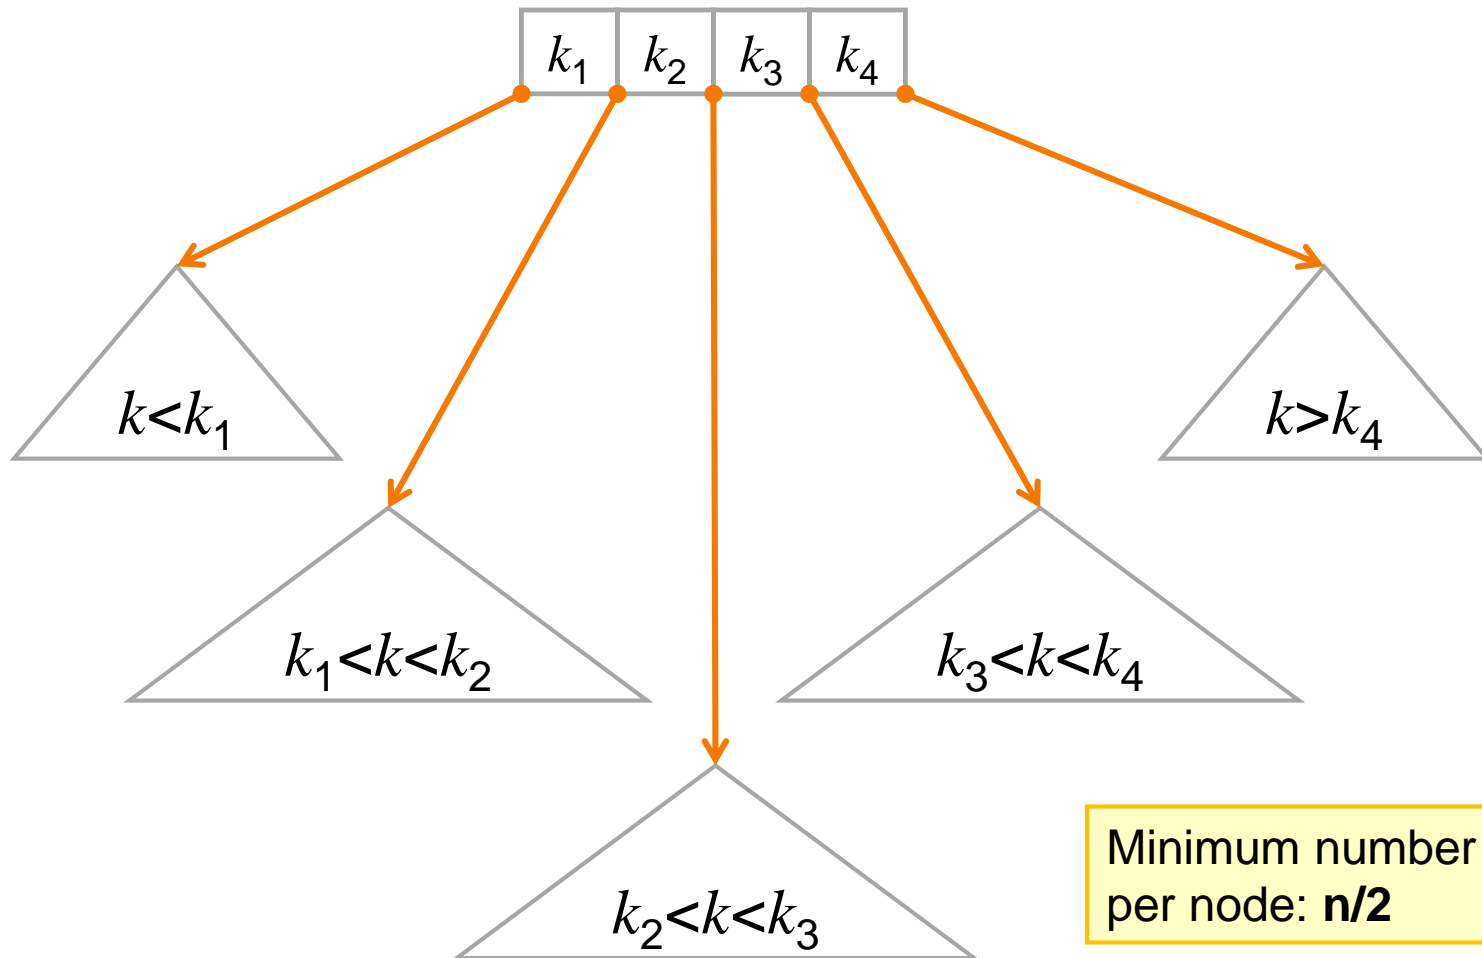
N-ary tree with search keys on nodes



Minimum number of keys
per node: $n/2$

B-Tree: General Idea

N-ary tree with search keys on nodes




Minimum number of keys
per node: $n/2$

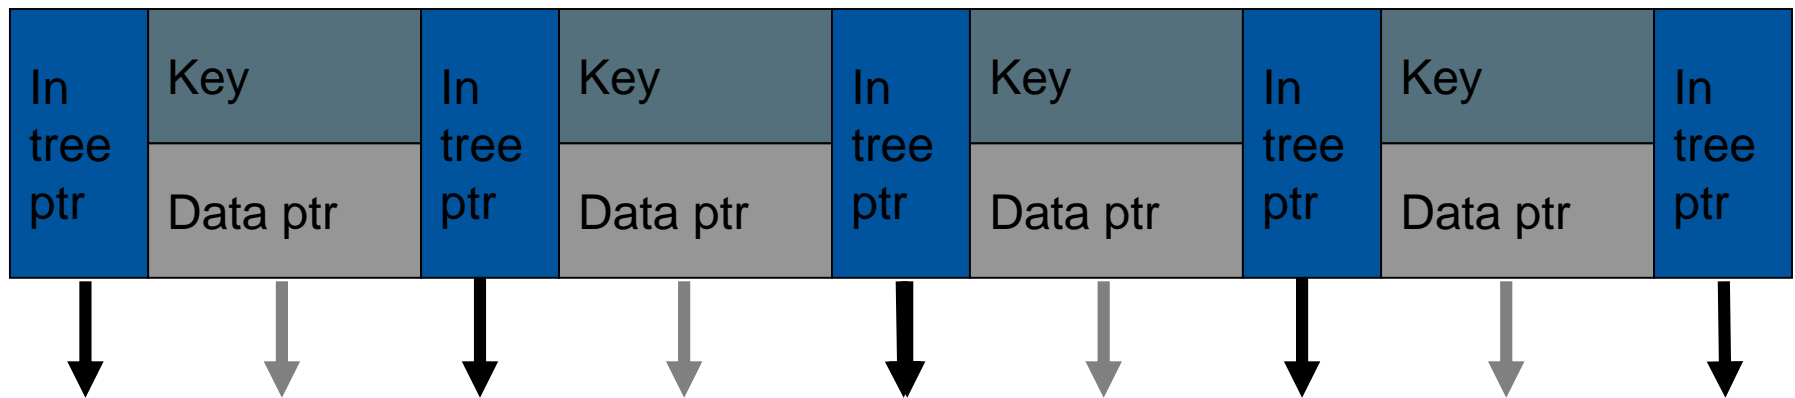
B-Tree: Definition (Knuth, 1997)

A B-tree of order m satisfies the following properties:

1. Every node (page) has at most m children
2. Each node (except the root and leaves) has at least $\lceil m/2 \rceil$ children
3. The root has at least 2 children (unless it is a leaf)
4. A non-leaf node with k children contains $k-1$ keys
5. All leaves appear in the same level



In reality



B-Tree: Definition (Knuth, 1997)

> Minimum number of keys of a B-Tree of order m and height h

>> The root has at least 1 key and 2 children

>> All internal nodes have at least $\lceil m/2 \rceil$ children and $\lceil m/2 \rceil - 1$ keys

>> All leaves have at least $(\lceil m/2 \rceil - 1)$ keys

>> Therefore, the minimum number of keys will be:

$$k = 2 * \lceil m/2 \rceil^h - 1$$

B-Tree: Definition (Knuth, 1997)

> Minimum number of keys of a B-Tree of order m and height h

root has 2 children
(minimum)

$$k = 1 + (\lceil m/2 \rceil - 1) * (2 + 2 * \lceil m/2 \rceil + 2 * \lceil m/2 \rceil^2 + \dots + 2 * \lceil m/2 \rceil^{(h-1)})$$

root has 1 key
(minimum)

$$= 1 + (\lceil m/2 \rceil - 1) * 2 * (1 + \lceil m/2 \rceil + \lceil m/2 \rceil^2 + \dots + \lceil m/2 \rceil^{(h-1)})$$

each children can
have $m/2$ children
(minimum)

$$= 1 + (\lceil m/2 \rceil - 1) * 2 * (\sum_{i=0}^{h-1} \lceil m/2 \rceil^i)$$

$$= 1 + (\lceil m/2 \rceil - 1) * 2 * \frac{\lceil m/2 \rceil^h - 1}{\lceil m/2 \rceil - 1}$$

minimum number
of keys

$$= 1 + 2 * (\lceil m/2 \rceil^h - 1)$$

$$= 2 * (\lceil m/2 \rceil^h) - 1$$

B-Tree: Search


a. Search between keys on a page

$k_1 \dots k_{m-1}$ (if m is large: binary search)

b. if not found on the page:

1. $x < k_1$: search must continue on page p_0
2. $k_i < x < k_{i+1}$ for $1 \leq i < m-1$: search must continue on page p_i
3. $k_{m-1} < x$: search must continue on page p_{m-1}

if there are no pages below the current one, the key does not exist



B-Tree: Insertion

Let p_i be the page where x should be inserted

if p_i has less than $m-1$ elements then

insert in p_i , in the proper position

if page p_i is already full then

1. allocates a new page p_k


2. distributes the keys as follows:

1. $\lceil m/2 - 1 \rceil$ smallest keys in p_i



2. $m - \lceil m/2 \rceil$ biggest keys in p_k

3. insert the median key (in $\lceil m/2 \rceil$) on the top page


(if page p_i is root: create new root with median)






Insert 10, 30, 50, 70, 90 (order 5)




Overflow > Split





Insert 20, 40 (order 5)




Insert 60, 80 (order 5)



Insert 100 (order 5)



Overflow > Split



Exercise

Draw a B-Tree step-by-step inserting the keys in the following order: **20, 11, 15, 3, 5, 7, 12, 15, 16, 19, 25, 30, 33, 37, 22, 23, 26, 31, 42, 35, 6, 47**

Use different orders $m=\{3, 5\}$


What is the height of the tree?

<https://www.cs.usfca.edu/~galles/visualization/BTree.html>

B-Tree: Remove


Removal must be performed on a **leaf node**

1. If the key to be removed is not in a leaf node then
 replace it with the **largest key in its left subtree (predecessor)**
 or the **smallest key in its right subtree (successor)**



B-Tree: Remove


2. if the key is on a leaf then
it must be removed;
if the leaf is **less than $m/2$ keys** then
a **merge** or a **rebalancing** must be performed.



B-Tree: Remove

Merge:

- > if, after removal, the page where the key was removed and its adjacent page have together less than m keys then
 - >> This page is merged with its adjacent one. The parent key that was between them goes to the page that was merged.
- > if the resulting page has less than $m/2$ keys then
 - >> This procedure is repeated until it reaches the root.




B-Tree: Remove

Rebalancing:

> if, after removal, the page where the key was removed and its adjacent page together have m keys or more then



>> Move the parent page key (the one “between” adjacent pages) to the page with less keys; then, move its adjacent page* key to the parent page.

> There is no propagation as the number of parent keys does not change.




* if the adjacent page is to the left of the page with less keys, the moved key is the largest on that page (*borrow from left*). If the adjacent page is to the right, the key moved is the smallest of that page (*borrow from right*).

B-Tree: Remove 70 (order 5)






merge
(# of keys < order)




B-Tree: Remove 20 (order 5)

after removing 30...



Meme for today's lecture! Keep practicing!



It seems a good idea... no names on gifts anymore.

References

- Chapter 14 and 19 of Introduction to Algorithms (by Thomas H. Cormen, Charles E. Leiserson, and Ronald L. Rivest)
- Chapter 10 of Lee K.D., Hubbard S. (2015) Balanced Binary Search Trees. In: Data Structures and Algorithms with Python. Undergraduate Topics in Computer Science. Springer, Cham.
- Knuth, Donald (1998), Sorting and Searching, The Art of Computer Programming, Volume 3 (Second ed.), Addison-Wesley, ISBN 0-201-89685-0.